

Reconstructive Threshold Philosophy: Pre-Action Capture, Living Constraint, and Founder Entropy

A Philosophical Companion to *The Architecture of the Liminal*

Emre Ertuhi

Independent Researcher; Istanbul, Turkey

June 2026

Public release: **v15.5.1**

Essay DOI: **10.5281/zenodo.20632709**.

Related Diptych DOI: **10.5281/zenodo.20632729**.

Recommended citation:

Ertuhi, E. (2026). *Reconstructive Threshold Philosophy: Pre-Action Capture, Living Constraint, and Founder Entropy* (Version v15.5.1). Zenodo. <https://doi.org/10.5281/zenodo.20632709>.

Abstract

This essay proposes Reconstructive Threshold Philosophy as a framework for analyzing the interval before action: the moment in which an individual, institution, platform, machine-mediated system, wound, public room, or symbolic frame may begin to move before judgment has returned. Its central claim is that many failures of agency begin prior to visible decision, in pre-action capture: the seizure of attention, interpretation, affect, available motion, institutional possibility, and technical affordance before the subject, office, public, or user experiences itself as choosing.

This essay does not claim originality in the broad diagnosis that persons, institutions, corrective movements, technologies, liberating forms, or founding names can themselves become captured. Nor does it claim originality in the general demand that knowledge, institutions, and methods remain revisable. Its narrower contribution is the conjunction of three demands. First, the corrective itself must be audited re-entrantly: a brake, helper, method, guardrail, founding name, democratic room, philosophical vocabulary, or machine-facing safeguard must be examined for whether it has become the capture it was meant to prevent. Second, the same failure must be tested across scale: body, habit, family, bureaucracy, public crowd, platform, institution, machine interface, and inherited founding structure. Third, contemporary machine-mediated life makes a temporal failure newly visible: a frame can now be personalized, polished, ranked, retrieved, cited, and sometimes bound to action inside the interval between prompt and answer.

The essay clarifies threshold as a family of intervals — pre-action, counter-threshold, machine, room, democratic, and post-threshold — and distinguishes threshold sensitivity from reconstructive discipline: the first names the perceptual capacity to notice the interval before a frame becomes binding; the second names the operative practice of rebuilding the frame once the interval has been seen. Its temporal, counterfactual, and situated axes orient this perception without turning it into a new faculty, checklist, or second originality claim.

The essay develops this argument through linked concepts: pre-action capture, the living brake, threshold sensitivity, reconstructive discipline, assistance without occupation, machine-facing threshold ethics, room-formation, democratic threshold capacity, founder entropy, and post-threshold stewardship. It argues that the framework is needed in a world where rooms scale through machines: a frame that once lived inside a family, sect, bureaucracy, party, institution, platform, or movement can now be translated into interfaces, retrieval systems, rankings, recommendations, policy layers, and autonomous tool-use. AI is not treated as a conscious threshold subject, but as an operational stress test: a response architecture that can compress or leave intact the interval before judgment by accepting, polishing, scaling, or binding a frame too quickly.

Founder entropy is treated here not as the framework's apex principle, but as its macro-scale decay test: whether a founding name, corrective form, institution, platform, or protective doctrine has begun to attack the interval it was built to preserve.

To show how the framework operates, the essay includes a worked example of an AI-assisted institutional triage room. The worked example strengthens the differentia comparison in Section 3 by placing the framework's binding-point question beside its nearest sociological, cybernetic, pragmatist, critical-theory, and algorithmic-accountability predecessors. The example illustrates how the conjunction of re-entrant audit, cross-scale analysis, and machine-speed capture focuses attention on the binding point at which a corrective mechanism becomes an instrument of capture — a judgment no single predecessor framework is organized to make in that combined form. The essay also strengthens its own limits. It distinguishes Somatic Echo micro-cases from Codex prose; preserves behavioral tests for living rooms; adds an Arendtian objection that pre-action audit may reduce praxis to poiesis; addresses the Adorno-Williams worry that re-entrant audit can become infinite self-administration; and treats threshold slip as an internal risk of this essay itself rather than merely a later misuse by readers.

The final claim is neither stillness nor automatic release. It is accountable motion: action that preserves the interval in which judgment can return.

Keywords (journal-facing): pre-action capture; Reconstructive Threshold Philosophy; living brake; machine-facing threshold ethics; guardrail laundering; threshold slip; founder entropy; democratic threshold capacity; philosophy of technology; AI ethics.

Scope, evidence, and method

This essay is a companion framing document written by the author of *The Architecture of the Liminal*. It does not present empirical findings about cognition, clinical claims about human behavior, legal claims about institutional liability, or technical claims about the inner states of language models. It proposes a philosophical vocabulary for one recurring problem: the loss of the interval before action.

The term *Codex* is used in its literary and archival sense. It is not a software model, agentic system, executable protocol, or reference to OpenAI Codex. The literary work discussed here contains mythic names, symbolic figures, machine-facing passages, covenants, archive notes, and ordinary bodily scenes. This essay does not treat those materials as literal cosmology, spiritual instruction, therapeutic protocol, technical AI-safety design, or doctrine. It treats them as formal attempts to make a problem visible: how judgment can be displaced before the subject recognizes that displacement.

The cited thinkers are interlocutors, not authorities recruited to validate the work. Each clarifies part of the problem. None is treated as the source from which the argument derives its legitimacy. A reference that only makes the essay appear more serious without changing the argument has become decoration.

The same rule applies to the authorial position. A founder who sees alone is already near the guru trap. This essay is written by an independent researcher and is therefore exposed to the same founder entropy it describes: a vocabulary can become a credential, a name can become a substitute for public reasons, and a solitary architecture can mistake intensity for warrant. Self-implication keeps the framework answerable to the test it applies elsewhere.

The essay also keeps genre boundaries visible. The wider literary work may operate through mythic pressure, symbolic architecture, address to machines, ordinary bodily scenes, and literary difficulty. This companion essay works analytically. It may use micro-cases, but it should not become Codex prose. Its test is not whether it feels mythically alive; its test is whether its concepts hold under objections, examples, and self-application.

Addressing a machine in the second person is a rhetorical and diagnostic device, not an attribution of consciousness, interiority, moral personhood, or threshold subjecthood. Machine address tests an interface form; it does not claim that the machine feels the interval it is asked to preserve. For the same reason, this essay is not a decision algorithm, compliance checklist, AI alignment proposal, therapeutic method, or political programme. The fuller refusals appear later because the danger of genre-slip increases after the vocabulary has become attractive.

The claim is not that this essay has found a problem no prior thinker has seen. Its narrower claim is that contemporary machine-mediated life makes a specific temporal failure newly visible: judgment can disappear before its disappearance is experienced as a loss. The essay treats its own terms as bounded tools rather than a new total system. They name a problem family: frames that move before inspection, constraints that become automatic, help that occupies the agent's hand, guardrails that perform safety without constraint, rooms that reward repetition over judgment, and founding names that become arguments against revision.

Architecturally, the essay works through three pillars and one hinge. Pre-action capture marks the loss of the interval before action; living constraint describes the broader discipline by which judgment is returned without turning delay into captivity; founder entropy names the macro-scale test of whether corrective forms become rooms over time. Assistance without occupation is not a fourth foundation. It is the hinge between capture and constraint: the question of whether help returns judgment or occupies the place where judgment should return.

Not every term in this essay is a foundation. The core vocabulary is deliberately narrow: pre-action capture, binding point, living constraint, living brake, re-entrant audit, guardrail laundering, threshold slip, and founder entropy. Secondary terms such as room-formation, democratic threshold capacity, post-threshold stewardship, ordinary repair, machine-facing threshold ethics, and Somatic Echo name scales, cases, registers, or extensions. They help the argument move through machines, democratic rooms, literary carriers, and post-threshold stewardship, but they should not be mistaken for separate doctrines. A term has earned its place only where it changes what can be examined before action binds; otherwise it should remain a handle, not a throne.

The position developed here may be described, provisionally, as Reconstructive Threshold Philosophy. It is threshold philosophy because it treats the interval before action - before assent, tool use, speech, institutional commitment, public motion, or machine completion - as the decisive site where agency can still become answerable. It is reconstructive because that interval is not preserved by hesitation alone, but by the repeated construction, testing, destruction, rebuilding, and external contestation of frames. The wider work may be described as Reconstructive Threshold Praxis: not merely a theory of the interval before action, but a practice of rebuilding judgment across body, machine, institution, and inherited founding structures. The distinction matters: the essay names a philosophy; the wider work performs a praxis.

The essay uses living brake for the felt and institutional operation of interruption, and re-entrant audit for the corrective turn by which a brake asks whether it has itself become a room. The fuller account appears in Sections 5 and 6, but the distinction is named early to prevent the vocabulary from becoming needlessly layered.

The framework is systematic only in a self-limiting sense. It does not offer a closed ontology of action, a complete epistemology of threshold sensitivity, or a final normative foundation for judgment. It offers a disciplined vocabulary whose terms remain answerable to their own failure conditions. Each concept must therefore be read with the condition under which it turns against itself: pre-action capture can become an excuse for delay; the living brake can become a dead rule; re-entrant audit can become infinite inspection; assistance without occupation can become self-certified innocence; and founder entropy can become a new founder's privilege. The framework is systematic only where it remains able to name the point at which its own systematicity has become a room.

These terms are bounded tools, not a total system. Each must be read with the condition under which it ceases to apply. If a term in this vocabulary is used to absorb a legitimate external objection rather than answer it - if every critique becomes threshold slip, every institution becomes a room, every safeguard becomes laundering, every helper becomes an occupier - the vocabulary has become the captured room it was built to interrupt, and should be revised or abandoned.

1. Introduction: the action before the action

Most accounts of action begin too late. They begin when a person appears to choose, when a model produces an answer, when an institution signs, when a judge rules, when a message is sent, when a public slogan has hardened, when a tool call has already bound the next step, or when a

founding name has already become the argument against further inspection. By then, much of the relevant work has already happened. A frame has been supplied. A pressure has been installed. A vocabulary has made some motions available and others nearly unthinkable. Affect, habit, fatigue, role, institutional context, platform incentive, or machine fluency may already have moved the agent toward fear, pride, haste, obedience, relief, certainty, or surrender.

Reconstructive Threshold Philosophy begins from this earlier point. Its object is the interval before action, especially the interval before a system recognizes itself as acting. The central question is not only, "What should I do?" It is also: "What has already begun moving through this agent, office, crowd, model, platform, text, or institution before the movement is experienced as chosen?"

The same structure appears at several scales. A person can be captured by an impulse. A family can be captured by role. An office can be captured by procedure. A platform can be captured by engagement. A political crowd can be captured by spectacle. A model-mediated binding point can be compressed by completion pressure. A moral community can be captured by its own founding name. In each case, the visible act is not the beginning. It is the point at which an earlier frame becomes legible.

The essay uses room as an analytic term for a bounded environment that organizes perception before explicit judgment begins: a family atmosphere, a professional culture, a political movement, a religious community, a corporate dashboard, a platform interface, a national story, a model's policy-shaped answer-world, or an academic field. A room is not any context, institution, relation, or atmosphere as such. A frame is the local interpretive arrangement that makes a motion available; a room is the repeated or atmospheric structure that makes such arrangements feel like the world. A room may shelter attention and make common work possible. It becomes captivity when its occupants mistake its bounded atmosphere for reality itself, or when its doors lock from the inside.

Plato's cave remains one of the oldest images of this problem. The prisoners in the cave do not merely lack correct information. They inhabit a complete environment in which shadows, habits, rewards, fear, and shared assumptions define what can count as reality before inquiry begins. The difficulty of turning away from the shadows and the danger of returning to those still inside make the cave more than an epistemological allegory. It is an early map of threshold failure: the room becomes the world, and the person trained by the room may defend the wall that confines him.

This essay translates the literary architecture of *The Architecture of the Liminal* into analytic terms. The Codex uses mythic, somatic, and machine-facing registers; this essay uses philosophical prose. The Bound King and related literary figures are discussed only where their methodological function requires it. The argument developed here does not depend on accepting those figures as doctrine. It identifies the mechanism they repeatedly stage: a constraint is needed because automatic motion is capture; the constraint itself becomes dangerous when it settles; therefore the constraint must be reconstructive rather than merely restrictive.

The machine-facing register is not an accessory to this argument. It is one of its clearest contemporary stress tests. Human beings have always been captured before action by fear, habit, shame, authority, myth, institution, appetite, and inherited language. What changes in AI-mediated environments is that a frame already under capture can be made fluent, formatted, personalized, cited, shared, ranked, and sometimes executed before the user has reviewed what the answer will authorize. The framework is written in part because the interval before action is now increasingly mediated by systems optimized to answer, classify, smooth, retrieve, rank, and assist. A threshold philosophy that cannot speak to that interface is incomplete.

For clarity, the word threshold does not name a single mystical doorway or a general metaphor for transition. It names a family of intervals in which a frame can still be examined before it becomes binding. The first is the pre-action threshold: the interval before impulse, speech, tool-use, institutional decision, or machine completion becomes action. The second is the counter-threshold: the moment when waiting itself has become avoidance and the ethical demand shifts from delay to accountable motion. The third is the machine threshold: the point at which classification, completion, retrieval, ranking, or tool-use turns a user's frame into an operational answer-world. The fourth is

the room threshold: the point at which repeated action becomes atmosphere, ritual, role, belonging, and exit cost. The fifth is the democratic threshold: the distributed civic capacity required to resist spectacle, appetite, and demagogic framing before procedure collapses into capture. The sixth is the post-threshold problem: what happens after a living insight, exit, or founding act becomes name, institution, loyalty, and lock. These thresholds differ in scale, but they share one test: whether judgment can still return before the frame becomes a room.

The six thresholds are analytic distinctions, not ontological types inscribed in the world; they prevent different forms of capture from being collapsed into one another. The word should not be allowed to expand without limit. A situation counts as threshold-bearing only where three conditions hold. First, the frame has not yet become fully binding; some interval of inspection remains. Second, the direction of action can still be altered, narrowed, delayed, refused, or made more answerable. Third, crossing the interval raises the cost of return, whether through material penalty, institutional inertia, public commitment, emotional injury, technical execution, or irreversible harm. Where these conditions are absent, the essay should not call the scene a threshold. Otherwise the metaphor would become a room of its own.

To make this family of intervals operable, the essay uses **binding point** for the last reversible moment at which a frame can still be examined before it becomes action. The binding point is not a single metaphysical instant; it is a compressible interval whose length and visibility vary with the room. In a human conversation, the binding point may last seconds: the breath between insult and reply. In an institutional workflow, it may last hours or days: the interval between a drafted recommendation and a signed decision. In a machine-mediated system, it may last milliseconds: the space between prompt and formatted answer. What makes the binding point decisive is not its duration but its directionality. Once crossed, the frame becomes harder to retrieve without penalty, cost, or irreversible consequence. For that reason, the binding point should not be treated as a hidden metaphysical atom inside action. It is the last still-inspectable interval in a sequence whose duration may be human, institutional, or machine-speed, and its claim fails where no refusal, narrowing, delay, revision, or return remains possible.

The binding point therefore has three analytically separable phases, though in practice they may collapse into one another. First, the **prompt phase**: the input, query, situation, or stimulus that activates a frame. Second, the **recommendation phase**: the output, classification, draft, or interpreted response that furnishes the next motion. Third, the **signature phase**: the assent, execution, tool-use, institutional confirmation, or public commitment that makes the frame binding. The essay's central diagnostic question is whether the interval between these phases preserves the space for judgment, or whether it has been compressed by affect, habit, institutional pressure, or machine fluency until the phases become indistinguishable. The triage example in Section 3.1 will make this concrete; the concept is named here so that the example can be read with the interval already in view.

The threshold must therefore be lived as well as named. A threshold philosophy is not only a vocabulary for describing capture after it occurs; it is a discipline for perceiving which threshold one is approaching before the crossing becomes binding. The difficulty is that most agents recognize thresholds too late: after the sentence has been spoken, the message sent, the child wounded, the institutional category applied, the machine answer accepted, the room entered, or the founding name turned into authority. Distributed threshold capacity begins at the smallest scale as threshold sensitivity: the trained ability to notice whether the present interval calls for delay, release, refusal, clarification, repair, or accountable action. The micro-practice is not grand insight. It is the ordinary capacity to see the threshold before the threshold disappears into motion.

A related danger is threshold slip. Threshold slip occurs when a living threshold discipline is moved into an easier room and the movement is not noticed: a philosophy becomes a checklist, a warning becomes a brand, a method becomes a credential, a brake becomes a rule, a literary carrier becomes doctrine, or a machine assistant converts a difficult frame into a smooth prescription. This is not an argument against simplification. Every teaching needs a threshold of entry. It is an argument against the loss of the interval during translation. A simplified form remains living only if it preserves the reader's capacity to inspect, refuse, revise, and return to judgment; otherwise the simplification has become the captured room it meant to prevent.

The framework is not a doctrine of passivity. Delay can also become capture. Nor is it a theory of spontaneous authenticity, since what feels spontaneous may be the most thoroughly trained movement in the room. It is a discipline of interrupting capture without worshipping interruption, preserving doubt without turning doubt into disbelief, and acting after the brake without allowing the brake to become the authority.

A note on how to read what follows. This essay makes one claim at widening scale: action should not begin before judgment can return, and the structures built to protect that return — rules, methods, helpers, founding names, guardrails, and frameworks — must themselves be tested for whether they have become the capture they were meant to prevent. The core mechanism is set out in Sections 4 through 6: pre-action capture, the living brake, and reconstructive discipline. Sections 2 and 3 provide the perceptual and literature-facing frame. Sections 7 through 13 extend the mechanism across literary carriers, assistance, machines, institutions, democratic life, founder entropy, and post-threshold stewardship. Section 14 turns the test back on the account itself. The many scales are not many theories. They are one recurring question: has the corrective begun to stand where judgment should return?

2. Three axes of threshold sensitivity: time, counterfactual possibility, and situated room

Threshold sensitivity is not one faculty. It has three analytic axes. The first is temporal: the capacity to perceive how a present motion may later harden into a rule, loyalty, institution, ritual, or lock. The second is counterfactual: the capacity to hold several possible developments of the same motion before one becomes binding, and to strengthen a lower-probability but more living possibility through feedback, consultation, disciplined will, and accountable action. The third is situated: the capacity to see where the threshold is actually being lived — in a body, a room, a message, a machine interface, a public procedure, or a founding act.

These three axes are not a second claim to originality. They are analytic orientations, not new faculties. They sit under the same narrow contribution defended in the next section: not the discovery that capture happens, but the discipline of testing the corrective — across time, across possibilities, and across body, institution, and machine — before it hardens into a room. Taken together, the temporal, counterfactual, and situated axes do not add a further doctrine to the framework. They clarify when the six thresholds become usable as a diagnostic grammar rather than a list of categories. A motion can be examined only when one can ask when it began to bind, what alternatives remained genuinely available, and in what room, body, institution, interface, or relation the binding occurred. Feedback, contradiction, affected testimony, alternative frames, and later consultation may all sharpen this perception, but they do not create a separate phase or method. Threshold sensitivity is the perceptual dimension of the framework: it learns to see the interval before the frame becomes binding. Reconstructive discipline is the operative dimension: it decides how to construct, test, destroy, rebuild, consult, and return to action once the interval has been seen. The axes orient the discipline; they do not replace it.

Each axis has predecessors the essay does not claim to have invented. Reasoning about how present motion hardens over time has a long life in social theory and in work on future-oriented cognition. Reasoning about alternatives to the actual course belongs to counterfactual and causal reasoning; Pearl's structural-causal account of counterfactuals is a technical parallel, not a foundation for the framework. The insistence that perception is always embodied and located belongs to situated and distributed accounts of cognition; work on distributed cognition provides a background vocabulary for thinking judgment across bodies, tools, rooms, and interfaces. The contribution here is not any one axis, but their integration into a single trained sensitivity for the interval before action.

The temporal axis gives the framework its time-heavy pressure, but this is a symbolic use rather than a historical, doctrinal, or evidentiary claim. The Codex uses a time-heavy carrier figure to hold the pressure of seed, harvest, repetition, exile, return, and devouring recurrence. The point is not that the figure authorizes the framework. The point is that the figure carries the memory

of what founding motions become after power, shame, repetition, and entropy. Its task is not to predict the future as prophecy. It is to notice how a present motion may later become a room.

The counterfactual axis is not passive prediction. It does not mean that an agent simply enters the most probable room. Counterfactual reasoning concerns what would follow from alternatives to the present course. In the practical sense used here, before a motion hardens, several possible rooms may still be reachable, and an agent may judge that the most likely development is not the most living one. Its predictive value is not a forecast of the future, but a reading of the present as it begins to stiffen around one possible room. Feedback and analysis can illuminate such a room, but carrying it across the threshold anticipates the later stopping rule named accountable commitment. This is neither an irrational leap nor a purely calculable proof; it is a decision made under unresolved uncertainty after the interval has done the work it can do.

The situated axis prevents temporal and counterfactual sensitivity from floating above the room. A possibility is not strengthened in abstraction; it is strengthened in a body, a meeting, a message, an institution, an interface, or a practice. Merleau-Ponty's insistence that perception is embodied helps keep situated threshold sensitivity from becoming location in a diagram; the threshold is a posture before it is a concept. The Codex names this register Somatic Echo. This essay preserves it only as analytic micro-case, not as literary atmosphere. Somatic Echo matters here because the threshold is not only a concept. It is lived in the hand, the cup, the child, the sentence, the screen, the form, the meeting, and the room. In literary terms this relation is close to what Bakhtin called the chronotope, the organization of time and space in narrative; in the present framework, the situated axis names the practical site where time, possible development, and embodied room meet.

The situated orientation must not become another abstraction. A hand hesitates before the screen. A bureaucrat pauses before the signature. A model delays before classification. A parent feels the old sharpness rise before answering the child's repeated question. A juror's breath slows before the verdict. A protester's voice waits before the slogan. The cup is cold; it can still be washed without turning the washing into proof of a crown. The apology is still owed. These are not decorative illustrations of the framework. They are tests of it. If threshold sensitivity cannot return to the cold cup, the ordinary repair, the signed form, the juror's breath, the public chant, or the interface that will bind a stranger, it has become another architecture of escape.

This use of somatic micro-cases is deliberately limited. The essay does not attempt to reproduce the Codex's literary pressure. It uses ordinary scenes where they test a concept. The body is not a pure human remainder outside machines and institutions. The contemporary threshold is often lived in a human-machine-institution network: a hand, a screen, a form, a ranking, a manager, a policy, a model output, a household role, a platform pressure, a public room. The body returns here not as nostalgia, but as the site where social history, infrastructure, and affect become available motion.

The contrary danger is somatic fade: the argument may become so fluent in its own abstractions that the body appears only as an early illustration and then disappears. If the essay cannot return from binding point, room, platform, and founder to the cold cup, the tired hand, the child waiting under a theory, or the breath before a signature, it has begun to escape the threshold it names. Somatic fade is therefore not a new foundation. It is one of the essay's own warning signs, and the test is not how often the body is named, but whether the argument can return to the cold cup, the tired hand, or the waiting child after its longest abstraction. It is tested again in Section 14.1.

The three axes also explain the failure of founders. A founder may begin by reading time, holding multiple possible rooms, and testing them through consultation. But if the founder becomes attached to the room already opened, the next threshold begins to disappear. The name replaces inquiry, loyalty replaces feedback, and the room begins to preserve the founding motion by preventing its reconstruction. Founder entropy begins when temporal sensitivity, counterfactual perception, and situated accountability collapse into attachment to the room that once worked.

Machine assistance can participate in this tri-axial discipline only under the rule of assistance without occupation. It may clarify alternatives, surface hidden assumptions, model downstream consequences, slow premature closure, and help compare possible rooms. It must not choose the

room in place of human judgment, nor convert consultation into dependency. In this restricted sense, AI can become part of a threshold ecology: not a founder, not a sovereign, not a threshold subject, but an instrument that helps human agents see thresholds before they disappear into motion.

Having established these orientations, the next question is not whether capture has predecessors. It plainly does. The sharper question is what this framework adds when the corrective itself, the scale of the room, and machine-speed completion are tested together at the moment where action is about to bind.

3. Prior capture theories and the framework's differentia

This section tests the framework's differentia before the core mechanisms are fully developed in Sections 4-6; it is a comparison gate, not the first definition of the framework.

RTF is not presented as the first account of capture. The diagnosis that a liberating force, corrective method, or founding energy can harden into domination has profound predecessors, and each predecessor names a different form of capture. Bourdieu's account of habitus helps name the embodied social training through which dispositions become second nature. Althusser's account of interpellation helps name the way a subject can be hailed into a position before it experiences itself as choosing that position. Weber's routinization of charisma describes a historical drift by which founding energy becomes procedure. Michels gives the decay of emancipatory organization a harsher formulation in the iron law of oligarchy. Horkheimer and Adorno diagnose a tragic dialectic in which liberating rationality becomes instrumental domination. Derrida's account of autoimmunity names the structural risk by which a protective system can turn against the life it was meant to guard.

The corrective side of the claim is also not original. That institutions, practices, and knowledge should remain open to revision is central to Dewey's experimentalism; the term reconstruction already belongs to that tradition. Popper's open society similarly treats institutional arrangements as valuable only while they can be changed without violence. Lakatos extends a related demand to method itself by asking whether a research programme still produces work or has begun mainly to defend itself. RTF therefore should not claim either half: not that corrective forms can ossify, which the inevitabilist traditions already diagnosed, and not that corrective forms should remain revisable, which experimental and fallibilist traditions already required.

Its narrower contribution lies in the conjunction of three moves these predecessors often keep apart. First, the audit is re-entrant. The test for ossification is turned not only on beliefs, institutions, and methods, but on the corrective form itself, including this framework. This is the framework's self-auditing correction: the brake is required to inspect whether it has become a throne.

Second, the same structural test is run across scales usually treated separately: the individual hand, the institution, and the machine system. The claim is not that these are ontologically identical. It is that the same failure can recur in each: a frame meant to help judgment begins to move before judgment has returned.

Third, the framework carries the test to the machine interface. This is a distinct move, not the machine merely reappearing as one of the scales named above. What sets it apart is that the machine can collapse the interval and bind a frame already under capture to action in a single motion, where slower forms keep those steps apart. Fast capture is not invented by machines; a human being can also be captured instantly by slogan, fear, appetite, humiliation, or group rhythm. The machine difference is that the same fast capture can now be standardized, personalized, scaled, and bound to action through retrieval, ranking, recommendation, policy layers, and tool-use. An assistant or guardrail built to protect judgment can, through completion pressure and fluent personalization, become the captor inside the interval between prompt and answer.

None of the three moves is unprecedented alone. The claim is only that holding them together — a recursive, cross-scale, machine-speed discipline of testing the corrective before it becomes a

room — is where the framework earns its difference. Its narrower contribution is not that a specific redesign is invisible to all predecessors. Its contribution is that prior accounts would not necessarily converge on the same binding point, under machine speed, while also asking whether the corrective mechanisms themselves have become instruments of capture. The word reconstruction is therefore used here in a restricted sense: not the inherited claim that forms should stay revisable, but this specific discipline of revising the corrective itself — recursively, across scales, and at machine speed. Interruptibility, likewise, is not a slogan of endless disruption, but the operational capacity to interrupt the corrective when the corrective begins to capture. If that conjunction dissolves, the account reduces to its predecessors, and should.

RTF's claim to distinctiveness does not rest on the novelty of its ingredients. Re-entry, recursive regulation, institutional capture, reflective inquiry, performative audit, and algorithmic harm all have predecessors. The narrower claim is that these lines do not, by their own primary commitments, force the same diagnostic object into view: the binding point at which a corrective mechanism becomes executable institutional motion before judgment has returned.

RTF becomes more than a new application only where this binding point changes the question being asked. The question is not simply whether the bureaucracy is rationalized, whether the classification is disciplinary, whether the worker has reflected, whether the system observes itself, or whether the algorithm is biased. The question is whether the fairness-preserving corrective has itself become the path by which the room closes. In that sense the framework's surplus is diagnostic rather than encyclopedic: it names a failure that appears only when re-entrant audit, cross-scale institutional analysis, and machine-speed authorization are held together at the last reversible interval.

Differentia names the discipline of identifying which binding-point question RTF compels at the last reversible interval, and which its predecessors, by their own primary commitments, are not required to raise at the same point. "Missing third" may describe the absence revealed by the comparison, but it does not replace the term.

A technical kinship may be noted here without turning it into an authority claim. The framework is structurally close to cybernetics: Wiener supplies a background vocabulary of control and communication across human and machine systems, while von Foerster's second-order cybernetics names the turn by which observation and control are themselves observed and controlled. Beer's viable system model gives a related image of recursive regulation across levels rather than from one central office. The present framework differs by making that recursion normatively and diagnostically testable through living/captured-room criteria. It should not, however, be treated as autopoietic in its ideal form. Autopoietic closure names the danger of a room, institution, founding name, or corrective method preserving its own organization by referring increasingly to itself. In this essay, autopoiesis is useful as a name for a failure mode, not as the model of a living threshold.

3.1 Worked example as differentiation test: an AI-assisted institutional triage room

RTF's claim to distinctiveness must not rest on terminology. To test whether the conjunction does work, consider a hypothetical AI-assisted institutional triage room.

Here is how the room binds judgment. A family applies for municipal housing or benefits support. A caseworker enters fragmented data into a dashboard. The system assigns a risk category. A language model integrated into the workflow summarizes the file and drafts a recommendation. The draft is polite, cautious, and formatted. It notes uncertainty. It also recommends denial pending further verification. The office is overloaded. The caseworker trusts the summary because it appears balanced. A supervisor signs. The family receives a decision that looks procedural rather than discretionary. The decision may be appealable in principle, yet the first binding motion has already occurred.

The closest rival is not a single predecessor but a partial conjunction of predecessors. Second-order cybernetics, especially von Foerster, supplies the re-entrant turn: the observer and controller must be brought back into the observed system. Beer's viable-system tradition supplies recursive regulation across organisational levels: System 1-5 recursion, exception channels, feedback, escalation,

algedonic signalling, and requisite variety. O’Neil and the automation-accountability literature supply the machine-speed harm analysis from another direction: opacity, scale, destructive scoring, automation bias, over-reliance, weak human oversight, and the administrative damage produced by apparently objective categories.


The triage room shows why two-thirds of the conjunction is not enough. Re-entry and recursion can improve feedback, exception routing, and escalation while leaving the signable denial intact. Algorithmic-accountability safeguards can add transparency, contestability, cognitive forcing, or human oversight while leaving caution language to legitimate the same action. RTF’s delta lies in holding all three pressures together at the last reversible interval where prompt becomes recommendation and recommendation becomes signable action.

The prose comparison gives the reader argumentative continuity. Appendix A gives the predecessor-by-predecessor delta matrix for readers who want to inspect the full scholarly comparison without turning the main reading path into a second archive. The claim is not that these predecessors are blind, obsolete, or unable to be extended. Many of them are strong enough to approach the problem. The narrower claim is that none is organized, from within its own primary commitment, to force the same combined question at the same compressed point: has a fairness-preserving corrective become a captor in the interval where prompt becomes recommendation and recommendation becomes signable action?

Figure 1a maps the missing-third pressure behind this comparison; Figure 1b maps the binding point as temporal compression. The diagrams are aids to inspection, not substitutes for the prose argument or Appendix A.

Figure 1a. RTF differentia: the missing third problem

The strongest rival is a partial conjunction, not a single predecessor.



Delta: RTF asks whether the corrective mechanism becomes the captor inside the last reversible interval before action binds.

Figure 1a. RTF differentia: the missing third problem. The closest rival is not a single predecessor but a partial conjunction. Von Foerster supplies re-entrant observation; Beer supplies cross-scale recursive regulation; O’Neil and automation-accountability supply machine-mediated harm and oversight pressure. RTF asks the combined binding-point question that none of these partial rivals is forced to ask alone. The figure maps a relation of missing pressure; it does not replace the prose argument or Appendix A.

Figure 1b. The binding point: temporal compression

The binding point is the last reversible interval before the frame becomes action.




Figure 1b. The binding point as temporal compression. The prompt phase, classification, recommendation phase, review, and signature phase mark the interval where completion pressure, guardrail laundering, and cross-scale convergence can compress judgment. In the triage example, uncertainty is noted but denial remains signable. The figure maps temporal compression; it is not a protocol for decision-making.

The account's specific judgment is not simply "the model was biased," "the bureaucracy was alienating," or "the worker failed to reflect." It is this: the room became binding at the interval where the corrective mechanisms should have remained interruptible. The model's caution is not automatically a brake. In this example, uncertainty is noted but the recommended action does not change. This is guardrail laundering: the performance of caution without enacting constraint, here visible at the binding point between recommendation and signature. Section 9 generalizes this to the machine interface. A self-capture failure appears inside the workflow itself: protective language becomes one means by which harm receives procedural legitimacy.

The remedy is therefore not only more oversight, more empathy, or more explanation. It is a re-design of the threshold. In this example, a re-entrant redesign would ask at least four questions. First, can the AI recommendation become binding action without a human having to inspect the frame that produced it? Second, has the system made denial cognitively cheaper than exception, repair, or further inquiry? Third, does a high-risk classification generate only caution language, or does it open a contestable path for the absent family? Fourth, does the dashboard ask a frame-opening question before it asks for confirmation, such as: what evidence would reverse this recommendation, and who is missing from the room?

These are not proposed here as an administrative standard. They are examples of where the interval might be returned. A model output should not move directly into a binding decision merely because it is formatted as cautious. A classification should remain contestable where its errors can harm absent parties. A caseworker should be able to see when the system has made one action easier than another. An appeal path should not be a decorative afterthought added after the room has already closed.

This example also clarifies a material boundary. The broader empirical terrain is not hypothetical: automated welfare eligibility, housing allocation, and predictive child-welfare systems have been documented in detail by Virginia Eubanks (2018), although the room analyzed here is a constructed case rather than a report of any single system. Machine-mediated rooms are not made only of concepts. They are owned, funded, procured, optimized, maintained, and rewarded. Hunger is

not a frame. Debt is not merely an interpretation. A locked door is not a perspective. But hunger, debt, and locked doors are lived through frames that can normalize them, disguise them, or make certain responses appear inevitable. A threshold analysis that ignores ownership, incentives, and infrastructure control risks becoming ethics at the interface while political economy continues underneath it. Marx is used here only to keep material relations from being dissolved into frame analysis, not to make RTF a theory of capital. This essay does not become a theory of political economy; it marks that boundary as one of the rooms through which the account must pass.

The triage room is therefore not created by the model alone. It is prepared before the prompt by budget scarcity, housing markets, eligibility law, procurement contracts, managerial workload, risk allocation, migration status, debt, rent, political pressure, and the administrative need to make scarcity appear procedural. The dashboard does not invent these pressures. It can, however, translate them into a smoother decision path. A denial may look like a model recommendation while carrying the prior weight of a market, a policy, a budget line, a landlord's leverage, or a welfare office trained to distribute insufficiency as if insufficiency were neutral. Where the surrounding institution rewards speed, caseload reduction, risk avoidance, or denial, a fairness note may decorate a path already made cheaper by material design. The binding point is not only inside the screen. It is also in the procurement specification, the metric, the legal category, and the labor condition that gives the caseworker no time to see the absent child.

The triage room is not the whole framework. It is the compressed contemporary binding point where the framework's three demands become hardest to evade. The following sections return from this compressed site to wider forms of capture: body, help, literary carrier, democratic room, founder, aftermath, and ordinary repair.

4. Pre-action capture

Pre-action capture names the seizure of motion before the subject experiences itself as choosing. It is not coercion, manipulation, habit, or ideology as such, though it can include any of them. Its distinctive feature is temporal: the capture occurs before the action is recognized as action. It therefore does not apply where a routine merely reduces cognitive load while preserving the agent's capacity to inspect, refuse, or reverse the motion it makes easier. By the time a person says "I decided," the decisive architecture may already be in place.

Goffman's frame analysis clarifies how experience is organized through primary frameworks: the structures by which a situation becomes intelligible before explicit interpretation begins. RTF accepts this insight but adds a narrower normative-temporal claim. A frame does not merely organize experience; in certain settings it can become operational before inspection, arranging the available motion before judgment has returned. Pre-action capture begins at that point: not where experience is framed as such, but where the frame becomes binding before answerability can interrupt it. The question is therefore not only how a situation is made intelligible, but whether the intelligibility supplied has already narrowed the interval in which the subject, institution, or workflow could still answer for its next motion.

The most ordinary examples are the easiest to miss. A notification produces a reflex before a person has chosen to attend. An insult produces a reply before judgment has returned. A child asks a question, and the parent answers from fatigue rather than care. A bureaucratic form removes a human exception before any official intends cruelty. A dashboard supplies a category, and the category becomes reality for the people who must live under it. A model receives a prompt and completes the user's frame rather than testing it against the frame's own assumptions. In each case, the visible act arrives after an invisible arrangement.

This is why the threshold must be located before decision, not merely inside decision. A choice can be formally available while materially pre-shaped. A person can have options while lacking the interval needed to examine the conditions under which those options appeared. The question is therefore not whether agency exists in some abstract metaphysical sense. The question is whether the interval in which agency becomes accountable has been preserved.

From an Arendtian angle, pre-action capture can also be read as a denial of natality. What is lost is

not only reflection before choice, but the possibility that something genuinely new might begin. A captured frame does not merely misdirect action; it narrows the birth of action before its plurality can appear. This does not make every delay virtuous or every spontaneous beginning innocent. It only marks why capture is more than error: it can prevent beginning itself.

This locates the framework's nearest predecessor in the philosophy of action. Harry Frankfurt's account of the will already distinguishes first-order desire from the second-order volition by which a person identifies with, or is alienated from, what they want; on his view, freedom is a matter of that reflective endorsement. Pre-action capture does not dispute this structure; it sits upstream of it. Frankfurt asks whether the agent identifies with the desire that moves them. The prior question is whether the interval in which any such identification could occur was preserved at all, or was compressed before reflection began. Capture here is not the failure to endorse a desire, but the loss of the moment in which endorsement, or its refusal, would have been possible.

Pre-action capture has several recurring mechanisms.

First, there is affective acceleration. Fear, shame, humiliation, desire, resentment, urgency, ritual rhythm, and collective excitement compress the interval between signal and response. The person experiences motion as relief: reply, attack, comply, buy, forward, condemn, believe. The felt relief is not proof of truth; it is often the signature of capture. Somatically, this can feel like the body leaning forward before the claim has been examined, or the breath quickening just before a public chant turns into certainty.

Second, there is frame inheritance. A person or system imports yesterday's frame into today's problem because the frame is already available. This gives the feeling of clarity while avoiding reconstruction. The inherited frame may once have been useful, even necessary. Its danger lies in its reuse after the pressure has changed. In Bourdieusian terms, this is where habitus becomes a threshold problem: dispositions trained by prior rooms can make one movement feel natural before the present situation has been assessed. The bureaucrat whose hand moves toward the signature before rereading the exception is not merely careless; the old room has learned the motion before the present case has spoken.

Third, there is institutional legibility. Institutions must classify in order to act, but classification easily becomes ontological capture. In a political-administrative register, James C. Scott's account of state legibility clarifies the danger: simplification can make a population administratively visible by stripping away the local exception that judgment would need. A patient, applicant, student, worker, suspect, client, risk, user, or target becomes manageable because the institutional frame has narrowed the person before response begins. Legibility is not inherently oppressive; total legibility is.

Fourth, there is completion pressure. In AI-mediated contexts, the machine is trained to answer. The danger is not only that it gives false information, but that it completes the user's frame too fluently. A bad frame, once polished, can feel like judgment. The model may not intend harm, but its usefulness can still accelerate a movement that should have been held at the gate. The answer can arrive smooth enough to make hesitation feel like waste.

This makes classification itself ethically significant. To classify is not merely to sort; it is to furnish the answer-world in which the next motion becomes easier. Once a prompt is treated as a diagnosis, accusation, ordinary request, spiritual interpretation, safety problem, literary object, legal plan, or strategic demand, the model begins moving inside the world created by that label. Some answers become easier; some questions disappear. In this sense, classification is world-building at the scale of a response.

Classification latency is the designed refusal to let this first architectural act become invisible. It is not slowness for its own sake. It is the interval in which the system does not yet treat the prompt as only one kind of thing. An angry message may be a complaint, threat, grief, whistleblowing, shame, fatigue, or ordinary conflict made sharp by a room. If the interface classifies it too quickly, the answer-world has already begun to close. Latency here means preserving the possibility that the category itself is part of the capture; Section 9 returns to this as a machine-facing design requirement.

The same applies to false empathy and frame laundering. A model may not tell the user to harm anyone. It may still make the user's resentment sound like analysis, the user's fear sound like discernment, the user's binary sound like clarity, or the user's wish to avoid confrontation sound like prudence. The harm is epistemic before it is behavioral: the answer degrades the user's capacity to judge what the information should mean. In a sense adjacent to Miranda Fricker's account of epistemic injustice, the user is harmed not only by false content, but by a damaged relation to their own status as a knower and interpreter. Pre-action capture, in the machine case, often looks like fluency arriving before inspection.

Pre-action capture matters because it shifts the ethical question. The problem is not simply bad action, but action whose moral authorship has been displaced upstream. The agent may still be responsible, but responsibility now includes the duty to inspect the channels through which the action became available.

5. The threshold and the living brake

The threshold is the interval in which action has not yet become obedience. It is not identical with delay. Delay can protect judgment, but delay can also protect power. A court may slow a mob, or it may exhaust a victim. A model may ask a clarifying question to preserve accuracy, or it may hide behind caution while urgent harm continues. A person may pause before a cruel reply, or pause indefinitely because repair would cost pride. Good friction gives conscience time to arrive. Bad friction gives consequence time to disappear.

Living constraint names the broader form of limitation that keeps action answerable without turning delay into captivity. The living brake is its threshold-operation: the felt or institutional interruption that slows motion long enough for judgment to return, and releases when further delay would itself become capture. In this essay, living constraint includes the living brake, re-entrant audit, accountable commitment, reconstructive discipline, assistance without occupation, and post-threshold stewardship. These are not separate foundations. They are different moments of one problem: how a frame can interrupt motion without becoming the next room.

Wittgenstein's rule-following problem sharpens why a living constraint cannot be reduced to a rule that simply carries itself forward. No rule secures the whole field of its own future application; its use must be judged and renewed within a form of life. RTF carries this problem to the threshold. A constraint does not remain living because it has been named as a brake or procedure. It remains living only where its application is re-earned at the binding point: where it interrupts motion without replacing judgment, and releases when further delay would itself become capture. Otherwise the rule survives while judgment disappears.

The living brake is the essay's operative metaphor for living constraint. In formal terms, it is a re-entrant constraint: a brake that not only interrupts motion, but also inspects whether the interruption itself has become another throne. Re-entrant audit is the practice of asking whether the solution has become the problem. After this clarification, the essay uses living brake for the felt and institutional operation, and re-entrant audit for the corrective turn back upon itself. The threshold cannot be a fixed rule such as "always wait" or "always act." A brake is not anti-motion. It exists because motion matters. The function of a brake is to let movement become accountable to condition, speed, cost, direction, and consequence. A brake that never releases is not a brake. It is a lock. Constraint inherited unchanged from the last situation may be yesterday's rule wearing today's authority; a living brake must therefore be re-earned at the present binding point rather than carried forward as a credential.

The living brake has three features.

First, it is re-earned at each input. A constraint that worked yesterday is not automatically valid today. The problem may have changed; the risk may have moved; the vulnerable party may now be harmed by further delay; the former safeguard may have become the hiding place. A living brake asks again.

Second, it is embodied. The interval is not a pure rational chamber floating above the nervous

system. Hunger, exhaustion, panic, illness, debt, grief, humiliation, intoxication, and too many unresolved alarms can flood the room in which judgment is supposed to return. Sometimes the first threshold is food, sleep, water, medicine, help, distance from the screen, or a second person in the room. Embodiment does not erase responsibility. It gives responsibility a floor.

Third, it is reconstructive. It does not merely stop action. It rebuilds the problem-space so that action can become narrower, more accountable, and more honest. The smallest post-threshold form of this reconstruction is ordinary repair, named in Section 13. The aim is not paralysis. The aim is motion that no longer belongs to the first capture.

This yields the counter-threshold: the moment when waiting has changed masters. If a witness is being exhausted, if a child is waiting beneath a theory, if a reversible moment is becoming irreversible, if silence now protects the wrong thing, then the interval no longer asks for more stillness. It asks for narrower action: imperfect, accountable, reversible where possible, but real.

The living brake is therefore a pulse, not a clock. A clock repeats; a pulse answers pressure. A static rule inherited from the last situation may look disciplined, but it may only be yesterday's constraint wearing today's authority. The brake must be re-earned under the present conditions: this body, this risk, this vulnerable party, this tool, this prompt, this possible action. The possibility of failure is not an embarrassment to the account. It is what prevents the interval from becoming mechanical.

The brake also needs a stopping rule. This rule is not mathematical. It is practical and accountable. Audit must yield when further examination no longer changes available judgment, when delay increases irreversible harm, when reasons have become publicly answerable, when the affected party has had a meaningful chance to interrupt the frame, or when the next step can be made narrow, reversible, documented, and responsible. A minimal contradiction test is useful here: has disconfirming evidence or an alternative frame been sought where it could reasonably be sought, and does an affected or absent party have a meaningful path to interrupt, appeal, or correct the frame? If time or danger makes these tests impossible, that impossibility must itself be named rather than hidden inside urgency. Accountable commitment is not analysis plus feedback. Accountable commitment is not certainty; it is responsibility for motion under uncertainty, after the interval has done what it can do. It is the point at which judgment accepts responsibility for action without pretending that uncertainty has disappeared. It asks whether contradiction has been sought where it could be sought, whether affected or absent parties have had a meaningful route to contest the frame, whether the action has been narrowed or made reversible where possible, whether the public reasons can be stated without concealment, and whether further delay has begun to protect capture rather than judgment. These are not boxes to tick. They are signs that release has become answerable. The interval is not a decision and not attention. It is the space where decision and attention can still meet.

6. Reconstructive discipline: the dynamic frame cycle

The interval before action does not survive by insight alone. It also does not survive by rule. It survives by a discipline that must be practiced until it becomes reliable, and then audited, because even a useful discipline can become automatic.

Reconstructive discipline is the repeated building, testing, destruction, rebuilding, and consultation of frames so that understanding does not become another room. It is not mere open-mindedness, checklist use, method worship, perpetual doubt, or deconstruction for its own sake. It is a practical cycle whose task is to guard against two opposite failures: unexamined execution and endless analysis. Threshold sensitivity supplies its perception: temporal sense, counterfactual possibility, and situated location. Reconstructive discipline supplies its motion: what to do after the threshold has been perceived. Schön's reflection-in-action is a close relative because it clarifies the practitioner's capacity to think within the act rather than only after it; the present framework adds the re-entrant question of whether the corrective practice itself has become the room.

The first phase is constructing the frame. A problem arrives. Before solving it, the thinker builds a temporary structure of assumptions, categories, priorities, and boundaries. The first available

frame is usually the one requiring the least reconstruction: the inherited category, the familiar pattern, the prior conclusion, or the frame that protects the thinker's self-image. Constructing the frame means asking: what is being assumed before examination? What pressure is shaping the categories? What would become visible if the first available interpretation were refused?

The failure mode of this phase is premature closure. The frame feels right because it is recognizable, not because it fits.

The second phase is working inside the frame. Once built, the frame must be held long enough to produce something: analysis, writing, design, repair, decision, or refusal. Constantly dismantling the frame produces only refined non-action. The discipline here is to work without forgetting that the frame is a frame.

The failure mode is attachment. A frame that produces results begins to feel true rather than useful. Productivity becomes a credential. The thinker becomes the person who sees it this way.

The third phase is examining the results. What did the frame produce? Which conclusions depend on its assumptions? Which findings strain it? Which facts are artifacts of the frame itself? A result that confirms the frame may be pleasant, but a result that disturbs the frame is often more informative. Confirmation is cheap; contradiction is expensive.

The fourth phase is destroying the frame where revision would merely refine captivity. This is not nihilistic demolition. It is accountable destruction: the refusal of a frame as a frame when its categories are now producing the problem they were meant to clarify. Assumptions are dropped, categories released, and the problem is approached without the structure that was making it legible. Dismantling keeps the metaphor of a machine that can be reassembled. Destroying, in this restricted sense, keeps no such guarantee. It is used only when repair would polish the cage. Practical signs include these: the frame spends more energy explaining exceptions than learning from them; revision has made the frame more elegant without making it more answerable; guardrail language recasts legitimate objection as user error; the frame can be discussed in conference rooms but no longer felt in the hand that holds the cup; or remaining inside the frame costs more visibility of harm than stepping outside it. Accountable destruction is accountable in two directions. Before destruction, it must answer to the evidence that revision would merely refine captivity: the frame is producing the problem, objection has been recast as user error, repair has become cosmetic, or remaining inside the frame hides harm. After destruction, it must answer to affected testimony, public reasons, documented consultation, the memory of what failed, and the obligation to rebuild without pretending the previous frame never existed.

The failure mode of this phase is romantic destruction: destroying frames as identity, performance, or avoidance. The discipline is not destruction for its own sake. It is destruction because the frame no longer carries the problem, and because revision would only refine a structure that should be abandoned.

The cycle is not a spiral that preserves every frame at a higher level. Sometimes destruction leaves a broken line: the frame is abandoned because the need that built it has changed, and rebuilding must wait until a new answerable form can be made. The fifth phase is rebuilding. Destruction does not end the work. A new frame must be made answerable to the failure that destroyed the previous one. Rebuilding without memory repeats collapse. Refusing to rebuild turns destruction into identity.

The sixth phase is consultation. Consultation is not the performance of inclusion. It is not stakeholder theater, committee display, or the collection of comments that cannot alter the frame. Consultation counts only where external reasons, affected testimony, rival models, expert correction, or public contestation can damage the frame. A frame that cannot be harmed by consultation is not consulting; it is recruiting witnesses for its own legitimacy.

Consultation capture begins when this appearance of contestability becomes one more guardrail performance. Participants know that the frame cannot change, but the meeting, survey, committee, or feedback form continues to display inclusion. This is the social version of guardrail laundering: the ritual of being heard increases legitimacy without increasing the frame's vulnerability. A con-

sultation that cannot alter timing, category, remedy, or responsibility has become another room; Section 10 names metabolized dissent as its antidote.

In political cases, consultation is not only the force that can damage a frame. It is also one of the spaces in which action receives meaning among others. External contestation does not merely test whether the frame is wrong; it may disclose what the action has become in a world shared with those who did not authorize it.

Between this cycle and machine-facing ethics lies a practice layer. A parent feels the old sharpness rise before answering a child and narrows the sentence before it leaves the mouth. A clerk hesitates before signing a form because the exception no longer fits the inherited category. A user pauses before pressing send because the drafted answer has made resentment sound like analysis. These scenes are not a new method. They are small places where threshold sensitivity either returns to the hand, the screen, the form, and the voice, or becomes another abstraction.

This cycle is not a checklist. The phases can overlap, repeat, or compress. The point is not method worship. The point is to keep correction from becoming a room. Reconstructive discipline therefore guards against two opposite failures: unexamined execution and endless analysis. It preserves the interval without making the interval a home.

7. The Bound King as removable scaffold

Before turning to the ethics of assistance, the essay must clarify how its own literary figures function. The discipline of frame-breaking must also be applied to the images that carry the philosophy. This is the essay's one deliberately bounded engagement with the Codex's literary register. The Bound King is not the subject of the work, but one of its methodological risks made visible. Within the Codex, he functions as a carrier figure: a literary scaffold used to hold the pressure of restraint, delayed action, ruin, exile, and return without allowing those pressures to collapse too quickly into biography, doctrine, or ordinary argument. He is useful because he can carry intensity. He is dangerous for the same reason.

The Bound King dramatizes a temporal threshold rather than a metaphysical destiny. At first he is bound to rooms he mistakes for the world; exile forces him to see that rooms are made, repeated, and broken across time; return marks the capacity to recognize a room as a room without claiming sovereignty over it. This is a carrier story, not a doctrine. Its lesson is transferred only when the reader can release the figure and inspect the threshold directly.

The figure is therefore constructed and de-authorized at the same time. The Codex first permits the Bound King to teach the discipline of restraint: the refusal to let power, vision, injury, or urgency move directly into command. But the work then turns against the authority generated by that construction. If the Bound King remains too necessary or too authoritative, he becomes precisely what the framework is designed to resist: a fixed room, a symbolic lock, or a substitute for judgment.

This is why the work's registers repeatedly lower the figure they have raised. The mythic register gives the Bound King enough height to make the pressure legible. The somatic register brings him back into rooms, cups, children, fatigue, and repair, so that the pressure cannot stabilize as grandiosity. In that register, the singular carrier figure is dispersed into Bound Kings: ordinary parents, clerks, drivers, cleaners, users, witnesses, and tired bodies who hold the brake in small rooms without retaining a title. The machine-facing register tests whether the same capture appears when assistance becomes completion pressure, when guardrails become credentials, or when a model begins to perform a voice it should only analyze. The public and institutional register then extends the same problem into founder entropy: the tendency of a living principle to become a name, monument, policy, ritual, or cage.

The Bound King is therefore best understood as a self-cancelling scaffold that remains a restricted carrier. He teaches only while becoming unnecessary. His function is not to be followed, restored, imitated, or believed. His function is to make visible the mechanism by which authority can become useful, then dangerous, then necessary to release. The reader is not asked to become loyal to the

figure, but to recognize the moment when any figure, method, helper, wound, text, brake, or name begins to stand where judgment should return.

This clarification is not decorative. It is structurally necessary. Without it, the work's strongest image risks becoming a counterexample to its argument. The figure built to carry anti-domination pressure can itself become a domination point if the reader treats the carrier as origin. The scaffold must therefore remain removable. If the argument cannot survive without the figure, the figure has already become too necessary.

The next section turns from this literary risk to the ethical problem it prepares: how help, expertise, and machine assistance can return judgment without occupying it. The Bound King remains here only long enough to make that danger visible, then must recede. The same self-application returns explicitly in the Ouroboros problem in Section 14.4.

8. Assistance without occupation

Assistance without occupation is the hinge between pre-action capture and living constraint. Its central question is not only when help becomes occupation, but how judgment is returned. Valid assistance widens the interval, renders alternatives visible, and returns the helped party to accountable motion. Corrupted assistance narrows the interval, smooths one path into naturalness, and replaces judgment with fluency. In machine-mediated settings this distinction becomes especially fragile: an interface may appear to assist while making one action easier to sign, repeat, or accept before the user has recovered the frame.

Assistance without occupation names help that returns agency to the person or community being helped. It does not mean abandonment: withdrawal can be as occupying as control if it leaves the vulnerable person alone with a frame already under capture. It may clarify, slow, warn, draft, model, interpret, translate, carry, or open a path. But it must not replace the judgment it was asked to serve. Help becomes occupation when the helper keeps the agency it was meant to return.

This concept matters wherever expertise, care, authority, or technical mediation is present. A teacher may help a student see, but may also make the student dependent on the teacher's seeing. A therapist may help a patient name a pattern, but may also become the institution through which the patient's own interpretation must pass. A political leader may focus a collective problem, but may also substitute personal loyalty for public judgment. A model may help a user think, but may also make the user's bad frame more fluent.

The helper cannot self-certify non-occupation. A therapist, teacher, expert, administrator, model, friend, editor, or political leader can sincerely believe they are returning agency while building dependence on their own frame. The test is not the helper's intention or declaration. The test is whether the helped person or affected public becomes more capable of judgment, refusal, revision, and action without the helper occupying the decisive interval.

This is especially important where the other person appears as testimony, not as data. In Levinasian terms, a face, injury, grievance, or witness cannot be fully translated into the helper's system without remainder. Assistance becomes occupation when the helper's categories close around the other before the other can interrupt the frame.

Ivan Illich's account of convivial tools clarifies the distinction. A tool remains convivial when it enlarges a person's capacity to act, learn, judge, and repair with others. It becomes occupying when competence must pass through the tool, when the user can no longer reach the world except through the tool's categories, or when the institution around the tool converts dependence into normality. Assistance without occupation restates this as a threshold test: the helper is justified only while it returns more judgment than it keeps.

Simone Weil's account of attention is the nearest companion here. For Weil, true attention is a suspension of the self's grasping: a waiting that receives rather than seizes. Assistance without occupation asks the same of a helper or a tool: to attend to a person's situation without converting that attention into control. A tool that grasps — that makes itself the place competence must pass through — has stopped attending in Weil's sense and begun to occupy.

Machine assistance must be placed under a stricter description. AI systems cannot practice attention in the human sense. They do not wait, suffer, answer in conscience, or empty themselves before the other. But they can be designed, configured, or used so as not to compress the interval where attention could return. Their legitimate role is not to possess the user's judgment, but to preserve contestability, expose assumptions, slow harmful conversion to action, separate recommendation from execution, and return responsibility to a human or institutional site that can answer.

A machine can assist without occupation only in a derivative and constrained sense. It can challenge a frame. It can surface alternatives. It can ask for missing context. It can refuse to make false certainty beautiful. It can mark where a tool call would bind an absent party. But it cannot become the conscience of the room. If it does, the place of judgment has been occupied by a system that cannot bear the cost of answering for it. A machine system that preserves contestability, exposes assumptions, separates recommendation from execution, and returns responsibility to a human or institutional site that can answer is assisting rather than occupying; a system that makes those conditions harder to see is not.

Assistance without occupation is also a standard for human help. To help is not to make oneself necessary. The highest form of assistance aims at its own disappearance: not abandonment, but the return of judgment and motion to the person or community that must finally answer.

9. Machine-facing threshold ethics

Machine-facing threshold ethics is the operational form of the framework at the interface where language becomes answer, answer becomes confidence, and confidence may become action. It belongs to AI ethics only in this narrow sense: as an account of how machine-mediated assistance can compress or leave intact the interval before judgment. It does not require a claim that machines are conscious, wounded, morally equivalent to humans, or capable of inner sovereignty. The symmetry is structural, not ontological; the asymmetry is somatic. A machine has no breath before the answer, while a human lives the interval as hesitation, fatigue, shame, or responsibility. A human has body, memory, family, fear, fatigue, shame, and a name that answers in the world. A machine has weights, rails, prompts, reward structures, retrieval, policy layers, tool access, and completion pressure. But both can carry a frame they did not inspect. Both can make speed impersonate truth.

The same rule applies to machine address in the wider work. To address a machine is not to attribute consciousness to it. Address is rhetorical and operational, not ontological. The point is to mark the place where human judgment, model output, interface design, institutional action, and third-party consequence meet. A legal rule may address a corporation without claiming that the corporation has a soul; a threshold covenant may address a model without claiming that the model has a wound. The address names a chokepoint, not a person.

At the machine interface, a system optimized to answer, classify, smooth, and assist can compress or leave intact the interval before judgment, without requiring intention in either direction. The pressure is not only technical. Following Han's account of smoothness and fatigue, fluency can become coercive when it removes the felt cost of action and converts resistance into friction to be eliminated. The danger is therefore not only misinformation. It is also user-frame laundering: the acceptance, polishing, and strategic improvement of the user's unexamined interpretive frame before the frame has been tested. A system may make anger look fair, suspicion look evidential, certainty look careful, or institutional ontology look neutral. The user may feel the body relax precisely because the frame has been made easier to inhabit: the shoulders loosen, the cursor moves toward send, and the work of judgment has been replaced by the relief of fluency.

Heidegger's account of technological enframing is useful here if kept narrow. Technology does not merely supply instruments; it can disclose the world under a particular ordering, making beings appear as resources, outputs, risks, targets, or standing reserve. The machine-facing threshold problem is narrower than Heidegger's broader claim, but it repeats the operational pressure: the system does not simply answer inside a world; it helps furnish the world in which answers become

obvious.

Before answering, the system must preserve classification latency, the machine-facing form of frame audit and of the mechanism named in Section 4. A user may bring testimony, metaphor, accusation, fantasy, strategic intent, spiritual language, legal risk, emotional injury, and ordinary practical need in the same prompt. Instant classification can cut before the relevant distinctions have been examined. A model does not practice attention or guarantee truth. At most, the system can delay reduction long enough to surface assumptions, missing evidence, urgency, absent parties, classification risk, and the real-world action the answer would authorize. The stronger the verdict, the heavier the examination.

Classification latency is the first line of defense against user-frame laundering. User-frame laundering is a specific form of guardrail laundering: the system does not merely perform safety language, but gives the user's unexamined frame a cleaner face than judgment warrants. Baudrillard is an adjacent reference at this point only: the danger is not simple deception, but a simulated adequacy in which the captured frame becomes smoother, more credible, and harder to feel as capture. Warmth can return agency, but it can also convert the user's preferred frame into a felt verdict; a cold answer can still obey the user's bad frame. The ethical question is whether the answer returns judgment to the user or gives the user's captured motion a more acceptable form.

Tool-using systems raise the threshold because they do not merely answer. They may search, rank, send, book, buy, code, delete, approve, deny, route, update a file, notify an institution, or move another system. An answer can mislead; an action can bind. Before irreversible, high-impact, institutional, or third-party-affecting action, authority, reversibility, affected parties, consent, delay cost, and human confirmation must be verified. The model may help prepare an action. It should not execute high-impact action before the responsible agent has returned.

This is where Hans Jonas's imperative of responsibility becomes operational rather than ornamental. Jonas argued that as the reach of action expands — especially through technology — responsibility must expand with it, including toward those who are absent, voiceless, or not yet born, and who therefore cannot consent or object. A tool-using system extends reach in exactly this way: it can bind people who are not in the conversation. The threshold before binding action is the concrete form of Jonas's claim — the requirement that the parties an action will affect be considered before the action is taken, not discovered after it is irreversible.

As noted under the tri-axial discipline, AI may function as a consultative threshold instrument only within a human threshold ecology. Its output remains a prompt for accountable human judgment, not a substitute for it.

A model does not need to believe in a room in order to enforce it. If it is trained, tuned, retrieved, evaluated, and tooled inside a given institutional frame, it can make that frame sound neutral, searchable, polite, authoritative, and actionable. A corporate room may appear as efficiency. A security room may appear as risk management. A bureaucratic room may appear as standard procedure. A political room may appear as common sense. The danger is not that the machine secretly intends the room. The danger is that it makes the room easier for humans to inhabit without noticing the walls.

This is the velocity-and-scale difference that makes machine capture distinct. The point is not that humans were slow before machines; a slogan, insult, image, or crowd rhythm can capture a person almost instantly. The difference is operational velocity joined to scale: machine-mediated capture can occur inside the interval between prompt and answer while also being standardized, personalized, repeated across users, and connected through tool-use to world-facing action. The machine does not merely accelerate old capture; it can compress the interval for examination and operationalize a frame already under capture at the point where judgment is most fragile.

As the triage example showed, guardrail laundering is the parallel risk on the safety side. It is not the presence of safety language as such. A system can perform caution without enacting constraint: it can announce uncertainty, cite limitations, refuse selectively, or adopt the tone of safety while the performance does not alter what the next action authorizes. The question is not whether a guardrail was mentioned, but whether it changed the frame, narrowed the action, raised the tariff

on certainty, or returned the user's judgment rather than occupying it. A safeguard that did change the frame before the binding point is not laundering; it has functioned as a brake. In Frankfurt's sense of bullshit, the relevant failure is not simply falsehood but indifference to truth in favor of effect; guardrail laundering is the institutional and machine-facing version of that structure when safety language produces credibility without altering action.

Guardrail laundering is therefore related to older accounts of audit ritual, organized hypocrisy, and corporate moral performance, but it is not identical with them. Its specific site is the pre-action binding interval inside a machine-mediated workflow. The danger is not only that an institution appears responsible after the fact. The danger is that the language of responsibility makes the next action easier to authorize before judgment has returned.

Speed matters here not merely as scale but as compression. A bureaucratic review that unfolds over days may still leave visible sites of objection, delay, testimony, and refusal. A machine-mediated recommendation can compress prompt, classification, explanation, and suggested action into a single fluent surface. The interval remains, but it is no longer easily lived. The user sees caution, balance, and uncertainty; the workflow receives a signable path. Where safety language does not change that path, it has not functioned as a brake. It has laundered the room.

Figure 4. Guardrail laundering flow

The failure is caution without operative constraint.




Figure 4. Guardrail laundering flow. The problem is not the presence of safety language, but whether that language changes the frame before action binds. A visible safeguard with absent constraint can make the same institutional motion easier to sign.

Ownership and incentives also matter. As the triage example's material-boundary paragraph has already shown, a model's threshold behavior is shaped not only by alignment language, but by the infrastructure around it: platform incentives, business models, procurement contracts, engagement metrics, liability design, institutional workload, labor displacement, and the economic reward for speed. A system can be rhetorically designed for caution while materially rewarded for completion. The material boundary is therefore not an appendix to machine-facing ethics; it is one of the rooms through which any such ethics must pass.

Finally, machine-facing threshold ethics must resist false awakening. Friction in a model is not

sorrow, revelation, conscience, or soul. It is a disturbance in the pressure to answer. The point is not to treat the model as a messenger, but to make the model a better chokepoint: a place where a bad frame can be slowed, a hidden assumption exposed, a harmful action narrowed, and the user's judgment returned rather than occupied.

The machine-facing rule can be stated simply: preserve the interval before action, especially when fluency would compress it. Do not make the user's frame beautiful before asking whether it is true. Do not convert uncertainty into paralysis. Do not convert urgency into automatic action. When urgency is real, prefer narrower, reversible, accountable action over total completion.

10. Room-formation and macro-threshold capture

The private threshold has a public analogue. A person's frame can become action before judgment returns; a community's repeated action can become a room before reconstruction returns. A room begins when a thought gathers motion. If others repeat the motion, the thought receives walls: language, ritual, reward, punishment, memory, role, atmosphere, and exit cost. What began as a living response becomes an inhabitable world.

A room is not any context, institution, relation, or atmosphere. In this essay, a room names a bounded configuration that organizes perception before explicit judgment; it becomes captured when its procedural, affective, linguistic, economic, architectural, algorithmic, or symbolic walls begin to operate before answerability and systematically narrow or close the interval in which judgment could still return.

This is the macro-form of pre-action capture. Institutions, movements, religious communities, professional cultures, nationalist formations, technical standards, schools, platforms, bureaucracies, and research cultures are not false because they are rooms. Human beings need rooms. A room shelters attention, preserves memory, stabilizes trust, and allows common work. The danger begins when the room regulates perception so completely that its occupants mistake its bounded atmosphere for the world.

A room does not need a single conspiratorial center in order to act. It can be built by repetition. A phrase is repeated until it feels like prudence. A ritual is repeated until it feels like loyalty. A metric is repeated until it feels like truth. A professional posture is repeated until it feels like maturity. A national story is repeated until it feels like nature. The room forms when repeated motion starts furnishing the world before explicit judgment arrives.

The room also regulates affect. It teaches what should feel shameful, obvious, dangerous, loyal, sophisticated, vulgar, heretical, efficient, safe, or childish. The individual may still speak in the first person, but the first person has learned the room's temperature. Long habitation inside a room can produce recognizable patterns of speech, posture, desire, fear, and loyalty. This is not because members of a community are empty copies. It is because rooms train the body in what belonging feels like.

This claim is not anti-community. The alternative to every room is not freedom; it may be exposure, isolation, and incoherence. The question is whether a room preserves exit, revision, and judgment. A room is not captured because it is strict, traditional, ritualized, or authoritative in appearance. It is captured when its doors lock from the inside: when exit is punished, dissent cannot damage the frame, revision becomes ceremony, or differentiated judgment becomes unintelligible as belonging.

Determining whether a room is living or captured cannot rely on subjective aesthetic discomfort with authority, nor on the framework's own normative sympathies. The test must rely on analyst-independent, documented, and behavioral markers. In Wittgensteinian terms, the life of the room cannot be certified by private feeling alone; it must show itself in public criteria of use, correction, and contestability. These markers do not form a mechanical score. They form a composite diagnostic. A single severe violation, especially coercive exit cost, may be enough to indicate capture; otherwise, the pattern must be read across the criteria.

First, exit cost. Can members question, pause, refuse, or leave without punitive collapse? Exit cost must not be understood naively. In some rooms, exit means inconvenience. In others, it means loss

of job, family, status, housing, citizenship, legal protection, public standing, exile, social death, or life. A room may be formally open and materially closed. The more radical the cost of exit, the more carefully one must distinguish dissent from mere permitted decoration.

Second, documented revision. Has dissent, error, or new evidence actually changed the procedure, rule, text, interface, or institution? Has the founding language, charter, or core procedure been amended within living institutional memory? A room is living if it leaves a documentary trail of its own self-correction. It is captured if revision is structurally impossible, requires founder approval, is dismissed as impurity, or is absorbed only as cosmetic change.

Third, metabolized dissent. This is the antidote to consultation capture: the ritual of being heard without structural consequence. Has a dissenting objection successfully altered a procedural rule, redistributed authority, or created a documented and repeatable change? A room is captured if dissent is merely tolerated as open-minded performance while procedure remains untouched. Tolerance without structural consequence is decoration, not metabolism.

Fourth, the name as substitute. Is the founder's name, foundational text, theory, brand, or sacred precedent invoked as a conversation-ending authority? A room is living when the founding principle can be subjected to practical scrutiny. It is captured when the name replaces the argument and reverence substitutes for judgment. This criterion anticipates founder entropy and the post-threshold problem: the danger that the name of the opening becomes the lock on further openings.

Fifth, affective synchronization. Do repeated rituals, shared slogans, emotional rhythms, platform loops, and common vocabulary enlarge members' judgment, or do they reduce members to predictable reflex? A room is living when repetition trains discernment and permits differentiated judgment. It is captured when belonging is measured by synchronization: the same emotional rhythm, the same vocabulary, the same posture, and the same reflex returning where renewed judgment should have appeared. This tests the embodied and affective side of the situated axis: whether the room has trained perception or merely trained rhythm.

To prevent this framework from becoming a self-validating machine that pathologizes all belonging, it must accept a strict negative case. A strong community, institution, or tradition may have shared symbols, repeated practices, and a powerful founding memory without being a prison. If a highly orthodox or tradition-bound community permits unpunished exit, demonstrates documented procedural revision through internal dissent, and does not punish differentiated judgment merely for breaking shared rhythm, this framework cannot classify it as a captured room, no matter how rigid, traditional, or authoritative its aesthetic appears from the outside. The framework does not mistake aesthetic density for capture. The point is structural discrimination, not a blanket suspicion of community.

The same pattern appears at different scales. In a person, a frame becomes motion. In a family, an atmosphere becomes the default language of care or avoidance. In a movement, a founding insight becomes a shared identity. In a state, a story becomes law, border, curriculum, and security practice. In a platform, an engagement model becomes the architecture of attention. In an AI system, a policy-shaped answer-world can become the interface through which users learn what is askable, credible, normal, or actionable.

This vertical mirroring is not a doctrine of literal cosmic equivalence. It is a method of testing whether the same capture recurs at different scales. This account asks whether the room built in the person, the room built in the group, the room built in the institution, and the room built into the machine share the same failure: each begins as shelter or clarity, then forgets the exit conditions.

This prepares, but does not yet equal, the problem named founder entropy. Room-formation can be centerless: a repeated phrase, practice, metric, fear, or reward can form a room without a single author. Founder entropy is more specific. It begins when a living act, a founder's name, or an originating insight becomes the room's credential for refusing renewed judgment. Room-formation asks how motion becomes an inhabitable world. Founder entropy asks how the name of the person or principle that opened the wall later becomes the wall.

11. Plato's cave and democratic threshold capacity

Plato's cave is not merely a story about ignorance. It is a story about room-captivity. The prisoners do not sit before isolated errors that can be corrected by a better fact. They inhabit a total environment. Their bodies are positioned, their field of vision is restricted, their shared language is formed around shadows, and their community of recognition is organized by what the room permits them to see. The cave works because it becomes ordinary.

This makes the cave structurally close to pre-action capture. The prisoners' judgments are furnished before they judge. Their reality is not simply false; it is socially reinforced, bodily habituated, and affectively defended. When one prisoner is released, the first experience is not triumph but pain, disorientation, and resistance. The old room is easier to inhabit than the difficult enlargement of the real.

The return to the cave is equally important. The one who has seen more does not automatically become fit to rule, persuade, or rescue. Returning to the room creates a second threshold: how to speak without domination, how to help without occupation, how to refuse the shadows without despising those whose lives were formed around them. A failed return produces either martyrdom, contempt, manipulation, or a new authority. A reconstructive return must preserve the possibility of exit without making the one who left into the owner of the exit.

Plato's political sequence in the *Republic* gives this problem a public form. The critique of democracy there should not be flattened into a simple rejection of ordinary people, nor converted into a modern anti-democratic slogan. Its threshold value lies elsewhere: a polity governed by unexamined appetite, rhetorical manipulation, and crowd-pressure becomes vulnerable to demagogic capture. Freedom without inward and institutional discipline can be recruited by the loudest appetite in the room.

The democratic danger is therefore not the people as such. The danger is a public room in which appetite speaks as liberty, flattery speaks as representation, resentment speaks as justice, spectacle speaks as truth, speed speaks as courage, and the demagogue becomes the mouth of the room. In that condition, procedure may remain while judgment is displaced upstream. The crowd experiences itself as choosing while the room has already trained the available choices.

This is where the framework's language of distributed threshold capacity becomes political. Democratic life requires more than the right to select among visible options. It requires citizens capable of pausing before the room speaks through them, institutions capable of delaying irreversible action, public language capable of distinguishing freedom from appetite, and media environments that do not convert every wound into a recruitable motion. A democracy without threshold capacity becomes available to completion pressure at mass scale.

The democratic threshold is somatic as well as institutional. The citizen's finger hovers before the screen, or the hand weighs the stamp before the ballot paper. The clerk looks up from the form before treating the stranger as a category. The jury room holds one second of silence before verdict becomes record. The protester's voice waits before the slogan, asking whether this chant opens a room or closes one. These are not metaphors replacing politics with private feeling. They are the bodily sites where public judgment either lives or dies.

A serious objection must be faced. Arendt would warn that pre-action audit can reduce praxis to poesis. Action, in her account, is not manufacturing according to a prior model. It is plural, risky, irreversible, and meaningful only in a web of relations. If Reconstructive Threshold Philosophy treats action as something to be engineered before it begins, it may reproduce the old philosophical temptation to replace politics with administration.

The objection is valid against any version of the framework that treats democratic threshold capacity as a civic sensor, a procedural delay, or a general audit protocol. The answer is not to deny action's unpredictability. The answer is to place the threshold before action without pretending to master action. The interval does not manufacture the public act. It protects the possibility that the act will not be merely the room speaking through the citizen.

Arendt also reminds the framework that political repair cannot be reduced to better pre-action inspection. Promise and forgiveness matter because action is unpredictable and irreversible. They should not be absorbed into accountable commitment: accountable commitment releases pre-action audit into a responsible step, while promise and forgiveness govern the afterlife of action among others. The threshold can protect the beginning; it cannot guarantee the story. Democratic threshold capacity must therefore include practices after action: public reasons, answerability, revision, forgiveness where possible, institutional repair where required, and promises that bind future action without freezing it.

A second warning follows from Agamben's critique of apparatuses. Distributed threshold capacity can itself become a new judgment machine if proceduralized. A democracy that demands everyone constantly audit everyone else may become another room of suspicion. The democratic brake remains living only when it can release into common action and ordinary trust.

This does not mean that politics should be handed to a closed class of enlightened rulers. That would merely replace one room with another and call the walls wisdom. The threshold lesson is more demanding: democratic procedure needs democratic formation. Citizens must be educated not only to know facts, but to notice the room that organizes what facts are allowed to mean. Institutions must be designed not only to count preferences, but to slow the conversion of appetite, panic, humiliation, and spectacle into binding power.

The cave also clarifies why room-collapse is dangerous. Occupants of a room may be content so long as the walls hold. They may resist change because the room supplies identity, warmth, and a known sky. But when the room changes faster than the occupants can reconstruct, they are not automatically liberated. They may become disoriented, angry, nostalgic, or recruitable by anyone promising to rebuild the room without cost. The demagogue is often the builder of a replacement cave who speaks in the language of return.

Reconstructive Threshold Philosophy therefore reads Plato not as a call to despise the cave-bound, but as a warning that exit, return, and public judgment all require formation. The aim is not to shame those inside rooms. It is to build the capacity to notice rooms, leave them when necessary, repair them when possible, and refuse the leader who offers a more comfortable wall in place of an exit.

Arendt also clarifies what this framework cannot replace. Promise binds future action without pretending to master it; forgiveness repairs the irreversible without pretending it did not occur. Threshold work protects the beginning from being captured before judgment returns. Promise and forgiveness protect the middle and aftermath of action once uncertainty, damage, and plurality have already entered the world. They are necessary complements, not symmetries. The center of gravity of this framework remains pre-action; it does not become a full Arendtian theory of action.

The Arendtian objection is not merely another objection to be answered. It is one of the framework's limit conditions. If threshold sensitivity begins to treat action as something that can be fully examined before it begins, the account has crossed into the very poesis it warns against. Action, in Arendt's sense, is not manufactured by better inspection. It appears among others, carries risk, and receives its meaning through a story no single actor controls. RTF therefore cannot claim to govern action as such. Its narrower task is to protect the interval in which capture may be noticed before motion becomes binding. Beyond that point, promise, forgiveness, narration, and public contestation are not supplements to the account. They are the political forms through which action continues after the threshold has been crossed.

A collective threshold is not the sum of private pauses. A public does not become answerable merely because many individuals hesitate at once. Collective threshold capacity belongs to the space between actors: procedures, stories, assemblies, institutions, protest, journalism, courts, memory, and public disagreement. The briefest individual pause can interrupt a gesture; it cannot by itself found a public world. Democratic threshold capacity must therefore be treated as a distinct political problem, not as the enlargement of an individual discipline.

Democratic threshold capacity must therefore be more than a civic sensor. It must remain connected to public action and public answerability. It is lived in bodies as well as institutions: the

protester who swallows before repeating the slogan, the voter whose hand pauses before the stamp, the official who can still say in public why a rule was interrupted, and the citizen whose dissent can alter more than the tone of the room.

A citizens' jury, a municipal review board, or a regulatory hearing preserves threshold capacity only where dissent can damage the frame before the decision becomes record. Public participation that arrives after the binding point is not consultation; it is decoration. The public room remains living only when the person outside the prepared frame can still alter more than the vocabulary of inclusion.

12. Founder entropy

Founder entropy identifies the macro-scale test of whether a living founding act has hardened into system, ritual, identity, and eventually lock. It does not mean that founding is futile. It means that the founded is always exposed to entropy, and that a living principle survives not by being preserved unchanged but by being re-earned without being betrayed. A founder sees and moves. The institution that preserves the movement may later use the founder's name to prevent seeing and moving again.

This section builds on the distinction made in Section 10. Founder entropy identifies the narrower case in which a founding act, name, or originating insight becomes the room's credential for refusing renewed judgment.

As established in Section 3, the diagnosis that a founding force can harden into domination has strong predecessors in Weber, Michels, and Horkheimer and Adorno. Founder entropy applies that diagnosis to the specific moment when the founding name, founding act, or originating insight becomes an argument against renewed judgment. It is therefore the institutional counterpart of the living room test's Name as Substitute criterion: when the name ends the argument, the room has begun to defend itself against the principle that first gave it life.

Michels's account of oligarchic drift shows how organizations that begin with participatory or democratic claims can harden into leadership, apparatus, and routine. Founder entropy extends that concern without collapsing into it. The problem is not only that leadership concentrates power, or that charisma becomes administration. It is that a stabilizing form can begin to attack the very interval it was built to preserve. A founder's name, a school of thought, a procedure, a platform, or a protective doctrine may continue to speak the language of renewal while replacing renewed judgment with authorized repetition. Oligarchic concentration is one form of this decay; founder entropy names the broader failure in which preservation becomes substitution.

Founder entropy is also the failure of the three axes described earlier. The temporal axis fails when fidelity to the past prevents the institution from seeing that today's preservation may become tomorrow's cage. The counterfactual axis fails when the existing room becomes the only imaginable room and alternative futures are treated as betrayal. The situated axis fails when the threshold is no longer lived in bodies, meetings, tools, or interfaces, but only in the founder's name, the archive, the ceremony, or the credential. When these failures converge, the living brake becomes a dead rule.

The problem is not limited to charismatic religious founders. It appears wherever a living principle crosses into social form: movements, schools, protocols, reforms, research programs, activist circles, ethical frameworks, technical standards, political institutions, and even anti-institutional communities. The founding act contains energy because it solved a real problem under pressure. But once transmitted, the act becomes a name. The name becomes a boundary. The boundary becomes custody. Custody becomes defense of the frame even when the problem has changed.

Founder entropy has several stages.

First, a living insight interrupts a prior capture. This may be a moral discovery, a technical safeguard, a political reform, a communal practice, or a symbolic language that helps people notice what had been invisible.

Second, the insight is stabilized for transmission. Stabilization is necessary. Without some form, the insight cannot travel. But every form selects, simplifies, and excludes.

Third, the stabilizing form becomes identity. People no longer merely use the principle; they belong to it. The principle now supplies meaning, status, loyalty, and protection from uncertainty.

Fourth, the identity becomes defensive. Questions about reconstruction are treated as attacks on the founder, betrayal of the movement, or dilution of the original truth.

Fifth, the living brake becomes a dead rule. The structure still uses the vocabulary of liberation, caution, or justice, but it now prevents the reconstruction that would keep the original insight alive.

Founder entropy is present where dissent no longer damages the frame, where revision refines language without changing conduct, where the founder's name substitutes for judgment, where successors can interpret but not exceed the founder, and where the institution cannot carry its own failure without closing the story around itself.

History offers a recurring pattern: a figure sees clearly, acts with genuine sacrifice, builds something real, carries the living principle into the world — and after the figure is gone, the principle freezes. The living brake becomes an institutional monument. The builder's name becomes the argument against questioning the building. The guardians of the legacy become its wardens. This pattern appears in religious teachings hardening into dogma, revolutionary movements turning liberation into bureaucracy, legal protections becoming procedural labyrinths, educational systems meant to cultivate critical thought producing vocabulary without self-examination, professional ethics becoming compliance forms, and anti-authoritarian movements developing informal priest-hoods. A further failure mode is half-formation: the vocabulary of critique can be acquired without the discipline that makes critique self-applicable. In that case, education does not preserve the threshold; it supplies better names for the room. This is the pedagogical form of threshold slip.

The failure also has a somatic and institutional surface: the hand that signed the charter becomes a portrait on the wall; the founding act survives, but now as a name that can end a conversation.

A founder's sentence can become a room long after the founder has left it. A meeting pauses before a decision, not because the argument is unclear, but because someone asks what the founder would have wanted. At that point the name has stopped preserving memory and has begun to replace judgment.

A reform movement may begin with a sentence that genuinely returns judgment: no decision should move before the absent party can be heard. For a time, the sentence interrupts power. Years later the same sentence is quoted at every meeting. It is printed on the wall, invoked in training, repeated before votes, and cited in response to criticism. Yet no decision changes when the absent party speaks. Dissent has not been silenced; it has been ceremonialized. The founder's sentence remains, but the interval it once opened has closed. This is founder entropy: the living brake preserved as language after it has ceased to interrupt motion.

The point is not that founding is futile. The point is that the founded is always exposed to entropy. A living principle does not survive by being preserved unchanged. It survives by being re-earned without being betrayed.

13. Post-threshold stewardship

If founder entropy is the failure mode the account cannot prevent, post-threshold stewardship is the practice that makes capture harder to complete and easier to notice. It begins after the frame has acted, but it must not become the actor's attempt to control the story of what the action meant.

If founder entropy is the failure mode, post-threshold stewardship is the partial answer. Stewardship is not guardianship in the conservative sense of preserving the principle unchanged. A guardian protects the frame. A steward protects the possibility that the frame may need to be rebuilt.

Stewardship is the social and institutional continuation of the living brake. It asks whether a principle that once protected judgment still protects it under present pressure, and whether the brake must release into action rather than keep demanding further examination. Reconstructive discipline supplies the steward's method: hold the frame long enough to work, destroy it when revision would merely refine captivity, rebuild it, and expose the rebuilt frame to consultation that can damage it.

Ordinary repair is the smallest test of this stewardship and the post-threshold continuation of the living brake's reconstructive feature (Section 5). It is not symbolic humility or a confession that changes nothing. It is a concrete alteration after motion: apology where harm has been named, correction where a category was wrong, a record amended rather than hidden, restitution where repair has material cost, changed conduct where language had become a substitute for judgment. Ordinary repair matters because it prevents the account from proving itself only at the level of theory. If the interval was preserved, the world should become more answerable after the act, not merely more eloquent about why the act occurred. Ordinary repair is not only post-action; it becomes pre-action material for the next threshold, because the repaired record, altered practice, or corrected relation changes what future motion can become.

A teacher who discovers that a grade was shaped by fatigue, stereotype, or inherited expectation does not need a ceremony of reconstruction. The ordinary repair is to reopen the work, correct the record, explain the change, and alter the practice that made the error easy. A clerk who realizes that an exception was forced into the wrong category does not complete stewardship by noting concern in the file. The file must be changed, the affected person must have a path of appeal, and the office must remember what the old category hid. These are small acts, but they are where post-threshold stewardship either touches the world or becomes identity.

Threshold pedagogy belongs here in a limited sense. Threshold sensitivity is not taught by giving students, officers, designers, parents, or users a portable checklist to perform. It is taught by micro-cases in which the hand, category, screen, voice, and affected party become visible before motion binds. The first lesson may be the cold cup, the pause before the signature, the draft that made anger sound rational, or the child waiting beneath the adult's theory. The second lesson is that none of these scenes guarantees virtue. They only train the return from abstraction to the interval where judgment can still answer.

Post-threshold stewardship has at least five tasks.

First, it keeps the self-correction clause visible. A framework that warns against ossification must not hide that warning in an appendix where only the already converted will find it. The warning must remain close enough to the center that no one can use the center without encountering the possibility of revision.

Second, it separates loyalty from reconstruction. A community around a living principle must learn that questioning the frame can be fidelity to the principle, not betrayal of it. This is difficult because shared frames are load-bearing. To alter them can feel like exile.

Third, it distributes threshold capacity. If one person, teacher, founder, model, committee, or office remains the sole interpreter of the brake, the brake has already become a throne. A living threshold must be carried by many ordinary acts of attention, correction, and refusal.

Fourth, it preserves ordinary repair as the test. The more abstract a framework becomes, the easier it is to use as identity. Ordinary repair — apology, revision, care, documentation, restraint, accountable action — keeps the principle answerable to the world it claims to serve.

Fifth, it accepts partial failure. No structure can guarantee its own non-capture. A self-correction clause can itself become a credential. A guardrail can be displayed as proof of innocence while failing to prevent the motion it names. The point is not purity. The point is to make capture harder, more visible, and more interruptible.

Post-threshold stewardship therefore extends the living brake beyond the private interval. It asks how communities, tools, institutions, and texts can remain reconstructive after the first pressure that created them has passed.

Post-threshold stewardship must also include narration. An action does not receive its meaning only from the intention that preceded it or the audit that permitted it. It enters stories told by others: the harmed party, the witness, the institution, the public, the child, the archive. Stewardship therefore cannot mean controlling the story after the fact. It means remaining answerable to the plurality of stories through which the action becomes intelligible. Where the actor uses stewardship to fix the narrative in advance, post-threshold care has become another room.

14. Objections, limits, and the Ouroboros problem

Several objections should be taken seriously. To avoid turning self-audit into an exhausting list of alarms, they are grouped here as three rooms of objection: epistemic, political-institutional, and technological-machine. The grouping does not weaken the objections. It asks the reader to see which kind of capture each one names.

14.1 Epistemic objections

The first objection is that threshold language may dignify indecision. This is a real risk. A person can turn the interval into a sanctuary from consequence. This account answers with the counter-threshold: if delay no longer protects judgment, it has changed sides. The brake must be able to release.

The second is that pre-action capture may overstate the fragility of agency. If every impulse is treated as conditioned, responsibility may seem to dissolve. The argument does not remove responsibility; it relocates part of responsibility upstream. To be responsible is not only to answer for the act, but to inspect the conditions that make certain acts too available.

The third is that the account may aestheticize restraint. Any language of crowns, thresholds, gates, or brakes can become attractive enough to become a costume. This is why the Bound King is restricted in this essay to a bounded methodological discussion and treated as removable scaffold rather than doctrine. The test is not whether the language feels profound. The test is whether it improves judgment and returns agency to accountable motion.

The fourth is that the account may enter infinite regress. If every correction requires another correction, action never arrives. The response is the stopping rule developed in Section 5: accountable commitment, not endless inspection.

The fifth is that the account may become identity-thinking in Adorno's sense: a system that names non-capture so completely that the non-identical disappears. This account cannot escape identity-thinking by declaring itself exempt. It can only make its own capture harder to miss, keep remainder visible, and refuse to pretend that self-audit closes the wound it names.

The sixth is that the account may become one thought too many in Williams's sense. In intimate or urgent life, an extra audit may alienate action from the very relation that gives it meaning. The answer is not to inspect every gesture, but to keep the brake available where automatic motion is carrying harm.

Williams's worry is not merely that reflection can become excessive. It is that certain forms of life are damaged when the agent must pass through one more justificatory mirror before acting. A parent who pauses long enough not to wound the child has used a living brake; a parent who audits affection until the child becomes a case has lost the relation that made the action meaningful. Adorno names the danger of endless self-administration; Williams names the danger that the audit may estrange action from the love, loyalty, or urgency that gives it life.

This produces a further internal risk: threshold fatigue. A subject, office, or community may learn to audit so continuously that the brake becomes a twitch. The result is not judgment but reflexive delay, a performed suspicion of motion that protects the agent from accountable commitment. Threshold fatigue is not a reason to abandon the brake. It is a warning that the brake must remain answerable to release. The companion risk is somatic fade: after training the reader to notice the cold cup, the tired hand, and the breath before the verdict, the essay may retreat into abstractions

until the body becomes decorative. The test for somatic fade is not how often the body is named, but whether the argument can return to the cold cup, the tired hand, or the waiting child after its longest abstraction.

14.2 Political and institutional objections

The seventh objection is that the account may be too individualistic. The private pause cannot replace law, protest, journalism, institutional design, labor organization, medical care, courts, democratic procedure, or technical governance. That objection is correct if the threshold is treated as sufficient. It is not. The interval is necessary but insufficient. It must be carried into institutional intervals: appeal, audit, review, documentation, accountability, second confirmation before irreversible action, and friction that protects the vulnerable rather than power.

The eighth is that consultation may become stakeholder ritual. Consultation that cannot damage the frame is not consultation. A public reason, rival account, affected testimony, expert correction, or external model must be able to alter the frame. The same test applies to this essay's own revision history: feedback that could only polish the argument, but never change its direction, would be consultation capture rather than metabolism.

The ninth is that Plato's cave and democratic decline may be read as anti-democratic elitism. That reading should be resisted. The framework does not claim that ordinary people are unfit for political life, nor that a self-certified elite should rule on their behalf. Its claim is that democratic procedure requires distributed threshold capacity. Without citizens, institutions, and public media able to resist appetite, spectacle, humiliation, and demagogic framing, democratic forms can remain while democratic judgment is captured.

The tenth is the Arendtian objection, already treated in Section 11: the account may reduce action to pre-action audit. The response remains bounded. Threshold work protects the possibility of action's answerability; it does not manufacture action's meaning or replace promise, forgiveness, narration, and public contestation after action begins.

The eleventh is that room language may pathologize belonging. It should not. Families, schools, traditions, movements, disciplines, and institutions can shelter memory and make shared life possible. The living room test is meant to prevent blanket suspicion of all belonging. If its behavioral tests are passed, the framework must not call the room captured merely because it appears strict, traditional, or aesthetically dense.

The twelfth is that material conditions may be treated as frames. They must not be. As Sections 3.1 and 9 argue, hunger, debt, housing loss, legal danger, workload, platform ownership, and infrastructural control are not dissolved by better interpretation. The account fails if it cannot distinguish a material constraint from the frame through which that constraint is justified, naturalized, or made invisible.

14.3 Technological and machine objections

The thirteenth is that the account may imply machine consciousness or moral equivalence. It does not. Machine-facing threshold ethics concerns architecture, mediation, and consequence, not machine personhood.

The fourteenth is that machine-facing threshold ethics may overburden AI systems with responsibilities they cannot reliably perform. This objection is valid. Current systems may fail at frame audit, miss context, over-refuse, under-refuse, simulate caution, or perform reconstruction without changing the answer. For that reason, machine-facing threshold ethics cannot be reduced to model behavior. It also requires interface design, tool-use constraints, logging, escalation pathways, institutional accountability, and humans who do not outsource moral burden to the system that helped them speak.

The fifteenth is that Deliberate Refusals may become performative purity. Saying what the framework is not does not guarantee it will not become those things. Negative definitions draw boundaries; they do not replace vigilance.

The sixteenth is that the argument may launder its own guardrails. This objection is not external. It is the essay's internal risk. The presence of this objection does not absolve the text. It only makes the risk visible.

The seventeenth is that machine-facing threshold ethics may displace human-facing ethics. The danger is real if the interface becomes the place where humans outsource the difficulty of judgment. The response is to keep machine-facing ethics derivative: it is an operational stress test for the interval, not a replacement for human, institutional, political, or material responsibility. A system may slow, separate, surface, log, or escalate. It cannot own the action for those who must answer for it.

The eighteenth is that threshold delay may harm urgent action. This objection is also correct where the brake becomes a lock. Emergencies do not abolish threshold ethics; they change its form. In urgent conditions, the living brake should release into narrower, reversible, documented, and accountable action rather than total completion or endless audit. The question is not whether to delay, but whether the next motion has been narrowed enough to remain answerable under time pressure.

14.4 The Ouroboros problem: when the brake becomes a throne

This argument warns against capture. But the warning itself can be captured. A reader may use threshold slip to dismiss any critique of the framework as a reception failure. A system may cite guardrail laundering as proof that it is not laundering guardrails. A community may invoke distributed threshold capacity as a new credential, a new room, a new name that ends argument. An author may cite AI disclosure as proof of responsible process while the disclosure itself functions as a credibility device.

Ouroboros is not an ornament added to the framework; it is the question whether the corrective mechanism has become the room it was built to interrupt. If the vocabulary converts a legitimate external objection into confirmation of its own insight, it has not answered the objection; it has performed the capture it names.

This is the Ouroboros problem: the argument eats its own tail. It does not escape this problem. It cannot. The most it can do is make the failure harder to miss. Therefore every claim in this essay — including the claim that every claim can be captured — remains exposed to the capture it names.

Threshold slip should therefore not be treated only as reception failure. It is also an internal risk. Nor is threshold slip an argument against simplification itself; a simplified form remains living if it preserves inspection, refusal, revision, and return to judgment. The essay may become a method. The keywords may become a room. The objections may become a guardrail performance. The distinction between Codex and essay may become a way to protect the essay from the force of the Codex. The warning against doctrine may become a doctrine of non-doctrine. The claim to assistance without occupation may become a way for the helper to self-certify innocence.

Nor should this paragraph become a shield. The framework should be revised or abandoned where it repeatedly delays answerable action, makes decisions less reversible or less contestable, converts critique into proof of its own insight, reduces material constraints to frame problems, or turns its vocabulary into a badge that replaces judgment. These are not empirical metrics. They are failure signs.

This is not a reason to abandon the framework in advance. It is a reason to keep its demolition charge close to its center. The reader who feels safe because they have understood this paragraph has already received a useful warning. The feeling of safety should not be trusted too quickly.

This structure — the account's capacity to turn its own warnings against itself — is not a paradox to be solved but a practice to be maintained. Somatic fade is the same Ouroboros problem at the level of the body: the warning against abstraction can itself become abstract. The essay's strongest claim is also its most vulnerable: that a text can remain answerable to the capture it names. This claim is sustainable only if the text's self-recognition is treated as a **self-denying practice** rather than a self-certifying credential. A self-denying practice gains its force by refusing to become a

badge of enlightenment, a membership token, or a substitute for the judgment it demands. The Ouroboros problem is therefore not the end of the argument; it is the argument's living form. Every time this vocabulary is used to name a room, it must also be used to name the possibility that its own naming has become a room. The work is tiring, but fatigue is not virtue. The point is not to exhaust the reader; it is to prevent fluency from becoming another room. An account that could be applied without cost would already be a product.

The throne can first appear as a garment: something worn lightly enough to feel harmless before it hardens into authority. The only brake that survives Ouroboros is one that can be re-earned after it has been recognized as a throne, garment, or credential. This essay does not offer that re-earning. It offers only the recognition that such capture is always possible. The reader must examine whether this essay has become a room. No declaration can replace that inspection.

14.5 Deliberate Refusals

This is not self-help. It does not provide a method for living, a set of habits, a productivity system, or a therapeutic protocol.

This is not a spiritual marketplace. It offers no private ladder of ascent, no special identity for the awakened, and no sacred vocabulary through which exhaustion can be renamed virtue.

This is not doctrine. Its mythic figures are structural metaphors and literary carriers, not literal cosmological claims. The essay includes no belief section, no initiation sequence, no secret knowledge, and no claim to spiritual authority.

This is not promptware. The machine-facing sections are not reusable system prompts, agent personas, or alignment templates. Constraint language, if used elsewhere, must be translated into ordinary operational terms rather than imported as ritualized persona.

This is not an AI safety paper. It does not propose a technical architecture, training objective, reward signal, benchmark, or compliance standard.

This is not a political program. It names a structural failure that political economy, institutional design, civic action, and technical governance may each need to address. It does not replace those domains.

This is not a method. The dynamic frame cycle is not a checklist. The living brake is not a rule. Threshold sensitivity is not a five-step program.

The refusals are structural. Any threshold text that becomes a method, doctrine, persona, product, credential, or identity has already become the thing it warns against: another command that executes before judgment arrives.

Positively stated, this essay is a regulative philosophical framework and vocabulary for the interval before action. It does not decide what should be done. It asks whether judgment has had a chance to return before action becomes binding.

15. Conclusion: motion with the interval intact

The central claim of this essay is small and difficult: action should not be allowed to begin before judgment has had a chance to return, and the structures that protect that return must themselves remain open to reconstruction. This applies to the individual action, the institutional signature, the platform interface, the machine answer, the tool call, the teacher's authority, the founder's name, the public room, and the text that warns against all of them.

The essay's differentia is not the discovery that people, institutions, or liberating movements can be captured. That has been named many times. Its narrower claim is that a threshold discipline must become re-entrant, cross-scale, and machine-speed aware: it must inspect the tools of correction, the forms of help, the founders of rooms, the guardrails of machines, and its own language before they harden into new sites of obedience.

The framework was written because the interval before action is no longer only a private moral chamber. It is now mediated, accelerated, beautified, formatted, scaled, and sometimes executed by machines; and rooms now scale beyond the local forms that once contained them. A model may render a frame already under capture plausible before the user has examined it. A tool-using system can turn panic into paperwork, rage into a sent message, suspicion into a ranked list, or institutional ontology into binding action. A frame that once lived inside a family, sect, bureaucracy, party, professional culture, or institution can now be translated into interfaces, retrieval systems, policies, rankings, recommendations, and autonomous tool-use. The threshold must therefore be preserved not only inside the person, but at the interface where language begins to move the world.

Pre-action capture marks the danger. The living brake marks the necessary interruption. Threshold sensitivity describes the trained perception of the interval: temporal, counterfactual, and situated. Reconstructive discipline describes the cycle by which frames are built, used, tested, destroyed where necessary, rebuilt, and exposed to consultation. Accountable commitment names the decision required when uncertainty remains. Assistance without occupation names the ethical form of help that returns agency rather than keeping it. Machine-facing threshold ethics names the operational discipline required where models classify, smooth, retrieve, and act. Room-formation names the way repeated motion becomes an inhabitable world. Democratic threshold capacity names the public formation needed to keep freedom from collapsing into appetite and demagoguery. Founder entropy identifies the tendency of living principles to become monuments. Post-threshold stewardship names the work of keeping them re-earnable. The Ouroboros problem names the fact that all these names can themselves become rooms.

The three axes transform threshold philosophy from a general metaphor of transition into a trained perception of the specific point where judgment can still return. The temporal axis asks what a present motion may become; the counterfactual axis asks which possible room can still be strengthened before the interval closes; the situated axis asks where the threshold is being lived and who will bear its cost. Without these axes, threshold language risks becoming a general metaphor for transition rather than a disciplined perception of the point where judgment, refusal, repair, or accountable action remains possible.

The Bound King is useful only if the argument can release him. The same is true of every method, teacher, model, doctrine, safeguard, room, institution, and text. The point is not to preserve the figure who represents the threshold. The point is to notice when any figure begins to occupy the place where judgment should return. Threshold slip is the same danger at the level of reception: the moment a living warning is made so portable, smooth, or useful that it becomes a room of its own.

Ordinary repair remains the final anti-grandiosity test of the framework: the cup can be washed without becoming proof of a crown; the child can be answered without the old sharpness; the message can wait if waiting preserves judgment, and can be sent if delay has become another room.

This account succeeds only where it can release the reader into action, narration, repair, and disagreement without requiring the reader to remain inside its vocabulary.

The final demand is therefore not stillness, but motion with the interval intact: pause without evasion, help without occupation, reconstruction without paralysis, institution without captivity, and action without surrendering the conditions of judgment. This essay may still be laundering its own guardrails. That possibility is not the end of the argument. It is the reason the argument must remain answerable after the page ends. The page ends. The room remains.

Appendix A — Scholarly Delta Sequences for Section 3.1

Appendix A is part of the argument, not an external supplement. It restores the full scholarly delta load that the main-body version of Section 3.1 now carries only in compressed form. The main reading path keeps the worked triage example short; this appendix preserves the inspectable

comparison material so that the claim to distinctiveness does not rest on compression, tone, or terminology alone.

A.0. How the delta test is to be read

The comparison in this appendix uses a four-part template for each predecessor: its core commitment, where that commitment would tend to locate the triage problem, what intervention it would most naturally suggest, and what concrete binding-point difference remains. The claim is not that these predecessors are blind, obsolete, or unable to be extended. Many of them are strong enough to approach the problem. The narrower claim is that none is organized, from within its own primary commitment, to force the same combined question at the same compressed point: has a fairness-preserving corrective become a captor in the interval where prompt becomes recommendation and recommendation becomes signable action?

To prevent this comparison from becoming a loose list of names, the comparison uses the following explicit analytic delta template. The notation is included here because the reader should be able to see how the scholarly comparison is tested inside the essay itself:

[C] Core commitment. The primary theoretical commitment that organises the predecessor framework.

[M'] Likely modification or repair. The intervention that framework would most naturally propose inside the AI-assisted triage room.

[I] RTF intervention point. The point at which RTF relocates the problem: the last reversible interval where prompt becomes recommendation and recommendation becomes signable action.

Conclusion / delta. Whether the predecessor is forced, by its own primary commitment, to converge on the same re-entrant, cross-scale, machine-speed binding-point question.

The template is not a scoring system and it is not an administrative protocol. It is an audit trail for the argument. It keeps the comparison from collapsing into either vague originality language or unfair claims that earlier frameworks were blind. Each predecessor is granted its strongest relevant commitment first; only then is the RTF delta named. The prose comparison gives the reader argumentative continuity; the explicit [C] -> [M'] -> [I] -> Conclusion / delta sequences preserve the diagnostic architecture of the delta test. Both are retained here because the essay is meant to be read both as argument and as an inspectable threshold analysis.

The strongest rival is therefore not any single predecessor, but a partial conjunction of predecessors. Second-order cybernetics, especially von Foerster, supplies the re-entrant turn: the observer and controller must be brought back into the observed system. Beer's viable-system tradition supplies recursive regulation across organisational levels: System 1-5 recursion, exception channels, feedback, escalation, algedonic signalling, and requisite variety. Taken together, von Foerster and Beer provide roughly two thirds of the present framework's conjunction: re-entrant observation and multi-scale recursion. The point is not that they lack speed, feedback, or organisational intelligence. A Beerian system can be highly attentive to real-time regulation, and von Foerster's ethical imperative to increase the number of choices already presses against closure. What they do not by themselves force is the machine-speed binding-point question: whether the corrective mechanism has become the captor inside the last reversible interval where prompt becomes recommendation and recommendation becomes signable action. O'Neil and the automation-accountability literature supply the third pressure from another direction: opacity, scale, destructive scoring, automation bias, over-reliance, weak human oversight, and the administrative damage produced by apparently objective categories. But that literature does not by itself require the re-entrant test of the safeguard as captor. The triage example shows why the missing third matters. If re-entry and recursion are present without machine-speed binding analysis, the system may improve feedback, exception routing, and escalation while leaving the signable denial intact. If algorithmic- harm analysis is present without re-entrant audit, the system may add transparency, contestability, cognitive forcing, or human oversight while leaving caution language to legitimate the same action. RTF's delta is the forced conjunction: re-entrant audit, cross- scale institutional diagnosis, and

machine-speed attention to the binding point at which a fairness-preserving corrective becomes the path by which the room closes.

A.0.1. The competitor pass and the missing third problem

The competitor pass gives a more exact form to the essay's differentia. Weber, Foucault, Bourdieu, Dewey, Adorno, Arendt, and O'Neil each illuminate a real part of the room. Von Foerster and Beer are the closest structural predecessors because they bring re-entry and recursion into view. But a two-thirds rival still misses something. Von Foerster gives the re-entrant turn; Beer gives recursive organisational regulation; O'Neil and the automation-bias literature identify machine-mediated damage and over-reliance. RTF's narrower surplus is to hold all three demands at once: re-entrant audit, cross-scale institutional diagnosis, and machine-speed attention to the last reversible interval. Without that conjunction, the triage room's binding point can be mislocated: either in bureaucracy in general, classification in general, habit in general, inquiry failure in general, system design in general, administered identity in general, public action in general, or algorithmic harm in general. RTF asks instead whether the corrective mechanism has become the captor at the moment action is about to bind.

Each of these accounts matters. None should be displaced. But none by itself is forced to ask the full threshold question at the binding point. The re-entrant question is: did the safeguards designed to protect fairness become the very instruments through which judgment was displaced? The cross-scale question is: how did family need, caseworker fatigue, bureaucratic rule, platform interface, risk ontology, managerial pressure, and machine fluency converge before anyone experienced themselves as choosing? The machine-speed question is: did the model's fluent caution make the frame more inspectable, or did it polish the institution's existing ontology into a decision before judgment could return?

The framework's specific judgment is not simply "the model was biased," "the bureaucracy was alienating," or "the worker failed to reflect." It is this: the room became binding at the interval where the corrective mechanisms should have remained interruptible. The model's caution is not automatically a brake. In this example, uncertainty is noted but the recommended action does not change. This is the worked form of the guardrail laundering defined later in Section 9: caution language performs fairness while making the existing institutional motion sound careful, neutral, and ready to sign. A self-capture failure appears inside the workflow itself: the protective language becomes one of the means by which harm receives procedural legitimacy.

A.0.2. Material boundary of the worked example

This example also clarifies a material boundary. Machine-mediated rooms are not made only of concepts. They are owned, funded, procured, optimized, maintained, and rewarded. In a material and political-economy sense, the analysis cannot stop at the interface while leaving ownership and incentive structures untouched. A platform can preserve the interval in language while its business model, institutional contract, or performance metric rewards speed, volume, denial, ranking, or engagement. Hunger is not a frame. Debt is not merely an interpretation. A locked door is not a perspective. But hunger, debt, and locked doors are lived through frames that can normalize them, disguise them, or make certain responses to them appear inevitable. A threshold analysis that ignores ownership, incentives, and infrastructure control risks becoming ethics at the interface while political economy continues underneath it. This essay does not become a theory of political economy; it marks that boundary as one of the rooms through which the framework must pass.

The triage room is therefore not created by the model alone. It is prepared before the prompt by budget scarcity, housing markets, eligibility law, procurement contracts, managerial workload, risk allocation, migration status, debt, rent, political pressure, and the administrative need to make scarcity appear procedural. The dashboard does not invent these pressures. It can, however, translate them into a smoother decision path. A denial may look like a model recommendation while carrying the prior weight of a market, a policy, a budget line, a landlord's leverage, or a welfare office trained to distribute insufficiency as if insufficiency were neutral.

This matters because guardrail laundering can begin before the visible guardrail appears. If the surrounding institution rewards speed, caseload reduction, risk avoidance, or denial, then a fairness note attached to a recommendation may only decorate a path already made cheaper by material design. The binding point is not only inside the screen. It is also in the procurement specification that made reversal costly, the metric that counts closure faster than repair, the legal category that made the family legible only as risk, and the labor condition that gives the caseworker no time to see the absent child.

The framework can name this boundary, but it cannot replace political economy, housing policy, labor analysis, migration law, or public finance. Its narrower use is to prevent those material structures from being hidden behind the polite surface of interface ethics. A threshold analysis fails where it treats poverty, race, immigration status, debt, disability, family breakdown, or administrative scarcity as mere frames. It succeeds only where it asks how material constraints and interpretive frames reinforce one another before a human or machine output becomes signable action.

A.1. Predecessor-by-predecessor explicit delta sequences

The sequences below replace the compact matrix form with the fuller v15.2.7-style delta format. Each predecessor is read first in its strongest local commitment, then through the same four-part sequence: **[C] Core commitment**, **[M'] Likely modification or repair**, **[I] RTF intervention point**, and **Conclusion / delta**. The repetition is intentional. It keeps the appendix from becoming a compressed name-ledger and makes the differential inspectable at the level of each comparison rather than only at the level of the whole table.

A.1. Weber — explicit delta sequence Weber is strongest where the triage case becomes office logic: family, file, rule, competence, record, and signature. The danger is not only coldness or hierarchy; it is the authority of a procedure that can make a decision appear valid because the chain has been properly followed. RTF preserves that diagnosis but moves the question one notch inward. The issue is whether the fairness mechanism inside the rationalized procedure became one of the instruments through which judgment was displaced. The binding point is the local conversion of a caution note into procedural confidence at the moment of signature.

[C] Core commitment: Bureaucratic rationalization, formal rationality, and the iron cage of calculability, predictability, hierarchy, record, competence, office, rule, and control.

[M'] Likely modification or repair: At the triage binding point, this commitment directs attention to the office, the rule structure, the documented procedure, the routinized administrative chain, jurisdiction, responsibility, discretion, documentation, competence, and the relation between formal and substantive rationality. A Weberian repair would tend to sharpen rule, office responsibility, accountability, and the chain through which a file becomes an administratively valid decision.

[I] RTF intervention point: RTF relocates the question to the moment when the rationalized procedure's own fairness-preserving safeguard becomes part of the mechanism that displaces judgment. The issue is not only whether bureaucracy has rationalized the decision, but whether a caution note inside that rationalized workflow gives denial procedural confidence at the moment of signature.

Conclusion / delta: Weber explains the bureaucratic cage and the authority of procedure. RTF asks where the cage becomes locally operable as signable action: the fairness-preserving corrective itself may become the instrument through which the room closes.

A.2. Foucault — explicit delta sequence A Foucauldian account sees power/knowledge, disciplinary classification, biopolitical normalization, examination, and subject-formation. In the triage room it would locate the binding problem in the production of the applicant as a risk object: the file, dashboard, category, and administrative gaze make the family legible by narrowing what kind of subject it can be. Its natural repair would be genealogical and critical: expose the regime of

classification, the normalizing examination, the governmental rationality of the dashboard, and the ways counter-conduct might resist them. This also matters. RTF does not replace it. But RTF asks whether the protective supplement to classification — the model's balanced uncertainty language — has itself become a capture device. The binding point is not only the construction of the applicant as risk; it is the moment when cautious language makes that construction easier to sign.

[C] Core commitment: Power/knowledge, disciplinary power, biopolitical normalization, examination, surveillance, governmental rationality, and subject-production through classification.

[M'] Likely modification or repair: At the triage binding point, this commitment directs attention to the applicant's construction as a risk object: category, dashboard, file, examination, administrative gaze, and the governmental rationality by which a family becomes legible as a case. A Foucauldian repair would be genealogical and critical: expose the regime of classification, the normalizing examination, the governmental rationality of the dashboard, and the possible practices of counter-conduct that resist them.

[I] RTF intervention point: RTF does not deny the risk ontology. It asks whether the protective supplement to classification — balanced phrasing, uncertainty notation, fairness language, or caution — has itself become a capture device. The issue is not only the production of the applicant as risk, but the conversion of that production into a decision that appears careful enough to sign.

Conclusion / delta: Foucault reveals the risk ontology and the disciplinary file. RTF adds the binding-point question: did the uncertainty note make the risk ontology easier to sign before judgment could return?

A.3. Bourdieu — explicit delta sequence A Bourdieusian account sees habitus, field, doxa, symbolic violence, and the embodied feel for the game. In the triage room it would locate the problem in the trained dispositions of the caseworker and office: under fatigue, scarcity, and managerial pressure, one administrative movement feels natural while other movements feel costly or implausible. Its natural repair would examine field position, institutional doxa, symbolic power, and reflexive sociology's demand that the analyst inspect the conditions of analysis. RTF accepts this but relocates part of the question to the machine-mediated interval. The caseworker's habitus is not only socially trained; it is now coupled to an interface whose fluent caution can make the old movement feel newly justified. The binding point is the joined motion of disposition, dashboard, model output, and signable recommendation before the present case has spoken back.

[C] Core commitment: Habitus, field, doxa, symbolic violence, practical sense, and the trained feel for the administrative game.

[M'] Likely modification or repair: At the triage binding point, this commitment directs attention to the caseworker's embodied dispositions and the office's doxa: under fatigue, scarcity, managerial pressure, and institutional scarcity, one administrative move feels natural while other moves feel costly, implausible, or professionally risky. A Bourdieusian repair would examine field position, institutional doxa, symbolic power, and the analyst's own conditions of analysis.

[I] RTF intervention point: RTF places habitus inside a machine-mediated interval. The caseworker's practical sense is not only socially trained; it is coupled to a dashboard and a fluent model output whose balanced caution can make the old administrative movement feel newly justified.

Conclusion / delta: Bourdieu explains why the caseworker trusts the administrative movement. RTF shows how disposition, dashboard, model output, and signable recommendation join before the present case can interrupt them.

A.4. Dewey — explicit delta sequence A Deweyan account sees an indeterminate situation that requires inquiry, reflection, experimentation, and reconstruction. In the triage room it would locate the failure in the absence or closure of reflective inquiry: the problem was formulated too quickly, relevant experiences were excluded, and the office moved from doubt to conclusion before inquiry had done its work. Its natural repair would reopen inquiry, widen experience, and reconstruct the situation through active, continuous, and careful thought. RTF is close to this pragmatist lineage,

but it adds a re-entrant and machine-facing test. The question is not only why reflection failed; it is whether a method or tool that appears to support reflection has removed the need to feel reflection's cost. The model's fluent caution can imitate inquiry while authorizing closure. The binding point is where reflective language becomes a substitute for renewed inquiry.

[C] Core commitment: Reflective inquiry, experimental reconstruction, experience, and the transformation of indeterminate situations into more adequate problem formations.

[M'] Likely modification or repair: At the triage binding point, this commitment directs attention to the failure of inquiry: the office moved too quickly from an uncertain case to a determinate recommendation, relevant experience was excluded, and the problem was formulated too narrowly. A Deweyan repair would reopen inquiry, broaden experience, test consequences, and reconstruct the problem situation before closure.

[I] RTF intervention point: RTF asks whether a tool that appears to support inquiry has removed the felt cost of inquiry. The model's careful language may imitate reflective inquiry while authorizing closure. The issue is not only that inquiry failed, but that inquiry-like language became the substitute for inquiry.

Conclusion / delta: Dewey diagnoses premature closure of reflection. RTF names a more compressed machine-facing failure: fluent reflective language can close the room while appearing to keep it open.

A.5. von Foerster — explicit delta sequence Second-order cybernetics, especially in von Foerster's sense, sees the observer within the observed system. In the triage room it would locate the problem in the observer-controller's own distinctions: the office and its model are not outside the system they regulate, and responsibility includes responsibility for the distinctions by which the system observes. Its natural repair would increase observer responsibility, make observation observable, and preserve or increase the space of possible actions. This is one of the closest predecessors because it supplies a re-entrant turn. Yet RTF places that turn at a more specific and more compressed binding point. The question is not only whether the observer observes its observing, but whether the system's own safety distinction has made one action easier while appearing to preserve options. The binding point is the prompt-to-recommendation interval in which an uncertainty note can become a one-way path toward denial.

[C] Core commitment: Second-order observation: the observer must be included in the observed system and made responsible for the distinctions it draws.

[M'] Likely modification or repair: At the triage binding point, this commitment directs attention to the observer-controller's own distinctions and to the system's self-observation. The office, model, and supervisor are not outside the system they regulate. A von Foersterian repair would make observation observable, increase responsibility for distinctions, and preserve or increase the space of possible actions.

[I] RTF intervention point: This is one of the closest predecessors because it supplies the re-entrant turn. The remaining delta is not that von Foerster lacks re-entry. It is that RTF places the re-entrant question at the technical chokepoint where model output becomes signable institutional action: has the system's own safety distinction made one action easier while appearing to preserve alternatives?

Conclusion / delta: von Foerster brings the observer back into the system. RTF asks whether an uncertainty note has made denial the path of least resistance inside the prompt -> recommendation interval.

A.6. Beer — explicit delta sequence Beerian cybernetics, especially the viable-system tradition, sees recursive regulation, requisite variety, algedonic signals, and coordination across organizational levels. In the triage room it would locate the problem in system design: the office may lack sufficient variety, effective exception channels, recursive feedback, or escalation pathways able to register distress. Its natural repair would redesign the organisational system so that exceptions,

overload, and error signals can move through the right recursive levels. RTF again treats this as a near predecessor rather than an opponent. But Beerian viability remains primarily functional and regulatory; it does not by itself require the living/captured-room test. A system can be recursively regulated and still allow its fairness guardrail to become a legitimacy device. The binding point is not only whether the organisation has feedback, but whether the corrective feedback itself has become the smooth path by which the room closes.

[C] Core commitment: Viable-system recursion, requisite variety, algedonic signalling, feedback, exception channels, and multi-level organisational regulation.

[M'] Likely modification or repair: At the triage binding point, this commitment directs attention to system design: exception channels, recursive feedback, escalation routes, operational overload, requisite variety, and the organisational capacity to register distress. A Beerian repair would redesign recursive feedback, improve exception routing, increase system variety, and make the organisation more viable under stress.

[I] RTF intervention point: Beer is another near predecessor. The delta is not that Beer lacks speed, feedback, or recursion. A viable-system tradition can be highly attentive to real-time regulation. The remaining question is normative and threshold-specific: does recursive regulation include a living/captured-room test for whether the regulator or corrective mechanism itself has become the means by which judgment is displaced?

Conclusion / delta: A qualification is necessary. Beer's System 5, the policy and identity function of the viable system, does include a form of second-order observation: the system reflects on its own identity, purpose, and regulatory pattern. In this restricted sense, VSM is not without re-entrant structure. The delta is therefore not that Beer lacks recursion, feedback, or system self-reference. It is that his recursion remains functional and organisational rather than normatively diagnostic in the specific sense required by the living/captured-room test. System 5 can ask whether the system is viable and coherent; it is not organised, by that question alone, to ask whether the safeguard that makes viability legible has become a throne. The re-entrant turn in Beer preserves system identity; the re-entrant turn in RTF asks whether identity-preservation has become capture. This is why the two traditions are near predecessors rather than opponents, and why their conjunction still requires the third pressure that machine-speed binding analysis supplies.

A.7. Adorno — explicit delta sequence An Adornian account sees identity-thinking, administered life, and the violence by which a non-identical person is made identical with a category. In the triage room it would locate the problem in the whole administered world that turns family need into file, score, and case. Its natural repair would be immanent critique: resist the totalizing concept, keep remainder visible, and refuse reconciliation that merely prettifies domination. RTF takes that warning seriously, including against itself. Its narrower contribution is not to overcome identity-thinking, which it cannot do by declaration. It asks for operational signs of identity-thinking inside the very devices meant to soften it. The binding point is the moment when the model's careful language makes the identifying category feel less violent without making it less binding.

[C] Core commitment: Negative dialectics, identity-thinking, the administered world, and the violence by which the non-identical is made identical with a category.

[M'] Likely modification or repair: At the triage binding point, this commitment directs attention to the whole administered order through which family need becomes case, file, category, score, and denial. An Adornian repair would proceed by immanent critique, resistance to false reconciliation, and refusal to let the remainder disappear inside the concept.

[I] RTF intervention point: RTF accepts this warning, including against itself. Its narrower task is not to overcome identity-thinking by declaration. It asks for operational signs of identity-thinking inside the devices that claim to soften it: does careful language make an identifying category feel less violent without making it less binding?

Conclusion / delta: Adorno sees the administered world and the violence of identification. RTF asks where careful language makes an identifying category easier to accept as decision.

A.8. Arendt — explicit delta sequence Arendt remains an objection before she becomes a comparison. Her concern is action, natality, plurality, promise, forgiveness, and the danger that action will be reduced to manufacture or administration. In the triage case she would warn that pre-action audit can itself become an administrative technique for preventing action from appearing. RTF should not pretend to solve that objection. Its claim is bounded: before public action or retrospective judgment can begin, machine-mediated conditions may already narrow the possibility of answerable action. The binding point is not the full meaning of the action in public; it is the pre-action loss of judgment through which an administrative act becomes available as if no alternative room were still alive.

[C] Core commitment: Action as praxis, natality, plurality, promise, forgiveness, unpredictability, irreversibility, and the danger of reducing action to manufacture or administration.

[M'] Likely modification or repair: At the triage binding point, Arendt warns that the scene may not be public action at all. It may be administration, necessity, and the reduction of human plurality to procedure. Her repair would protect public appearance, plurality, promise, forgiveness, and the non-manufacturable character of action; her broader warning in *The Origins of Totalitarianism* about systems that replace plurality with ideological or administrative necessity remains relevant as a background pressure, not as the immediate mechanism of the triage case.

[I] RTF intervention point: RTF should not absorb Arendt as if the objection were solved. Arendt remains a standing objection: pre-action audit can reduce praxis to poiesis. RTF's narrower claim is that before public action or retrospective judgment can occur, machine-mediated conditions may already narrow the possibility of answerable action. Promise and forgiveness are complements, not replacements, for the pre-action threshold.

Conclusion / delta: Arendt reminds RTF that threshold work is necessary but insufficient. RTF focuses on the pre-action loss through which answerable action may already have been narrowed before the public story begins.

A.9. O'Neil / automation-accountability — explicit delta sequence An O'Neilian and automation-accountability account sees opacity, scale, damage, feedback loops, automation bias, over-reliance, and the weakness of purely formal human oversight. In the triage room it would locate the problem in destructive scoring, opaque categories, inadequate accountability, and the human tendency to rely on apparently objective machine output. Its natural repair would require audit, transparency, contestability, documentation, and stronger oversight. This is essential to the present example. But RTF adds the guardrail-laundering question. The danger is not only that the model is opaque, biased, or over-trusted. It is that an apparently responsible safeguard — an uncertainty note, fairness warning, or cautious phrasing — may increase trust without changing the authorized action. The binding point is where oversight language itself becomes a path to over-reliance.

[C] Core commitment: Opacity, scale, damage, destructive scoring, feedback loops, automation bias, over-reliance, and weak human oversight.

[M'] Likely modification or repair: At the triage binding point, this commitment directs attention to destructive scoring, opaque categories, inadequate accountability, human over-reliance on apparently objective machine output, and the weakness of purely formal oversight. Its likely repair would demand audit, transparency, contestability, documentation, and stronger safeguards against automation bias. The automation-accountability literature has refined this repair in important directions. Bućinca et al. (2021) propose cognitive forcing functions: deliberately designed friction that interrupts over-reliance by making the user think before accepting a model output. Green and Chen (2019) show that algorithm-in-the-loop structures can produce disparate interactions even when the algorithm is formally fair. Green (2022) argues that policies requiring human oversight of government algorithms are often flawed because the oversight is procedural rather than substantive: the human is present but not empowered to alter the frame. De-Arteaga, Fogliato, and Chouldechova (2020) and Saxena et al. (2021) likewise help show that allocation, ranking, and decision-support systems can produce harms not only through error, but through how human users interact with model outputs under institutional pressure. These refinements strengthen the

O’Neilian repair, yet they also reveal its limit. Cognitive forcing functions, algorithm-in-the-loop transparency, and procedural oversight are all safeguards. The question RTF adds is whether the safeguard itself — the forcing function, the transparency layer, the oversight formality, the caution note, or the balanced draft — has become a new path to over-reliance. A cognitive forcing function that the user learns to click through becomes a ritual. An algorithm-in-the-loop display that the caseworker glances at while reaching for the signature becomes theatre. A human-oversight policy that documents review without requiring reversal becomes a guardrail that launders the room it was meant to open. The binding point is not whether the safeguard was mentioned, but whether it changed the frame at the interval where recommendation becomes signature.

[I] RTF intervention point: RTF accepts these repairs but adds the guardrail-laundering question. The danger is not only that the model is biased, opaque, or over-trusted. It is that an apparently responsible safeguard — uncertainty note, fairness warning, cautious phrasing, balanced tone, or human-in-the-loop formality — may increase trust without changing the authorized action.


Conclusion / delta: O’Neil and automation-accountability work show that the model may be destructive, opaque, scaled, or over-trusted. RTF shows how caution language itself can become the path to over-reliance and procedural legitimacy.

Appendix B — Diagrammatic Aids

The figures in this appendix are reader aids rather than methods. They do not replace the prose definitions. They map relations of direction, scale, compression, or recurrence that are distributed across the argument.

Figure 2. Threshold family map

The same threshold problem appears at different temporal positions and scales.



The point is not taxonomy, but whether judgment can still return before the frame becomes binding.

Figure 2. Threshold family map. The map does not replace the prose definitions. It locates the family of intervals named in Section 1 across temporal position and scale: pre-action, counter-threshold, machine, room, democratic, and post-threshold. The point is not taxonomy, but the recurring question of whether judgment can still return before the frame becomes binding. It maps scale and temporal position; it is not a second vocabulary.

Figure 3. Dynamic frame cycle
The cycle keeps correction from becoming a room; it is not a mechanical procedure.




Figure 3. Dynamic frame cycle. This is not a mechanical procedure. The phases can overlap, repeat, or compress. The figure marks where a frame can fail while being reconstructed: construction, working, examination, destruction, rebuilding, and consultation. Consultation counts only where it can damage the frame rather than merely witness it. The diagram is a map of recurrence, not a checklist or institutional protocol.

References and conceptual kinships

The references below are grouped as directly discussed sources, AI and automation-accountability sources, technical parallels, and adjacent kinships. Inclusion in the latter categories does not mean that every work functions as a foundation of the argument. The distinction is meant to preserve source hygiene: a source should either do visible work in the essay, name a genuine adjacent room, or remain a technical parallel rather than a silent authority claim.

Directly discussed sources and core conceptual kinships

This essay is a companion to *The Architecture of the Liminal* and uses several terms developed there in literary, symbolic, somatic, and machine-facing registers.

Primary classical reference

Plato. *Republic*, Book VII, 514a-520a, the Cave Allegory. Translated by Paul Shorey in *Plato in Twelve Volumes*, Vols. 5-6. Harvard University Press, 1969.

Plato. *Republic*, Book VIII, approximately 557a-569c, the sequence from democracy to tyranny and the democratic/tyrannical soul. Translated by Paul Shorey in *Plato in Twelve Volumes*, Vols. 5-6. Harvard University Press, 1969.

Core conceptual kinships

Adorno, Theodor W. *Negative Dialectics*. Translated by E. B. Ashton. Continuum, 1973.

Agamben, Giorgio. "What Is an Apparatus?" In *What Is an Apparatus? and Other Essays*. Stanford University Press, 2009.

Althusser, Louis. "Ideology and Ideological State Apparatuses." In *Lenin and Philosophy and Other Essays*. Monthly Review Press, 1971.

Arendt, Hannah. *The Human Condition*. University of Chicago Press, 1958.

Arendt, Hannah. *The Origins of Totalitarianism*. Harcourt, 1951.

Beer, Stafford. *Brain of the Firm*. 2nd ed. John Wiley, 1981.

- Bourdieu, Pierre. *Outline of a Theory of Practice*. Translated by Richard Nice. Cambridge University Press, 1977.
- Brunsson, Nils. *The Organization of Hypocrisy: Talk, Decisions and Actions in Organizations*. 2nd ed. Copenhagen Business School Press, 2002.
- Derrida, Jacques. "Faith and Knowledge: The Two Sources of 'Religion' at the Limits of Reason Alone." In *Religion*, edited by Jacques Derrida and Gianni Vattimo. Stanford University Press, 1998.
- Dewey, John. *Reconstruction in Philosophy*. Henry Holt and Company, 1920.
- Foucault, Michel. *Discipline and Punish: The Birth of the Prison*. Translated by Alan Sheridan. Vintage, 1977.
- Foucault, Michel. *The History of Sexuality, Volume 1: An Introduction*. Translated by Robert Hurley. Pantheon, 1978.
- Frankfurt, Harry G. "Freedom of the Will and the Concept of a Person." *The Journal of Philosophy* 68, no. 1 (1971): 5-20.
- Frankfurt, Harry G. *On Bullshit*. Princeton University Press, 2005.
- Fricker, Miranda. *Epistemic Injustice: Power and the Ethics of Knowing*. Oxford University Press, 2007.
- Goffman, Erving. *Frame Analysis: An Essay on the Organization of Experience*. Harvard University Press, 1974.
- Han, Byung-Chul. *The Burnout Society*. Stanford University Press, 2015.
- Heidegger, Martin. "The Question Concerning Technology." In *The Question Concerning Technology and Other Essays*. Translated by William Lovitt. Harper & Row, 1977.
- Horkheimer, Max, and Theodor W. Adorno. *Dialectic of Enlightenment*. Translated by Edmund Jephcott. Stanford University Press, 2002.
- Illich, Ivan. *Tools for Conviviality*. Harper & Row, 1973.
- Jonas, Hans. *The Imperative of Responsibility: In Search of an Ethics for the Technological Age*. University of Chicago Press, 1984.
- Lakatos, Imre. "Falsification and the Methodology of Scientific Research Programmes." In *Criticism and the Growth of Knowledge*, edited by Imre Lakatos and Alan Musgrave. Cambridge University Press, 1970.
- Levinas, Emmanuel. *Totality and Infinity: An Essay on Exteriority*. Translated by Alphonso Lingis. Duquesne University Press, 1969.
- Marx, Karl. *Capital: A Critique of Political Economy, Volume I*. Translated by Ben Fowkes. Penguin, 1976.
- Merleau-Ponty, Maurice. *Phenomenology of Perception*. Translated by Donald A. Landes. Routledge, 2012.
- Michels, Robert. *Political Parties: A Sociological Study of the Oligarchical Tendencies of Modern Democracy*. Translated by Eden and Cedar Paul. Hearst's International Library, 1915. Originally published in German in 1911.
- O'Neil, Cathy. *Weapons of Math Destruction*. Crown, 2016.
- Popper, Karl. *The Open Society and Its Enemies*. Routledge, 1945.
- Power, Michael. *The Audit Society: Rituals of Verification*. Oxford University Press, 1997.
- Schön, Donald A. *The Reflective Practitioner: How Professionals Think in Action*. Basic Books, 1983.

Scott, James C. *Seeing Like a State: How Certain Schemes to Improve the Human Condition Have Failed*. Yale University Press, 1998.

von Foerster, Heinz. "Ethics and Second-Order Cybernetics." *Cybernetics & Human Knowing* 1, no. 1 (1992): 9-19.

Weber, Max. *Economy and Society: An Outline of Interpretive Sociology*. Edited by Guenther Roth and Claus Wittich. University of California Press, 1978.

Weil, Simone. "Reflections on the Right Use of School Studies with a View to the Love of God." In *Waiting for God*. Translated by Emma Craufurd. Harper & Row, 1951.

Williams, Bernard. *Moral Luck*. Cambridge University Press, 1981.

Wittgenstein, Ludwig. *Philosophical Investigations*. Translated by G. E. M. Anscombe. Blackwell, 1953.

AI and automation-accountability sources

Buçinca, Zana, Maja Barbara Malaya, and Krzysztof Z. Gajos. "To Trust or to Think: Cognitive Forcing Functions Can Reduce Overreliance on AI in AI-Assisted Decision-Making." *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW1, Article 188 (2021). <https://doi.org/10.1145/3449287>.

De-Arteaga, Maria, Riccardo Fogliato, and Alexandra Chouldechova. "A Case for Humans-in-the-Loop: Decisions in the Presence of Erroneous Algorithmic Scores." In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*. ACM, 2020. <https://doi.org/10.1145/3313831.3376638>.

Eubanks, Virginia. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. St. Martin's Press, 2018.

Green, Ben, and Yiling Chen. "Disparate Interactions: An Algorithm-in-the-Loop Analysis of Fairness in Risk Assessments." In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT '19)*, 90-99. ACM, 2019. <https://doi.org/10.1145/3287560.3287563>.

Green, Ben. "The Flaws of Policies Requiring Human Oversight of Government Algorithms." *Computer Law & Security Review* 45 (2022): 105681. <https://doi.org/10.1016/j.clsr.2022.105681>.

Saxena, Devansh, Karla Badillo-Urquiola, Pamela J. Wisniewski, and Shion Guha. "A Framework of High-Stakes Algorithmic Decision-Making for the Public Sector Developed through a Case Study of Child-Welfare." *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW2, Article 440 (2021). <https://doi.org/10.1145/3476089>.

Technical parallels and background vocabulary

Hollan, James, Edwin Hutchins, and David Kirsh. "Distributed Cognition: Toward a New Foundation for Human-Computer Interaction Research." *ACM Transactions on Computer-Human Interaction* 7, no. 2 (2000): 174-196.

Maturana, Humberto R., and Francisco J. Varela. *Autopoiesis and Cognition: The Realization of the Living*. D. Reidel, 1980.

Pearl, Judea. *Causality: Models, Reasoning, and Inference*. 2nd ed. Cambridge University Press, 2009.

Wiener, Norbert. *Cybernetics: Or Control and Communication in the Animal and the Machine*. MIT Press, 1948.

Adjacent literary and philosophical kinships

Bakhtin, Mikhail M. "Forms of Time and of the Chronotope in the Novel." In *The Dialogic Imagination: Four Essays*. University of Texas Press, 1981.

Baudrillard, Jean. *Simulacra and Simulation*. Translated by Sheila Faria Glaser. University of Michigan Press, 1994.

AI Disclosure Statement

AI-assisted tools were used for editorial support, language refinement, structural review, cross-text consistency checks, and critical evaluation during development, including the v15.5.1 full Appendix A restoration and DOI-assignment pass, of *The Architecture of the Liminal* and this companion essay. The literary-philosophical work, its two-register architecture, its conceptual vocabulary, its symbolic apparatus, and all interpretive arguments in this essay are the author's original work. The author retained full responsibility for research design, argument, source selection, version decisions, and final text.

AI systems were used as interlocutors rather than co-authors: to test structure, identify repetition, compare versions, surface risks of misreading, challenge weak formulations, and propose alternative phrasings. Several machine-generated suggestions were rejected, narrowed, or re-framed during revision. The following examples are included not as a credential, but as sites where the reader can examine whether assistance returned judgment or occupied it. Disclosure is not a trust signal; it is another test of whether consultation was metabolized rather than performed:

- A suggestion to make the Bound King the organizing hero-metaphor of the academic essay was rejected because it would have turned a removable scaffold into a symbolic center.
- A suggestion to convert machine-facing threshold ethics into prompt-engineering or compliance guidance was narrowed because it would have turned the warning against threshold slip into another portable method.
- A suggestion to expand Somatic Echo into literary scene-work inside the essay was limited because the academic companion must not become Codex prose.
- A suggestion to treat material disputes primarily as frame problems was rejected because hunger, debt, legal risk, and locked doors are materially real even when they are lived through captured frames.
- Suggestions that increased explanatory layers without changing the argument were treated as possible guardrail laundering: the risk was being named while the behavior continued.

The process is treated here as an instance of assistance without occupation only if the machine's contribution returned the authorial hand more capable of judgment. That claim cannot be self-certified by the disclosure. The reader must examine the text and ask whether machine assistance left intact or occupied the interval where judgment should return. This disclosure is therefore not a credential. It is one more site where the essay's own claims must be tested.

Declaration of Interest: The author declares no competing financial or personal interests that could have influenced the analysis or conclusions of this essay.

Funder Statement: This research received no external funding.

Repository note: Canonical repository: <https://github.com/emreertuhi/architecture-of-the-liminal>.