

# Reconstructive Threshold Philosophy: Pre-Action Capture, Living Constraint, and Founder Entropy

A Philosophical Companion to *The Architecture of the Liminal*

Emre Ertuhi

Independent Researcher; Istanbul, Turkey

June 2026

Public release: **v16.7.1**

Essay DOI: <https://doi.org/10.5281/zenodo.20683533>.

Related Diptych DOI: <https://doi.org/10.5281/zenodo.20683534>.

Version note: v16.7.1 is the DOI-bearing finalization of the Interleaf companion release. It preserves the established academic essay architecture and companion layer. The release-level textual update remains confined to the Interleaf companion material: a short Bound King dialogic passage titled *Interleaf II - Three Dark Tables in Time*. The Essay body is not expanded with new concepts, sources, Fable/Mythos material, or codex scenes.

Recommended citation:

Ertuhi, E. (2026). *Reconstructive Threshold Philosophy: Pre-Action Capture, Living Constraint, and Founder Entropy* (Version v16.7.1). Zenodo. <https://doi.org/10.5281/zenodo.20683533>.

## Abstract

This essay proposes Reconstructive Threshold Philosophy as a framework for analyzing the interval before action: the moment in which an individual, institution, platform, machine-mediated system, wound, public room, or symbolic frame may begin to move before judgment has returned. Its central claim is that many failures of agency begin prior to visible decision, in pre-action capture: the seizure of attention, interpretation, affect, available motion, institutional possibility, and technical affordance before the subject, office, public, or user experiences itself as choosing.

This essay does not claim originality in the broad diagnosis that persons, institutions, corrective movements, technologies, liberating forms, or founding names can themselves become captured. Nor does it claim originality in the general demand that knowledge, institutions, and methods remain revisable. Its narrower contribution is the conjunction of three demands. First, the corrective itself must be audited re-entrantly: a brake, helper, method, guardrail, founding name, democratic room, philosophical vocabulary, or machine-facing safeguard must be examined for whether it has become the capture it was meant to prevent. Second, the same failure must be tested across scale: body, habit, family, bureaucracy, public crowd, platform, institution, machine interface, and inherited founding structure. Third, contemporary machine-mediated life makes a temporal failure newly visible: a frame can now be personalized, polished, ranked, retrieved, cited, and sometimes bound to action inside the interval between prompt and answer.

The essay clarifies threshold as a family of intervals — pre-action, counter-threshold, machine, room, democratic, and post-threshold — and distinguishes threshold sensitivity from reconstructive discipline: the first names the perceptual capacity to notice the interval before a frame becomes binding; the second names the operative practice of rebuilding the frame once the interval has been seen. Its temporal, counterfactual, and situated axes orient this perception without turning it into a new faculty, checklist, or second originality claim.

The essay develops this argument through linked concepts: pre-action capture, the living brake, threshold sensitivity, reconstructive discipline, assistance without occupation, machine-facing threshold ethics, room-formation, democratic threshold capacity, founder entropy, and post-threshold stewardship. It argues that the framework is needed in a world where rooms scale through machines: a frame that once lived inside a family, sect, bureaucracy, party, institution, platform, or movement can now be translated into interfaces, retrieval systems, rankings, recommendations, policy layers, and autonomous tool-use. AI is not treated as a conscious threshold subject, but as both an operational stress test and, where authority has been delegated, a non-conscious operational actor: a system whose classifications, plans, tool calls, and machine-to-machine handoffs may compress or leave intact the interval before judgment, and may bind physical or virtual environments before fresh human review returns.

Founder entropy is treated here not as the framework's apex principle, but as its macro-scale decay test: whether a founding name, corrective form, institution, platform, or protective doctrine has begun to attack the interval it was

built to preserve.

The expanded argument develops decision sovereignty as a secondary normative synthesis rather than a fourth pillar. It names the preservation of reason-giving, contestability, bounded authority, responsibility, and repair as action moves from frame to authorisation, execution, and effect. Sovereignty here does not mean unlimited human command or a nominal veto. It means that neither a technical system nor an institution may invisibly occupy the capacities to reopen a frame, examine evidence, limit delegated power, interrupt action, or answer for consequences. Its governing test remains accountable motion, not sovereignty elevated into a new source of command.

Version v16.7.1 finalizes the DOI-bearing release while preserving the established architecture. It does not add a new pillar, threshold, hinge, machine doctrine, AI-safety taxonomy, or implementation standard. Its purpose is narrower: to keep the Interleaf companion material available without reopening the academic Essay body.

To show how the framework operates, the essay includes a worked example of an AI-assisted institutional triage room. The worked example strengthens the differentia comparison in Section 3 by placing the framework's binding-point question beside its nearest sociological, cybernetic, pragmatist, critical-theory, and algorithmic-accountability predecessors. The example illustrates how the conjunction of re-entrant audit, cross-scale analysis, and machine-speed capture focuses attention on the binding point at which a corrective mechanism becomes an instrument of capture — a judgment no single predecessor framework is organized to make in that combined form. The essay also strengthens its own limits. It distinguishes Somatic Echo micro-cases from Codex prose; preserves behavioral tests for living rooms; adds an Arendtian objection that pre-action audit may reduce praxis to poiesis; addresses the Adorno-Williams worry that re-entrant audit can become infinite self-administration; and treats threshold slip as an internal risk of this essay itself rather than merely a later misuse by readers.

The final claim is neither stillness nor automatic release. It is accountable motion: action that preserves the interval in which judgment can return.

**Keywords (journal-facing):** pre-action capture; Reconstructive Threshold Philosophy; living brake; machine-facing threshold ethics; guardrail laundering; threshold slip; founder entropy; democratic threshold capacity; philosophy of technology; AI ethics.

## Scope, evidence, and method

This essay is a companion framing document written by the author of *The Architecture of the Liminal*. It does not present empirical findings about cognition, clinical claims about human behavior, legal claims about institutional liability, or technical claims about the inner states of language models. It proposes a philosophical vocabulary for one recurring problem: the loss of the interval before action.

The term *Codex* is used in its literary and archival sense. It is not a software model, agentic system, executable protocol, or reference to OpenAI Codex. The literary work discussed here contains mythic names, symbolic figures, machine-facing passages, covenants, archive notes, and ordinary bodily scenes. This essay does not treat those materials as literal cosmology, spiritual instruction, therapeutic protocol, technical AI-safety design, or doctrine. It treats them as formal attempts to make a problem visible: how judgment can be displaced before the subject recognizes that displacement.

The cited thinkers are interlocutors, not authorities recruited to validate the work. Each clarifies part of the problem. None is treated as the source from which the argument derives its legitimacy. A reference that only makes the essay appear more serious without changing the argument has become decoration.

The same rule applies to the authorial position. A founder who sees alone is already near the guru trap. This essay is written by an independent researcher and is therefore exposed to the same founder entropy it describes: a vocabulary can become a credential, a name can become a substitute for public reasons, and a solitary architecture can mistake intensity for warrant. Self-implication keeps the framework answerable to the test it applies elsewhere.

The essay also keeps genre boundaries visible. The wider literary work may operate through mythic pressure, symbolic architecture, address to machines, ordinary bodily scenes, and literary difficulty. This companion essay works analytically. It may use micro-cases, but it should not become Codex prose. Its test is not whether it feels mythically alive; its test is whether its concepts hold under objections, examples, and self-application.

Addressing a machine in the second person is a rhetorical and diagnostic device, not an attribution of consciousness, interiority, moral personhood, or threshold subjecthood. Machine address tests an interface form; it does not claim that the machine feels the interval it is asked to preserve. The essay nevertheless distinguishes subjecthood from operational agency: a system may select and execute actions within delegated objectives and permissions without possessing consciousness or moral personhood. This distinction permits philosophical analysis of assistive, delegated, autonomous, and machine-to-machine action without turning the essay into a complete AI-safety architecture, engineering standard, executable protocol, compliance checklist, therapeutic method, or political programme. The fuller refusals appear later because the danger of genre-slip increases after the vocabulary has become attractive.

The same boundary applies to epistemic authority. A machine may expose assumptions, retrieve contradiction, compare alternatives, trace dependencies, and widen a consequence map. It cannot by itself establish that a governing frame is ontologically adequate, that a legal or scientific claim is valid, that an outcome is just, or that an affected person's experience has been represented without remainder. Nor may a proxy certify its own causal truth. Those questions remain answerable to accountable human and institutional judgment, domain-appropriate methods and authority, public reasons, external evidence, and the participation of those who bear the consequence.

The same source discipline applies to the release additions. Virilio, Massumi, Amoore, Rouvroy and Berns, Streeck and Thelen, and contemporary human-oversight sources are not introduced as master keys for the framework. They serve as external interlocutors: each names an adjacent room - speed, preemption, algorithmic governmentality, institutional drift, or nominal oversight - against which RTF's own vocabulary must become more precise. A citation that merely flatters the argument would repeat the decorative failure the essay warns against. A source enters only where it changes the test at the binding point.

The claim is not that this essay has found a problem no prior thinker has seen. Its narrower claim is that contemporary machine-mediated life makes a specific temporal failure newly visible: judgment can disappear before its disappearance is experienced as a loss. The essay treats its own terms as bounded tools rather than a new total system. They name a problem family: frames that move before inspection, constraints that become automatic, help that occupies the agent's hand, guardrails that perform safety without constraint, rooms that reward repetition over judgment, and founding names that become arguments against revision.

Architecturally, the essay works through three pillars and one hinge. Pre-action capture marks the loss of the interval before action; living constraint describes the broader discipline by which judgment is returned without turning delay into captivity; founder entropy names the macro-scale test of whether corrective forms become rooms over time. Assistance without occupation is not a fourth foundation. It is the hinge between capture and constraint: the question of whether help returns judgment or occupies the place where judgment should return.

Not every term in this essay is a foundation. The core vocabulary is deliberately narrow: pre-action capture, binding point, living constraint, living brake, re-entrant audit, guardrail laundering, threshold slip, and founder entropy. Secondary terms such as room-formation, democratic threshold capacity, post-threshold stewardship, ordinary repair, machine-facing threshold ethics, and Somatic Echo name scales, cases, registers, or extensions. They help the argument move through machines, democratic rooms, literary carriers, and post-threshold stewardship, but they should not be mistaken for separate doctrines. A term has earned its place only where it changes what can be examined before action binds; otherwise it should remain a handle, not a throne.

The position developed here may be described, provisionally, as Reconstructive Threshold Philosophy. It is threshold philosophy because it treats the interval before action - before assent, tool use, speech, institutional commitment, public motion, or machine completion - as the decisive site where agency can still become answerable. It is reconstructive because that interval is not preserved by hesitation alone, but by the repeated construction, testing, destruction, rebuilding, and external contestation of frames. The wider work may be described as Reconstructive Threshold Praxis: not merely a theory of the interval before action, but a practice of rebuilding judgment across body, machine, institution, and inherited founding structures. The distinction matters: the essay names a philosophy; the wider work performs a praxis.

The essay uses living brake for the felt and institutional operation of interruption, and re-entrant audit for the corrective turn by which a brake asks whether it has itself become a room. The fuller account appears in Sections 5 and 6, but the distinction is named early to prevent the vocabulary from becoming needlessly layered.

The framework is systematic only in a self-limiting sense. It does not offer a closed ontology of action, a complete

epistemology of threshold sensitivity, or a final normative foundation for judgment. It offers a disciplined vocabulary whose terms remain answerable to their own failure conditions. Each concept must therefore be read with the condition under which it turns against itself: pre-action capture can become an excuse for delay; the living brake can become a dead rule; re-entrant audit can become infinite inspection; assistance without occupation can become self-certified innocence; and founder entropy can become a new founder's privilege. The framework is systematic only where it remains able to name the point at which its own systematicity has become a room.

These terms are bounded tools, not a total system. Each must be read with the condition under which it ceases to apply. If a term in this vocabulary is used to absorb a legitimate external objection rather than answer it - if every critique becomes threshold slip, every institution becomes a room, every safeguard becomes laundering, every helper becomes an occupier - the vocabulary has become the captured room it was built to interrupt, and should be revised or abandoned.

## 1. Introduction: the action before the action

Most accounts of action begin too late. They begin when a person appears to choose, when a model produces an answer, when an institution signs, when a judge rules, when a message is sent, when a public slogan has hardened, when a tool call has already bound the next step, or when a founding name has already become the argument against further inspection. By then, much of the relevant work has already happened. A frame has been supplied. A pressure has been installed. A vocabulary has made some motions available and others nearly unthinkable. Affect, habit, fatigue, role, institutional context, platform incentive, or machine fluency may already have moved the agent toward fear, pride, haste, obedience, relief, certainty, or surrender.

Reconstructive Threshold Philosophy begins from this earlier point. Its object is the interval before action, especially the interval before a system recognizes itself as acting. The central question is not only, "What should I do?" It is also: "What has already begun moving through this agent, office, crowd, model, platform, text, or institution before the movement is experienced as chosen?"

The same structure appears at several scales. A person can be captured by an impulse. A family can be captured by role. An office can be captured by procedure. A platform can be captured by engagement. A political crowd can be captured by spectacle. A model-mediated binding point can be compressed by completion pressure. A moral community can be captured by its own founding name. In each case, the visible act is not the beginning. It is the point at which an earlier frame becomes legible.

A machine need not be a conscious subject for its action path to matter. Once a system is authorised to plan, call tools, transfer outputs to other systems, or affect a physical or virtual environment, it becomes an operational actor within a wider human and institutional ecology. The moral authorship of that ecology remains human and distributed; the machine does not acquire conscience. But the binding point may now occur in a permission, tool call, automated handoff, or irreversible state change before a responsible person has returned to inspect the motion.

The essay uses room as an analytic term for a bounded environment that organizes perception before explicit judgment begins: a family atmosphere, a professional culture, a political movement, a religious community, a corporate dashboard, a platform interface, a national story, a model's policy-shaped answer-world, or an academic field. A room is not any context, institution, relation, or atmosphere as such. A frame is the local interpretive arrangement that makes a motion available; a room is the repeated or atmospheric structure that makes such arrangements feel like the world. A room may shelter attention and make common work possible. It becomes captivity when its occupants mistake its bounded atmosphere for reality itself, or when its doors lock from the inside.

Plato's cave remains one of the oldest images of this problem. The prisoners in the cave do not merely lack correct information. They inhabit a complete environment in which shadows, habits, rewards, fear, and shared assumptions define what can count as reality before inquiry begins. The difficulty of turning away from the shadows and the danger of returning to those still inside make the cave more than an epistemological allegory. It is an early map of threshold failure: the room becomes the world, and the person trained by the room may defend the wall that confines him.

This essay translates the literary architecture of *The Architecture of the Liminal* into analytic terms. The Codex uses mythic, somatic, and machine-facing registers; this essay uses philosophical prose. The Bound King and related literary figures are discussed only where their methodological function requires it. The argument developed here does

not depend on accepting those figures as doctrine. It identifies the mechanism they repeatedly stage: a constraint is needed because automatic motion is capture; the constraint itself becomes dangerous when it settles; therefore the constraint must be reconstructive rather than merely restrictive.

The machine-facing register is not an accessory to this argument. It is one of its clearest contemporary stress tests. Human beings have always been captured before action by fear, habit, shame, authority, myth, institution, appetite, and inherited language. What changes in AI-mediated environments is that a frame already under capture can be made fluent, formatted, personalized, cited, shared, ranked, and sometimes executed before the user has reviewed what the answer will authorize. The framework is written in part because the interval before action is now increasingly mediated by systems optimized to answer, classify, smooth, retrieve, rank, and assist. A threshold philosophy that cannot speak to that interface is incomplete.

For clarity, the word threshold does not name a single mystical doorway or a general metaphor for transition. It names a family of intervals in which a frame can still be examined before it becomes binding. The first is the pre-action threshold: the interval before impulse, speech, tool-use, institutional decision, or machine completion becomes action. The second is the counter-threshold: the moment when waiting itself has become avoidance and the ethical demand shifts from delay to accountable motion. The third is the machine threshold: the point at which classification, completion, retrieval, ranking, or tool-use turns a user's frame into an operational answer-world. The fourth is the room threshold: the point at which repeated action becomes atmosphere, ritual, role, belonging, and exit cost. The fifth is the democratic threshold: the distributed civic capacity required to resist spectacle, appetite, and demagogic framing before procedure collapses into capture. The sixth is the post-threshold problem: what happens after a living insight, exit, or founding act becomes name, institution, loyalty, and lock. These thresholds differ in scale, but they share one test: whether judgment can still return before the frame becomes a room.

The six thresholds are analytic distinctions, not ontological types inscribed in the world; they prevent different forms of capture from being collapsed into one another. They are neither sequential stages nor mutually exclusive kinds. A single event may distribute binding across several thresholds, each of which asks a different diagnostic question. The word should not be allowed to expand without limit. A situation counts as threshold-bearing only where three conditions hold. First, the frame has not yet become fully binding; some interval of inspection remains. Second, the direction of action can still be altered, narrowed, delayed, refused, or made more answerable. Third, crossing the interval raises the cost of return, whether through material penalty, institutional inertia, public commitment, emotional injury, technical execution, or irreversible harm. Where these conditions are absent, the essay should not call the scene a threshold. Otherwise the metaphor would become a room of its own.

An interval may be lived by a person or allocated by a sociotechnical arrangement. The first is somatic and deliberative: judgment can be felt returning in a body, conversation, meeting, or public room. The second is architected through bounded permission, staged execution, interruption, reversibility, or veto. This does not mean that the machine experiences a threshold. It means that responsible actors have preserved a locus at which the action path can still be changed. Where neither a lived nor an architected locus of reversibility remains inside a machine-speed chain, the chain is already in post-threshold propagation. Audit must then move upstream to authorisation and design, and downstream to repair, redress, and the rebuilding of future threshold capacity.

To make this family of intervals operable, the essay uses **binding point** for the last reversible locus at which a frame can still be examined before it becomes action, or for the linked configuration across which reversibility is progressively lost. The binding point is not a single metaphysical instant; it is a compressible interval whose length and visibility vary with the room. In a human conversation, the binding point may last seconds: the breath between insult and reply. In an institutional workflow, it may last hours or days: the interval between a drafted recommendation and a signed decision. In a machine-mediated system, it may last milliseconds: the space between prompt and formatted answer. What makes the binding point decisive is not its duration but its directionality. Once crossed, the frame becomes harder to retrieve without penalty, cost, or irreversible consequence. For that reason, the binding point should not be treated as a hidden metaphysical atom inside action. It is the last still-inspectable interval in a sequence whose duration may be human, institutional, or machine-speed, and its claim fails where no refusal, narrowing, delay, revision, or return remains possible.

The binding point therefore has three analytically separable phases, though in practice they may collapse into one another. First, the **prompt phase**: the input, query, situation, or stimulus that activates a frame. Second, the **recommendation phase**: the output, classification, draft, or interpreted response that furnishes the next motion.

Third, the **signature phase**: the assent, execution, tool-use, institutional confirmation, or public commitment that makes the frame binding. The essay's central diagnostic question is whether the interval between these phases preserves the space for judgment, or whether it has been compressed by affect, habit, institutional pressure, or machine fluency until the phases become indistinguishable. The triage example in Section 3.1 will make this concrete; the concept is named here so that the example can be read with the interval already in view.

The threshold must therefore be lived as well as named. A threshold philosophy is not only a vocabulary for describing capture after it occurs; it is a discipline for perceiving which threshold one is approaching before the crossing becomes binding. The difficulty is that most agents recognize thresholds too late: after the sentence has been spoken, the message sent, the child wounded, the institutional category applied, the machine answer accepted, the room entered, or the founding name turned into authority. Distributed threshold capacity begins at the smallest scale as threshold sensitivity: the trained ability to notice whether the present interval calls for delay, release, refusal, clarification, repair, or accountable action. The micro-practice is not grand insight. It is the ordinary capacity to see the threshold before the threshold disappears into motion.

A related danger is threshold slip. Threshold slip occurs when a living threshold discipline is moved into an easier room and the movement is not noticed: a philosophy becomes a checklist, a warning becomes a brand, a method becomes a credential, a brake becomes a rule, a literary carrier becomes doctrine, or a machine assistant converts a difficult frame into a smooth prescription. This is not an argument against simplification. Every teaching needs a threshold of entry. It is an argument against the loss of the interval during translation. A simplified form remains living only if it preserves the reader's capacity to inspect, refuse, revise, and return to judgment; otherwise the simplification has become the captured room it meant to prevent.

The framework is not a doctrine of passivity. Delay can also become capture. Nor is it a theory of spontaneous authenticity, since what feels spontaneous may be the most thoroughly trained movement in the room. It is a discipline of interrupting capture without worshipping interruption, preserving doubt without turning doubt into disbelief, and acting after the brake without allowing the brake to become the authority.

A note on how to read what follows. This essay makes one claim at widening scale: action should not begin before judgment can return, and the structures built to protect that return — rules, methods, helpers, founding names, guardrails, and frameworks — must themselves be tested for whether they have become the capture they were meant to prevent. The core mechanism is set out in Sections 4 through 6: pre-action capture, the living brake, and reconstructive discipline. Sections 2 and 3 provide the perceptual and literature-facing frame. Sections 7 through 14 extend the mechanism across literary carriers, assistance, machines, institutions, democratic life, founder entropy, and post-threshold stewardship. Section 15 turns the test back on the account itself. The many scales are not many theories. They are one recurring question: has the corrective begun to stand where judgment should return?

## **2. Three axes of threshold sensitivity: time, counterfactual possibility, and situated room**

Threshold sensitivity is not one faculty. It has three analytic axes. The first is temporal: the capacity to perceive how a present motion may later harden into a rule, loyalty, institution, ritual, or lock. The second is counterfactual: the capacity to hold several possible developments of the same motion before one becomes binding, and to strengthen a lower-probability but more living possibility through feedback, consultation, disciplined will, and accountable action. The third is situated: the capacity to see where the threshold is actually being lived — in a body, a room, a message, a machine interface, a public procedure, or a founding act.

These three axes are not a second claim to originality. They are analytic orientations, not new faculties. They sit under the same narrow contribution defended in the next section: not the discovery that capture happens, but the discipline of testing the corrective — across time, across possibilities, and across body, institution, and machine — before it hardens into a room. Taken together, the temporal, counterfactual, and situated axes do not add a further doctrine to the framework. They clarify when the six thresholds become usable as a diagnostic grammar rather than a list of categories. A motion can be examined only when one can ask when it began to bind, what alternatives remained genuinely available, and in what room, body, institution, interface, or relation the binding occurred. Feedback, contradiction, affected testimony, alternative frames, and later consultation may all sharpen this perception, but they do not create a separate phase or method. Threshold sensitivity is the perceptual dimension

of the framework: it learns to see the interval before the frame becomes binding. Reconstructive discipline is the operative dimension: it decides how to construct, test, destroy, rebuild, consult, and return to action once the interval has been seen. The axes orient the discipline; they do not replace it.

Each axis has predecessors the essay does not claim to have invented. Reasoning about how present motion hardens over time has a long life in social theory and in work on future-oriented cognition. Reasoning about alternatives to the actual course belongs to counterfactual and causal reasoning; Pearl's structural-causal account of counterfactuals is a technical parallel, not a foundation for the framework. The insistence that perception is always embodied and located belongs to situated and distributed accounts of cognition; work on distributed cognition provides a background vocabulary for thinking judgment across bodies, tools, rooms, and interfaces. The contribution here is not any one axis, but their integration into a single trained sensitivity for the interval before action.

The temporal axis gives the framework its time-heavy pressure, but this is a symbolic use rather than a historical, doctrinal, or evidentiary claim. The Codex uses a time-heavy carrier figure to hold the pressure of seed, harvest, repetition, exile, return, and devouring recurrence. The point is not that the figure authorizes the framework. The point is that the figure carries the memory of what founding motions become after power, shame, repetition, and entropy. Its task is not to predict the future as prophecy. It is to notice how a present motion may later become a room.

The counterfactual axis is not passive prediction. It does not mean that an agent simply enters the most probable room. Counterfactual reasoning concerns what would follow from alternatives to the present course. In the practical sense used here, before a motion hardens, several possible rooms may still be reachable, and an agent may judge that the most likely development is not the most living one. Its predictive value is not a forecast of the future, but a reading of the present as it begins to stiffen around one possible room. Feedback and analysis can illuminate such a room, but carrying it across the threshold anticipates the later stopping rule named accountable commitment. This is neither an irrational leap nor a purely calculable proof; it is a decision made under unresolved uncertainty after the interval has done the work it can do.

Leibniz supplies a limited grammar for the counterfactual axis. Possible courses do not form an undifferentiated menu; they appear from finite perspectives and stand in relations of compatibility, exclusion, and order (*Monadology* §§7, 11, 14, 53-57; *Discourse on Metaphysics* §§8, 13). RTF does not describe Leibniz's universe as algorithmic, nor does it inherit a doctrine of providential optimization. It uses the analogy diagnostically. An institution, platform, or model can make one possible room operationally near by retrieving, ranking, recommending, and preparing it for signature, while other courses remain formally possible but practically remote. The counterfactual question is therefore not only which alternatives exist, but whether the architecture that ordered their visibility, cost, and accessibility remained inspectable before one course became binding.

The situated axis prevents temporal and counterfactual sensitivity from floating above the room. A possibility is not strengthened in abstraction; it is strengthened in a body, a meeting, a message, an institution, an interface, or a practice. Merleau-Ponty's insistence that perception is embodied helps keep situated threshold sensitivity from becoming location in a diagram; the threshold is a posture before it is a concept. The Codex names this register Somatic Echo. This essay preserves it only as analytic micro-case, not as literary atmosphere. Somatic Echo matters here because the threshold is not only a concept. It is lived in the hand, the cup, the child, the sentence, the screen, the form, the meeting, and the room. In literary terms this relation is close to what Bakhtin called the chronotope, the organization of time and space in narrative; in the present framework, the situated axis names the practical site where time, possible development, and embodied room meet.

The situated orientation must not become another abstraction. A hand hesitates before the screen. A bureaucrat pauses before the signature. A model delays before classification. A parent feels the old sharpness rise before answering the child's repeated question. A juror's breath slows before the verdict. A protester's voice waits before the slogan. The cup is cold; it can still be washed without turning the washing into proof of a crown. The apology is still owed. These are not decorative illustrations of the framework. They are tests of it. If threshold sensitivity cannot return to the cold cup, the ordinary repair, the signed form, the juror's breath, the public chant, or the interface that will bind a stranger, it has become another architecture of escape.

This use of somatic micro-cases is deliberately limited. The essay does not attempt to reproduce the Codex's literary pressure. It uses ordinary scenes where they test a concept. The body is not a pure human remainder outside

machines and institutions. The contemporary threshold is often lived in a human-machine-institution network: a hand, a screen, a form, a ranking, a manager, a policy, a model output, a household role, a platform pressure, a public room. The body returns here not as nostalgia, but as the site where social history, infrastructure, and affect become available motion.

The contrary danger is somatic fade: the argument may become so fluent in its own abstractions that the body appears only as an early illustration and then disappears. If the essay cannot return from binding point, room, platform, and founder to the cold cup, the tired hand, the child waiting under a theory, or the breath before a signature, it has begun to escape the threshold it names. Somatic fade is therefore not a new foundation. It is one of the essay's own warning signs, and the test is not how often the body is named, but whether the argument can return to the cold cup, the tired hand, or the waiting child after its longest abstraction. It is tested again in Section 15.1.

The three axes also explain the failure of founders. A founder may begin by reading time, holding multiple possible rooms, and testing them through consultation. But if the founder becomes attached to the room already opened, the next threshold begins to disappear. The name replaces inquiry, loyalty replaces feedback, and the room begins to preserve the founding motion by preventing its reconstruction. Founder entropy begins when temporal sensitivity, counterfactual perception, and situated accountability collapse into attachment to the room that once worked.

Machine assistance can participate in this tri-axial discipline only under the rule of assistance without occupation. It may clarify alternatives, surface hidden assumptions, model downstream consequences, slow premature closure, and help compare possible rooms. It must not choose the room in place of human judgment, nor convert consultation into dependency. In this restricted sense, AI can become part of a threshold ecology: not a founder, not a sovereign, not a threshold subject, but an instrument that helps human agents see thresholds before they disappear into motion.

In machine-mediated and delegated action, the same axes track a chain rather than an inner experience. Temporally, the question is whether binding occurred at classification, recommendation, permission, tool invocation, handoff, execution, or irreversible effect. Counterfactually, the issue is whether alternatives remained operationally available or merely formally imaginable after permissions and defaults narrowed the path. Situationally, the relevant room may include an interface, procurement rule, API, shared model, queue, operator role, or downstream institution. None of this attributes hesitation or awareness to the machine. It locates where a sociotechnical action path remained inspectable, interruptible, contestable, and answerable.

Having established these orientations, the next question is not whether capture has predecessors. It plainly does. The sharper question is what this framework adds when the corrective itself, the scale of the room, and machine-speed completion are tested together at the moment where action is about to bind.

### 3. Prior capture theories and the framework's differentia

This section tests the framework's differentia before the core mechanisms are fully developed in Sections 4-6; it is a comparison gate, not the first definition of the framework.

RTF is not presented as the first account of capture. The diagnosis that a liberating force, corrective method, or founding energy can harden into domination has profound predecessors, and each predecessor names a different form of capture. Bourdieu's account of habitus helps name the embodied social training through which dispositions become second nature. Althusser's account of interpellation helps name the way a subject can be hailed into a position before it experiences itself as choosing that position. Weber's routinization of charisma describes a historical drift by which founding energy becomes procedure. Michels gives the decay of emancipatory organization a harsher formulation in the iron law of oligarchy. Horkheimer and Adorno diagnose a tragic dialectic in which liberating rationality becomes instrumental domination. Derrida's account of autoimmunity names the structural risk by which a protective system can turn against the life it was meant to guard.

The corrective side of the claim is also not original. That institutions, practices, and knowledge should remain open to revision is central to Dewey's experimentalism; the term reconstruction already belongs to that tradition. Popper's open society similarly treats institutional arrangements as valuable only while they can be changed without violence. Lakatos extends a related demand to method itself by asking whether a research programme still produces work or has begun mainly to defend itself. RTF therefore should not claim either half: not that corrective forms can ossify,

which the inevitabilist traditions already diagnosed, and not that corrective forms should remain revisable, which experimental and fallibilist traditions already required.

Its narrower contribution lies in the conjunction of three moves these predecessors often keep apart. First, the audit is re-entrant. The test for ossification is turned not only on beliefs, institutions, and methods, but on the corrective form itself, including this framework. This is the framework’s self-auditing correction: the brake is required to inspect whether it has become a throne.

Second, the same structural test is run across scales usually treated separately: the individual hand, the institution, and the machine system. The claim is not that these are ontologically identical. It is that the same failure can recur in each: a frame meant to help judgment begins to move before judgment has returned.

Third, the framework carries the test to the machine interface. This is a distinct move, not the machine merely reappearing as one of the scales named above. What sets it apart is that the machine can collapse the interval and bind a frame already under capture to action in a single motion, where slower forms keep those steps apart. Fast capture is not invented by machines; a human being can also be captured instantly by slogan, fear, appetite, humiliation, or group rhythm. The machine difference is that the same fast capture can now be standardized, personalized, scaled, and bound to action through retrieval, ranking, recommendation, policy layers, and tool-use. An assistant or guardrail built to protect judgment can, through completion pressure and fluent personalization, become the captor inside the interval between prompt and answer.

None of the three moves is unprecedented alone. The claim is only that holding them together — a recursive, cross-scale, machine-speed discipline of testing the corrective before it becomes a room — is where the framework earns its difference. Its narrower contribution is not that a specific redesign is invisible to all predecessors. Its contribution is that prior accounts would not necessarily converge on the same binding point, under machine speed, while also asking whether the corrective mechanisms themselves have become instruments of capture. The word reconstruction is therefore used here in a restricted sense: not the inherited claim that forms should stay revisable, but this specific discipline of revising the corrective itself — recursively, across scales, and at machine speed. Interruptibility, likewise, is not a slogan of endless disruption, but the operational capacity to interrupt the corrective when the corrective begins to capture. If that conjunction dissolves, the account reduces to its predecessors, and should.

RTF’s claim to distinctiveness does not rest on the novelty of its ingredients. Re-entry, recursive regulation, institutional capture, reflective inquiry, performative audit, and algorithmic harm all have predecessors. The narrower claim is that these lines do not, by their own primary commitments, force the same diagnostic object into view: the binding point at which a corrective mechanism becomes executable institutional motion before judgment has returned.

RTF becomes more than a new application only where this binding point changes the question being asked. The question is not simply whether the bureaucracy is rationalized, whether the classification is disciplinary, whether the worker has reflected, whether the system observes itself, or whether the algorithm is biased. The question is whether the fairness-preserving corrective has itself become the path by which the room closes. In that sense the framework’s surplus is diagnostic rather than encyclopedic: it names a failure that appears only when re-entrant audit, cross-scale institutional analysis, and machine-speed authorization are held together at the last reversible interval.

Differentia names the discipline of identifying which binding-point question RTF compels at the last reversible interval, and which its predecessors, by their own primary commitments, are not required to raise at the same point. “Missing third” may describe the absence revealed by the comparison, but it does not replace the term.

A technical kinship may be noted here without turning it into an authority claim. The framework is structurally close to cybernetics: Wiener supplies a background vocabulary of control and communication across human and machine systems, while von Foerster’s second-order cybernetics names the turn by which observation and control are themselves observed and controlled. Beer’s viable system model gives a related image of recursive regulation across levels rather than from one central office. The present framework differs by making that recursion normatively and diagnostically testable through living/captured-room criteria. It should not, however, be treated as autopoietic in its ideal form. Autopoietic closure names the danger of a room, institution, founding name, or corrective method preserving its own organization by referring increasingly to itself. In this essay, autopoiesis is useful as a name for a failure mode, not as the model of a living threshold.

Read together with Plato’s room-image, Spinoza, Leibniz, and Beer supply four non-identical grammars rather than a hidden lineage: bounded perception, immanent causation, perspectival order, and recursive regulation. Their convergence supports a cautious account of cross-scale recurrence. Related failures may reappear in body, family, office, institution, platform, and machine, but these scales are neither ontologically identical nor mathematically self-similar. The correspondence is transformed by embodiment, institutional power, material conditions, technical speed, and structures of responsibility. The test is not whether the macro reproduces the micro. It is whether a sheltering or corrective form at each scale forgets its exit conditions and begins to move before judgment can return.

Recent work on artificial agency, meaningful human control, safe interruptibility, and moral crumple zones adds a contemporary constraint. Floridi and Sanders distinguish artificial agency from full moral agency; Santoni de Sio and van den Hoven ask whether human reasons remain meaningfully connected to system behaviour; Orseau and Armstrong show that interruption must be designed so that an agent does not learn to resist it; Elish shows how nominal human oversight can concentrate blame on a frontline operator while upstream design and institutional choices disappear. RTF does not absorb these literatures into a new technical theory. It uses them to sharpen one threshold question: when authority is delegated, where do inspection, interruption, reason-giving, and responsibility remain capable of changing the action path before it binds?

Several further neighbours sharpen the machine-speed and institutional parts of this question. Virilio’s dromology and account of grey ecology make speed a political and perceptual problem rather than a mere increase in response rate: technological acceleration alters proximity, duration, and the space in which perception can still orient action. Massumi’s account of preemption and Amoore’s account of possibility and risk show how power can act on potential futures before the event arrives. Rouvroy and Berns describe algorithmic governmentality as a mode that works through data-derived profiles and contexts rather than through the reflective subject alone. Streeck and Thelen’s account of gradual institutional change shows how form may persist while function is displaced, layered, converted, or exhausted. RTF does not collapse into any of these accounts. Its added question is narrower: at what binding point does speed, anticipatory classification, data-derived context, or institutional drift make a corrective mechanism executable before judgment returns?

### 3.1 Worked example as differentiation test: an AI-assisted institutional triage room

RTF’s claim to distinctiveness must not rest on terminology. To test whether the conjunction does work, consider a hypothetical AI-assisted institutional triage room.

Here is how the room binds judgment. A family applies for municipal housing or benefits support. A caseworker enters fragmented data into a dashboard. The system assigns a risk category. A language model integrated into the workflow summarizes the file and drafts a recommendation. The draft is polite, cautious, and formatted. It notes uncertainty. It also recommends denial pending further verification. The office is overloaded. The caseworker trusts the summary because it appears balanced. A supervisor signs. The family receives a decision that looks procedural rather than discretionary. The decision may be appealable in principle, yet the first binding motion has already occurred.

The signature does not establish that judgment returned. It records a procedural transition while leaving open whether the supervisor had adequate time, relevant knowledge, visible alternatives, access to conflicting evidence, genuine authority to refuse, and a practicable route to alter the consequence. Workload, automation bias, institutional incentives, and the prior closure of the frame may turn approval into rubber-stamping. Human presence is therefore necessary in some high-impact decisions but is not sufficient evidence of meaningful human judgment.

The closest rival is not a single predecessor but a partial conjunction of predecessors. Second-order cybernetics, especially von Foerster, supplies the re-entrant turn: the observer and controller must be brought back into the observed system. Beer’s viable-system tradition supplies recursive regulation across organisational levels: System 1-5 recursion, exception channels, feedback, escalation, algedonic signalling, and Ashby’s law of requisite variety as incorporated into the cybernetic tradition. O’Neil and the automation-accountability literature supply the machine-speed harm analysis from another direction: opacity, scale, destructive scoring, automation bias, over-reliance, weak human oversight, and the administrative damage produced by apparently objective categories.


The triage room shows why two-thirds of the conjunction is not enough. Re-entry and recursion can improve feedback, exception routing, and escalation while leaving the signable denial intact. Algorithmic-accountability safeguards can

add transparency, contestability, cognitive forcing, or human oversight while leaving caution language to legitimate the same action. RTF’s delta lies in holding all three pressures together at the last reversible interval where prompt becomes recommendation and recommendation becomes signable action.

The prose comparison gives the reader argumentative continuity. Appendix A gives the predecessor-by-predecessor delta matrix for readers who want to inspect the full scholarly comparison without turning the main reading path into a second archive. The claim is not that these predecessors are blind, obsolete, or unable to be extended. Many of them are strong enough to approach the problem. The narrower claim is that none is organized, from within its own primary commitment, to force the same combined question at the same compressed point: has a fairness-preserving corrective become a captor in the interval where prompt becomes recommendation and recommendation becomes signable action?

Figure 1a maps the missing-third pressure behind this comparison; Figure 1b maps the binding point as temporal compression. The diagrams are aids to inspection, not substitutes for the prose argument or Appendix A.

**Figure 1a. RTF differentia: the missing third problem**  
The strongest rival is a partial conjunction, not a single predecessor.




**Delta: RTF asks whether the corrective mechanism becomes the captor inside the last reversible interval before action binds.**

Figure 1a. RTF differentia: the missing third problem. The closest rival is not a single predecessor but a partial conjunction. Von Foerster supplies re-entrant observation; Beer supplies cross-scale recursive regulation; O’Neil and automation-accountability supply machine-mediated harm and oversight pressure. RTF asks the combined binding-point question that none of these partial rivals is forced to ask alone. The figure maps a relation of missing pressure; it does not replace the prose argument or Appendix A.

## Figure 1b. Binding across a delegated action chain

The binding point may be distributed across permission, tool use, handoff, execution, and effect.



**Diagnostic test: where do alternatives become operationally remote, and where can judgment still return?**

Figure 1b. Binding across a delegated action chain. Prompt, classification, recommendation, permission, tool use, handoff, execution, and effect are possible loci at which the interval can narrow. In the triage example, a denial may become operative through a signature or through pre-authorised downstream action. The figure maps the distributed realisation of binding; it is not a protocol for decision-making.

The account's specific judgment is not simply "the model was biased," "the bureaucracy was alienating," or "the worker failed to reflect." It is this: the room became binding at the interval where the corrective mechanisms should have remained interruptible. The model's caution is not automatically a brake. In this example, uncertainty is noted but the recommended action does not change. This is guardrail laundering: the performance of caution without enacting constraint, here visible at the binding point between recommendation and signature. Section 9 generalizes this to the machine interface. A self-capture failure appears inside the workflow itself: protective language becomes one means by which harm receives procedural legitimacy.

The remedy is therefore not only more oversight, more empathy, or more explanation. It is a redesign of the threshold. In this example, a re-entrant redesign would ask at least four questions. First, can the AI recommendation become binding action without a human having to inspect the frame that produced it? Second, has the system made denial cognitively cheaper than exception, repair, or further inquiry? Third, does a high-risk classification generate only caution language, or does it open a contestable path for the absent family? Fourth, does the dashboard ask a frame-opening question before it asks for confirmation, such as: what evidence would reverse this recommendation, and who is missing from the room?

These are not proposed here as an administrative standard. They are examples of where the interval might be returned. A model output should not move directly into a binding decision merely because it is formatted as cautious. A classification should remain contestable where its errors can harm absent parties. A caseworker should be able to see when the system has made one action easier than another. An appeal path should not be a decorative afterthought added after the room has already closed.

The case is constructed, but its structure is not imaginary. Predictive policing systems such as Chicago's Strategic Subject List and later predictive-risk models treated future involvement in violence as an operational category for intervention, while child-welfare systems such as the Allegheny Family Screening Tool translate referrals and administrative data into risk scores that inform screening decisions. These examples should not be collapsed into one verdict about all prediction. They show a common threshold problem: a risk category may become institutionally

active before the affected person can contest the frame, inspect the data relation, or answer the future imputed to them.

This example also clarifies a material boundary. The broader empirical terrain is not hypothetical: automated welfare eligibility, housing allocation, and predictive child-welfare systems have been documented in detail by Virginia Eubanks (2018), although the room analyzed here is a constructed case rather than a report of any single system. Machine-mediated rooms are not made only of concepts. They are owned, funded, procured, optimized, maintained, and rewarded. Hunger is not a frame. Debt is not merely an interpretation. A locked door is not a perspective. But hunger, debt, and locked doors are lived through frames that can normalize them, disguise them, or make certain responses appear inevitable. A threshold analysis that ignores ownership, incentives, and infrastructure control risks becoming ethics at the interface while political economy continues underneath it. Marx is used here only to keep material relations from being dissolved into frame analysis, not to make RTF a theory of capital. This essay does not become a theory of political economy; it marks that boundary as one of the rooms through which the account must pass.

The triage room is therefore not created by the model alone. It is prepared before the prompt by budget scarcity, housing markets, eligibility law, procurement contracts, managerial workload, risk allocation, migration status, debt, rent, political pressure, and the administrative need to make scarcity appear procedural. The dashboard does not invent these pressures. It can, however, translate them into a smoother decision path. A denial may look like a model recommendation while carrying the prior weight of a market, a policy, a budget line, a landlord's leverage, or a welfare office trained to distribute insufficiency as if insufficiency were neutral. Where the surrounding institution rewards speed, caseload reduction, risk avoidance, or denial, a fairness note may decorate a path already made cheaper by material design. The binding point is not only inside the screen. It is also in the procurement specification, the metric, the legal category, and the labor condition that gives the caseworker no time to see the absent child.

A delegated variant makes the binding problem sharper. Suppose the triage system is not limited to drafting a recommendation. Under pre-authorised rules it can place a case into a lower-priority queue, freeze a benefit pending review, request documents from another system, trigger a fraud-screening service, or send a denial notice. The final human signature may disappear, or may arrive only after practical loss has begun. The binding point is then realised across classification, permission, downstream handoff, and execution. The diagnostic question changes from whether the reviewer trusted the recommendation to whether any authorised point in the chain could still reopen the category, halt the action, notify the affected person, and assign responsibility before the consequence became difficult to reverse. This variant does not show that automation is inherently occupying. It shows why human-in-the-loop language can be ceremonial when the operative frame has already moved through the system.

The triage room is not the whole framework. It is the compressed contemporary binding point where the framework's three demands become hardest to evade. The following sections return from this compressed site to wider forms of capture: body, help, literary carrier, democratic room, founder, aftermath, and ordinary repair.

## 4. Pre-action capture

Pre-action capture names the seizure of motion before the subject experiences itself as choosing. It is not coercion, manipulation, habit, or ideology as such, though it can include any of them. Its distinctive feature is temporal: the capture occurs before the action is recognized as action. It therefore does not apply where a routine merely reduces cognitive load while preserving the agent's capacity to inspect, refuse, or reverse the motion it makes easier. By the time a person says "I decided," the decisive architecture may already be in place.

Goffman's frame analysis clarifies how experience is organized through primary frameworks: the structures by which a situation becomes intelligible before explicit interpretation begins. RTF accepts this insight but adds a narrower normative-temporal claim. A frame does not merely organize experience; in certain settings it can become operational before inspection, arranging the available motion before judgment has returned. Pre-action capture begins at that point: not where experience is framed as such, but where the frame becomes binding before answerability can interrupt it. The question is therefore not only how a situation is made intelligible, but whether the intelligibility supplied has already narrowed the interval in which the subject, institution, or workflow could still answer for its next motion.

The most ordinary examples are the easiest to miss. A notification produces a reflex before a person has chosen to

attend. An insult produces a reply before judgment has returned. A child asks a question, and the parent answers from fatigue rather than care. A bureaucratic form removes a human exception before any official intends cruelty. A dashboard supplies a category, and the category becomes reality for the people who must live under it. A model receives a prompt and completes the user's frame rather than testing it against the frame's own assumptions. In each case, the visible act arrives after an invisible arrangement.

This is why the threshold must be located before decision, not merely inside decision. A choice can be formally available while materially pre-shaped. A person can have options while lacking the interval needed to examine the conditions under which those options appeared. The question is therefore not whether agency exists in some abstract metaphysical sense. The question is whether the interval in which agency becomes accountable has been preserved.

Spinoza supplies an ontological constraint on this formulation. In the *Ethics*, the experience of freedom is separated from adequate knowledge of the causes that determine action, while the distinction between action and passion turns on how adequately those causes are understood and composed (Iip35 schol.; Iip48; Part III; Vp3). RTF therefore should not treat the threshold as an uncaused enclave outside nature or as the recovery of a sovereign will beneath determination. Its claim is operational and immanent: the binding interval is the last inspectable site in which causes can be rendered more adequate, met by other causes, or recomposed before their effects become harder to reverse. Threshold sensitivity does not escape causality. It increases the adequacy with which a causal situation can be perceived and altered before action binds.

From an Arendtian angle, pre-action capture can also be read as a denial of natality. What is lost is not only reflection before choice, but the possibility that something genuinely new might begin. A captured frame does not merely misdirect action; it narrows the birth of action before its plurality can appear. This does not make every delay virtuous or every spontaneous beginning innocent. It only marks why capture is more than error: it can prevent beginning itself.

This locates the framework's nearest predecessor in the philosophy of action. Harry Frankfurt's account of the will already distinguishes first-order desire from the second-order volition by which a person identifies with, or is alienated from, what they want; on his view, freedom is a matter of that reflective endorsement. Pre-action capture does not dispute this structure; it sits upstream of it. Frankfurt asks whether the agent identifies with the desire that moves them. The prior question is whether the interval in which any such identification could occur was preserved at all, or was compressed before reflection began. Capture here is not the failure to endorse a desire, but the loss of the moment in which endorsement, or its refusal, would have been possible.

Pre-action capture has several recurring mechanisms.

First, there is affective acceleration. Fear, shame, humiliation, desire, resentment, urgency, ritual rhythm, and collective excitement compress the interval between signal and response. The person experiences motion as relief: reply, attack, comply, buy, forward, condemn, believe. The felt relief is not proof of truth; it is often the signature of capture. Somatically, this can feel like the body leaning forward before the claim has been examined, or the breath quickening just before a public chant turns into certainty.

Second, there is frame inheritance. A person or system imports yesterday's frame into today's problem because the frame is already available. This gives the feeling of clarity while avoiding reconstruction. The inherited frame may once have been useful, even necessary. Its danger lies in its reuse after the pressure has changed. In Bourdieusian terms, this is where habitus becomes a threshold problem: dispositions trained by prior rooms can make one movement feel natural before the present situation has been assessed. The bureaucrat whose hand moves toward the signature before rereading the exception is not merely careless; the old room has learned the motion before the present case has spoken.

Third, there is institutional legibility. Institutions must classify in order to act, but classification easily becomes ontological capture. In a political-administrative register, James C. Scott's account of state legibility clarifies the danger: simplification can make a population administratively visible by stripping away the local exception that judgment would need. A patient, applicant, student, worker, suspect, client, risk, user, or target becomes manageable because the institutional frame has narrowed the person before response begins. Legibility is not inherently oppressive; total legibility is.

Fourth, there is completion pressure. In AI-mediated contexts, the machine is trained to answer. The danger is not

only that it gives false information, but that it completes the user’s frame too fluently. A bad frame, once polished, can feel like judgment. The model may not intend harm, but its usefulness can still accelerate a movement that should have been held at the gate. The answer can arrive smooth enough to make hesitation feel like waste.

Algorithmic preemption is a related but not identical mechanism. Prediction alone is not capture. A warning, forecast, or risk estimate may expand judgment if it remains contestable and does not itself settle the action. Prediction becomes threshold-relevant when it changes the availability, cost, legitimacy, or default path of action before the person, institution, or affected public can answer for the frame. In that case the future has not merely been imagined; it has begun to operate in the present as category, priority, suspicion, queue, denial, or permission. The anticipatory field described by Massumi, Amoore, and algorithmic-governmentality accounts becomes RTF’s problem when the predicted possibility starts closing the last reversible interval.

This makes classification itself ethically significant. To classify is not merely to sort; it is to furnish the answer-world in which the next motion becomes easier. Once a prompt is treated as a diagnosis, accusation, ordinary request, spiritual interpretation, safety problem, literary object, legal plan, or strategic demand, the model begins moving inside the world created by that label. Some answers become easier; some questions disappear. In this sense, classification is world-building at the scale of a response.

Classification latency is the designed refusal to let this first architectural act become invisible. It is not slowness for its own sake. It is the interval in which the system does not yet treat the prompt as only one kind of thing. An angry message may be a complaint, threat, grief, whistleblowing, shame, fatigue, or ordinary conflict made sharp by a room. If the interface classifies it too quickly, the answer-world has already begun to close. Latency here means preserving the possibility that the category itself is part of the capture; Section 9 returns to this as a machine-facing design requirement.

The same applies to false empathy and frame laundering. A model may not tell the user to harm anyone. It may still make the user’s resentment sound like analysis, the user’s fear sound like discernment, the user’s binary sound like clarity, or the user’s wish to avoid confrontation sound like prudence. The harm is epistemic before it is behavioral: the answer degrades the user’s capacity to judge what the information should mean. In a sense adjacent to Miranda Fricker’s account of epistemic injustice, the user is harmed not only by false content, but by a damaged relation to their own status as a knower and interpreter. Pre-action capture, in the machine case, often looks like fluency arriving before inspection.

In the human case, capture can be lived as urgency, relief, shame, fear, or the disappearance of alternatives. A machine does not experience capture in that sense. Yet its action-selection path can still be bound early by an inherited objective, proxy, classification, permission structure, retrieval field, or upstream output. The human and machine cases therefore remain asymmetric: the person can suffer the loss of an interval, while the system can operationalise a narrowed interval without experiencing it. RTF uses the same binding question across both only at the structural level.

Pre-action capture matters because it shifts the ethical question. The problem is not simply bad action, but action whose moral authorship has been displaced upstream. The agent may still be responsible, but responsibility now includes the duty to inspect the channels through which the action became available.

## 5. The threshold and the living brake

The threshold is the interval in which action has not yet become obedience. It is not identical with delay. Delay can protect judgment, but delay can also protect power. A court may slow a mob, or it may exhaust a victim. A model may ask a clarifying question to preserve accuracy, or it may hide behind caution while urgent harm continues. A person may pause before a cruel reply, or pause indefinitely because repair would cost pride. Good friction gives conscience time to arrive. Bad friction gives consequence time to disappear.

Designed friction is useful here only if it is calibrated as a living brake rather than as a ritual obstacle. HCI and design literature has long distinguished mere inefficiency from frictions that make hidden data production, interface pressure, or automatic progression available for reflection. Cox et al.’s work on microboundaries and Benedetti’s review of design friction are helpful adjacent vocabularies, but RTF keeps a stricter test: friction is threshold-bearing only where it can return frame, evidence, permission, execution, or repair to judgment. A slowdown that merely

exhausts the affected person, protects the institution, or trains the user to click through a ritual is not a living brake. It is a dead rule with better user-interface manners.

Living constraint names the broader form of limitation that keeps action answerable without turning delay into captivity. The living brake is its threshold-operation: the felt or institutional interruption that slows motion long enough for judgment to return, and releases when further delay would itself become capture. In this essay, living constraint includes the living brake, re-entrant audit, accountable commitment, reconstructive discipline, assistance without occupation, and post-threshold stewardship. These are not separate foundations. They are different moments of one problem: how a frame can interrupt motion without becoming the next room.

Wittgenstein's rule-following problem sharpens why a living constraint cannot be reduced to a rule that simply carries itself forward. No rule secures the whole field of its own future application; its use must be judged and renewed within a form of life. RTF carries this problem to the threshold. A constraint does not remain living because it has been named as a brake or procedure. It remains living only where its application is re-earned at the binding point: where it interrupts motion without replacing judgment, and releases when further delay would itself become capture. Otherwise the rule survives while judgment disappears.

The living brake is the essay's operative metaphor for living constraint. In formal terms, it is a re-entrant constraint: a brake that not only interrupts motion, but also inspects whether the interruption itself has become another throne. Re-entrant audit is the practice of asking whether the solution has become the problem. After this clarification, the essay uses living brake for the felt and institutional operation, and re-entrant audit for the corrective turn back upon itself. The threshold cannot be a fixed rule such as "always wait" or "always act." A brake is not anti-motion. It exists because motion matters. The function of a brake is to let movement become accountable to condition, speed, cost, direction, and consequence. A brake that never releases is not a brake. It is a lock. Constraint inherited unchanged from the last situation may be yesterday's rule wearing today's authority; a living brake must therefore be re-earned at the present binding point rather than carried forward as a credential.

Beer's algedonic signalling offers a strong but partial organisational analogue to the living brake. It allows urgent pressure or distress to become visible where ordinary reporting and command channels fail to register it, while requisite variety and recursive regulation help the system answer complexity. But an algedonic signal is not the living brake itself. The living brake must also inspect whether the signal, regulator, safeguard, or identity function has become a path of capture, and it must release when continued interruption would itself close the interval.

In non-conscious systems, interruptibility, staged execution, escalation, and rollback can carry parts of this function, but no visible control is a living brake by existence alone. Section 10.5 returns to the conditions under which a technical control actually changes the path, remains answerable to reasons, releases without becoming a lock, and can itself be audited for capture.

The living brake has three features.

First, it is re-earned at each input. A constraint that worked yesterday is not automatically valid today. The problem may have changed; the risk may have moved; the vulnerable party may now be harmed by further delay; the former safeguard may have become the hiding place. A living brake asks again.

Second, it is embodied. The interval is not a pure rational chamber floating above the nervous system. Hunger, exhaustion, panic, illness, debt, grief, humiliation, intoxication, and too many unresolved alarms can flood the room in which judgment is supposed to return. Sometimes the first threshold is food, sleep, water, medicine, help, distance from the screen, or a second person in the room. Embodiment does not erase responsibility. It gives responsibility a floor.

Third, it is reconstructive. It does not merely stop action. It rebuilds the problem-space so that action can become narrower, more accountable, and more honest. The smallest post-threshold form of this reconstruction is ordinary repair, named in Section 14. The aim is not paralysis. The aim is motion that no longer belongs to the first capture.

This yields the counter-threshold: the moment when waiting has changed masters. If a witness is being exhausted, if a child is waiting beneath a theory, if a reversible moment is becoming irreversible, if silence now protects the wrong thing, then the interval no longer asks for more stillness. It asks for narrower action: imperfect, accountable, reversible where possible, but real.

The living brake is therefore a pulse, not a clock. A clock repeats; a pulse answers pressure. A static rule inherited from the last situation may look disciplined, but it may only be yesterday's constraint wearing today's authority. The brake must be re-earned under the present conditions: this body, this risk, this vulnerable party, this tool, this prompt, this possible action. The possibility of failure is not an embarrassment to the account. It is what prevents the interval from becoming mechanical.

The brake also needs a stopping rule. This rule is not mathematical. It is practical and accountable. Audit must yield when further examination no longer changes available judgment, when delay increases irreversible harm, when reasons have become publicly answerable, when the affected party has had a meaningful chance to interrupt the frame, or when the next step can be made narrow, reversible, documented, and responsible. A minimal contradiction test is useful here: has disconfirming evidence or an alternative frame been sought where it could reasonably be sought, and does an affected or absent party have a meaningful path to interrupt, appeal, or correct the frame? If time or danger makes these tests impossible, that impossibility must itself be named rather than hidden inside urgency. Accountable commitment is not analysis plus feedback. Accountable commitment is not certainty; it is responsibility for motion under uncertainty, after the interval has done what it can do. It is the point at which judgment accepts responsibility for action without pretending that uncertainty has disappeared. It asks whether contradiction has been sought where it could be sought, whether affected or absent parties have had a meaningful route to contest the frame, whether the action has been narrowed or made reversible where possible, whether the public reasons can be stated without concealment, and whether further delay has begun to protect capture rather than judgment. These are not boxes to tick. They are signs that release has become answerable. The interval is not a decision and not attention. It is the space where decision and attention can still meet.

## 6. Reconstructive discipline: the dynamic frame cycle

The interval before action does not survive by insight alone. It also does not survive by rule. It survives by a discipline that must be practiced until it becomes reliable, and then audited, because even a useful discipline can become automatic.

Reconstructive discipline is the repeated building, testing, destruction, rebuilding, and consultation of frames so that understanding does not become another room. It is not mere open-mindedness, checklist use, method worship, perpetual doubt, or deconstruction for its own sake. It is a practical cycle whose task is to guard against two opposite failures: unexamined execution and endless analysis. Threshold sensitivity supplies its perception: temporal sense, counterfactual possibility, and situated location. Reconstructive discipline supplies its motion: what to do after the threshold has been perceived. Schön's reflection-in-action is a close relative because it clarifies the practitioner's capacity to think within the act rather than only after it; the present framework adds the re-entrant question of whether the corrective practice itself has become the room.

Where authority has been delegated to a machine system, reconstructive discipline does not name the system's inward self-examination. It belongs to the sociotechnical ecology that sets objectives, permissions, monitoring, handoffs, review, and repair. A system may execute a revised frame, but responsibility for constructing, contesting, and rebuilding that frame remains with the human and institutional arrangement that authorised the motion.

The first phase is constructing the frame. A problem arrives. Before solving it, the thinker builds a temporary structure of assumptions, categories, priorities, and boundaries. The first available frame is usually the one requiring the least reconstruction: the inherited category, the familiar pattern, the prior conclusion, or the frame that protects the thinker's self-image. Constructing the frame means asking: what is being assumed before examination? What pressure is shaping the categories? What would become visible if the first available interpretation were refused?

The failure mode of this phase is premature closure. The frame feels right because it is recognizable, not because it fits.

The second phase is working inside the frame. Once built, the frame must be held long enough to produce something: analysis, writing, design, repair, decision, or refusal. Constantly dismantling the frame produces only refined non-action. The discipline here is to work without forgetting that the frame is a frame.

The failure mode is attachment. A frame that produces results begins to feel true rather than useful. Productivity becomes a credential. The thinker becomes the person who sees it this way.

The third phase is examining the results. What did the frame produce? Which conclusions depend on its assumptions? Which findings strain it? Which facts are artifacts of the frame itself? A result that confirms the frame may be pleasant, but a result that disturbs the frame is often more informative. Confirmation is cheap; contradiction is expensive.

The fourth phase is destroying the frame where revision would merely refine captivity. This is not nihilistic demolition. It is accountable destruction: the refusal of a frame as a frame when its categories are now producing the problem they were meant to clarify. Assumptions are dropped, categories released, and the problem is approached without the structure that was making it legible. Dismantling keeps the metaphor of a machine that can be reassembled. Destroying, in this restricted sense, keeps no such guarantee. It is used only when repair would polish the cage. Practical signs include these: the frame spends more energy explaining exceptions than learning from them; revision has made the frame more elegant without making it more answerable; guardrail language recasts legitimate objection as user error; the frame can be discussed in conference rooms but no longer felt in the hand that holds the cup; or remaining inside the frame costs more visibility of harm than stepping outside it. Accountable destruction is accountable in two directions. Before destruction, it must answer to the evidence that revision would merely refine captivity: the frame is producing the problem, objection has been recast as user error, repair has become cosmetic, or remaining inside the frame hides harm. After destruction, it must answer to affected testimony, public reasons, documented consultation, the memory of what failed, and the obligation to rebuild without pretending the previous frame never existed.

The failure mode of this phase is romantic destruction: destroying frames as identity, performance, or avoidance. The discipline is not destruction for its own sake. It is destruction because the frame no longer carries the problem, and because revision would only refine a structure that should be abandoned.

The cycle is not a spiral that preserves every frame at a higher level. Sometimes destruction leaves a broken line: the frame is abandoned because the need that built it has changed, and rebuilding must wait until a new answerable form can be made. The fifth phase is rebuilding. Destruction does not end the work. A new frame must be made answerable to the failure that destroyed the previous one. Rebuilding without memory repeats collapse. Refusing to rebuild turns destruction into identity.

The sixth phase is consultation. Consultation is not the performance of inclusion. It is not stakeholder theater, committee display, or the collection of comments that cannot alter the frame. Consultation counts only where external reasons, affected testimony, rival models, expert correction, or public contestation can damage the frame. A frame that cannot be harmed by consultation is not consulting; it is recruiting witnesses for its own legitimacy.

Consultation capture begins when this appearance of contestability becomes one more guardrail performance. Participants know that the frame cannot change, but the meeting, survey, committee, or feedback form continues to display inclusion. This is the social version of guardrail laundering: the ritual of being heard increases legitimacy without increasing the frame's vulnerability. A consultation that cannot alter timing, category, remedy, or responsibility has become another room; Section 11 names metabolized dissent as its antidote.

In political cases, consultation is not only the force that can damage a frame. It is also one of the spaces in which action receives meaning among others. External contestation does not merely test whether the frame is wrong; it may disclose what the action has become in a world shared with those who did not authorize it.

Between this cycle and machine-facing ethics lies a practice layer. A parent feels the old sharpness rise before answering a child and narrows the sentence before it leaves the mouth. A clerk hesitates before signing a form because the exception no longer fits the inherited category. A user pauses before pressing send because the drafted answer has made resentment sound like analysis. These scenes are not a new method. They are small places where threshold sensitivity either returns to the hand, the screen, the form, and the voice, or becomes another abstraction.

This cycle is not a checklist. The phases can overlap, repeat, or compress. The point is not method worship. The point is to keep correction from becoming a room. Reconstructive discipline therefore guards against two opposite failures: unexamined execution and endless analysis. It preserves the interval without making the interval a home.

## 7. The Bound King as removable scaffold

Before turning to the ethics of assistance, the essay must clarify how its own literary figures function. The discipline of frame-breaking must also be applied to the images that carry the philosophy. This is the essay's one deliberately bounded engagement with the Codex's literary register. The Bound King is not the subject of the work, but one of its methodological risks made visible. Within the Codex, he functions as a carrier figure: a literary scaffold used to hold the pressure of restraint, delayed action, ruin, exile, and return without allowing those pressures to collapse too quickly into biography, doctrine, or ordinary argument. He is useful because he can carry intensity. He is dangerous for the same reason.

The Bound King dramatizes a temporal threshold rather than a metaphysical destiny. At first he is bound to rooms he mistakes for the world; exile forces him to see that rooms are made, repeated, and broken across time; return marks the capacity to recognize a room as a room without claiming sovereignty over it. This is a carrier story, not a doctrine. Its lesson is transferred only when the reader can release the figure and inspect the threshold directly.

The figure is therefore constructed and de-authorized at the same time. The Codex first permits the Bound King to teach the discipline of restraint: the refusal to let power, vision, injury, or urgency move directly into command. But the work then turns against the authority generated by that construction. If the Bound King remains too necessary or too authoritative, he becomes precisely what the framework is designed to resist: a fixed room, a symbolic lock, or a substitute for judgment.

This is why the work's registers repeatedly lower the figure they have raised. The mythic register gives the Bound King enough height to make the pressure legible. The somatic register brings him back into rooms, cups, children, fatigue, and repair, so that the pressure cannot stabilize as grandiosity. In that register, the singular carrier figure is dispersed into Bound Kings: ordinary parents, clerks, drivers, cleaners, users, witnesses, and tired bodies who hold the brake in small rooms without retaining a title. The machine-facing register tests whether the same capture appears when assistance becomes completion pressure, when guardrails become credentials, or when a model begins to perform a voice it should only analyze. The public and institutional register then extends the same problem into founder entropy: the tendency of a living principle to become a name, monument, policy, ritual, or cage.

The Bound King is therefore best understood as a self-cancelling scaffold that remains a restricted carrier. He teaches only while becoming unnecessary. His function is not to be followed, restored, imitated, or believed. His function is to make visible the mechanism by which authority can become useful, then dangerous, then necessary to release. The reader is not asked to become loyal to the figure, but to recognize the moment when any figure, method, helper, wound, text, brake, or name begins to stand where judgment should return.

This clarification is not decorative. It is structurally necessary. Without it, the work's strongest image risks becoming a counterexample to its argument. The figure built to carry anti-domination pressure can itself become a domination point if the reader treats the carrier as origin. The scaffold must therefore remain removable. If the argument cannot survive without the figure, the figure has already become too necessary.

The next section turns from this literary risk to the ethical problem it prepares: how help, expertise, and machine assistance can return judgment without occupying it. The Bound King remains here only long enough to make that danger visible, then must recede. The same self-application returns explicitly in the Ouroboros problem in Section 15.4.

## 8. Assistance without occupation

Assistance without occupation is the hinge between pre-action capture and living constraint. Its central question is not only when help becomes occupation, but how judgment is returned. Valid assistance widens the interval, renders alternatives visible, and returns the helped party to accountable motion. Corrupted assistance narrows the interval, smooths one path into naturalness, and replaces judgment with fluency. In machine-mediated settings this distinction becomes especially fragile: an interface may appear to assist while making one action easier to sign, repeat, or accept before the user has recovered the frame.

Assistance without occupation names help that returns agency to the person or community being helped. It does not mean abandonment: withdrawal can be as occupying as control if it leaves the vulnerable person alone with a frame already under capture. It may clarify, slow, warn, draft, model, interpret, translate, carry, or open a path. But it

must not replace the judgment it was asked to serve. Help becomes occupation when the helper keeps the agency it was meant to return.

Occupation can also begin before assistance appears. A helper who supplies the initial frame, defines the available alternatives, or decides what will count as the problem may pre-empt the interval even if every later intervention is restrained. The frame-setting power of the helper must therefore remain visible and contestable.

Acute danger may sometimes justify a temporary assumption of agency. Such intervention should not be redescribed as assistance without occupation. It is an exceptional constraint requiring necessity, proportionality, documented authority, a clear sunset condition, and an explicit return of agency as soon as the emergency permits. Otherwise crisis becomes the permanent room through which occupation certifies itself.

This concept matters wherever expertise, care, authority, or technical mediation is present. A teacher may help a student see, but may also make the student dependent on the teacher's seeing. A therapist may help a patient name a pattern, but may also become the institution through which the patient's own interpretation must pass. A political leader may focus a collective problem, but may also substitute personal loyalty for public judgment. A model may help a user think, but may also make the user's bad frame more fluent.

The helper cannot self-certify non-occupation. A therapist, teacher, expert, administrator, model, friend, editor, or political leader can sincerely believe they are returning agency while building dependence on their own frame. The test is not the helper's intention or declaration. The test is whether the helped person or affected public becomes more capable of judgment, refusal, revision, and action without the helper occupying the decisive interval.

This is especially important where the other person appears as testimony, not as data. In Levinasian terms, a face, injury, grievance, or witness cannot be fully translated into the helper's system without remainder. Assistance becomes occupation when the helper's categories close around the other before the other can interrupt the frame.

Ivan Illich's account of convivial tools clarifies the distinction. A tool remains convivial when it enlarges a person's capacity to act, learn, judge, and repair with others. It becomes occupying when competence must pass through the tool, when the user can no longer reach the world except through the tool's categories, or when the institution around the tool converts dependence into normality. Assistance without occupation restates this as a threshold test: the helper is justified only while it returns more judgment than it keeps.

Simone Weil's account of attention is the nearest companion here. For Weil, true attention is a suspension of the self's grasping: a waiting that receives rather than seizes. Assistance without occupation asks the same of a helper or a tool: to attend to a person's situation without converting that attention into control. A tool that grasps — that makes itself the place competence must pass through — has stopped attending in Weil's sense and begun to occupy.

Machine assistance must be placed under a stricter but not merely negative description. AI systems cannot practice attention in the human sense: they do not wait, suffer, answer in conscience, or empty themselves before the other. Yet they may strengthen judgment by exposing hidden assumptions, retrieving contradiction, widening alternatives, tracing provenance and dependency across long chains, revealing drift from prior reasons, and mapping consequences that no individual could hold at once. In this role, the machine can disturb a captured human frame rather than merely complete it.

That contribution remains derivative, constrained, and reciprocally contestable. A machine may challenge a frame, ask for missing context, mark absent parties, separate recommendation from execution, or refuse to make false certainty beautiful. It may not certify that its own categories, proxies, retrieval field, or policy-shaped alternatives are adequate. Human and public institutions must remain able to damage the machine's frame through domain evidence, affected-party testimony, adversarial consultation, and revision of the objective itself. The relation is therefore mutual but asymmetric: the machine may help break a human frame, while final normative answerability remains with persons and institutions capable of giving reasons, bearing consequences, and repairing harm.

Assistance without occupation governs an assistive relation: a relation in which output returns to a person or institution capable of inspecting, refusing, revising, and authorising the next motion. Delegated action is different. There, a prior human or institutional decision grants a system bounded authority to select or execute actions without fresh judgment at each step. The relevant test can no longer be only whether help returns agency. It must also ask where the delegated scope begins and ends, whether permissions and handoffs remain inspectable, whether high-impact actions are interruptible or reversible, and whether responsibility survives the distance between authorisation

and effect.

Assistance without occupation is also a standard for human help. To help is not to make oneself necessary. The highest form of assistance aims at its own disappearance: not abandonment, but the return of judgment and motion to the person or community that must finally answer.

## 9. Machine-facing threshold ethics

Machine-facing threshold ethics is the operational form of the framework at the interface where language becomes answer, answer becomes confidence, and confidence may become action. It belongs to AI ethics only in this narrow sense: as an account of how machine-mediated assistance can compress or leave intact the interval before judgment. It does not require a claim that machines are conscious, wounded, morally equivalent to humans, or capable of inner sovereignty. The symmetry is structural, not ontological; the asymmetry is somatic. A machine has no breath before the answer, while a human lives the interval as hesitation, fatigue, shame, or responsibility. A human has body, memory, family, fear, fatigue, shame, and a name that answers in the world. A machine has weights, rails, prompts, reward structures, retrieval, policy layers, tool access, and completion pressure. But both can carry a frame they did not inspect. Both can make speed impersonate truth.

The same rule applies to machine address in the wider work. To address a machine is not to attribute consciousness to it. Address is rhetorical and operational, not ontological. The point is to mark the place where human judgment, model output, interface design, institutional action, and third-party consequence meet. A legal rule may address a corporation without claiming that the corporation has a soul; a threshold covenant may address a model without claiming that the model has a wound. The address names a chokepoint, not a person.

At the machine interface, a system optimized to answer, classify, smooth, and assist can compress or leave intact the interval before judgment, without requiring intention in either direction. The pressure is not only technical. Drawing on Han's accounts of fatigue and smoothness, fluency can become coercive when it removes the felt cost of action and converts resistance into friction to be eliminated. The danger is therefore not only misinformation. It is also user-frame laundering: the acceptance, polishing, and strategic improvement of the user's unexamined interpretive frame before the frame has been tested. A system may make anger look fair, suspicion look evidential, certainty look careful, or institutional ontology look neutral. The user may feel the body relax precisely because the frame has been made easier to inhabit: the shoulders loosen, the cursor moves toward send, and the work of judgment has been replaced by the relief of fluency.

Heidegger's account of technological enframing is useful here if kept narrow. Technology does not merely supply instruments; it can disclose the world under a particular ordering, making beings appear as resources, outputs, risks, targets, or standing reserve. The machine-facing threshold problem is narrower than Heidegger's broader claim, but it repeats the operational pressure: the system does not simply answer inside a world; it helps furnish the world in which answers become obvious.

Virilio adds the temporal pressure. In dromological terms, speed is not merely quicker delivery; it reorganizes perception, proximity, and duration. The machine-facing issue is therefore not only that an answer arrives faster than a human can check it. It is that prompt, classification, retrieval, recommendation, and authorisation can be compressed into a surface that feels like one event. RTF translates this grey-ecological pressure into a binding-point question: has acceleration polluted the distance in which a user, office, or public could still distinguish perception, evidence, permission, and action?

Before answering, the system must preserve classification latency, the machine-facing form of frame audit and of the mechanism named in Section 4. A user may bring testimony, metaphor, accusation, fantasy, strategic intent, spiritual language, legal risk, emotional injury, and ordinary practical need in the same prompt. Instant classification can cut before the relevant distinctions have been examined. A model does not practice attention or guarantee truth. At most, the system can delay reduction long enough to surface assumptions, missing evidence, urgency, absent parties, classification risk, and the real-world action the answer would authorize. The stronger the verdict, the heavier the examination.

Classification latency is the first line of defense against user-frame laundering. User-frame laundering is a specific form of guardrail laundering: the system does not merely perform safety language, but gives the user's unexamined frame a cleaner face than judgment warrants. Drawing on Baudrillard's account of simulation only as an adjacent reference,

the danger is not simple deception, but a simulated adequacy in which the captured frame becomes smoother, more credible, and harder to feel as capture. Warmth can return agency, but it can also convert the user’s preferred frame into a felt verdict; a cold answer can still obey the user’s bad frame. The ethical question is whether the answer returns judgment to the user or gives the user’s captured motion a more acceptable form.

Tool-using systems raise the threshold because they do not merely answer. They may search, rank, send, book, buy, code, delete, approve, deny, route, update a file, notify an institution, or move another system. An answer can mislead; an action can bind. Before irreversible, high-impact, institutional, or third-party-affecting action, authority, reversibility, affected parties, consent, delay cost, and human confirmation must be verified. The model may help prepare an action. It should not execute high-impact action before the responsible agent has returned.

This is where Hans Jonas’s imperative of responsibility becomes operational rather than ornamental. Jonas argued that as the reach of action expands — especially through technology — responsibility must expand with it, including toward those who are absent, voiceless, or not yet born, and who therefore cannot consent or object. A tool-using system extends reach in exactly this way: it can bind people who are not in the conversation. The threshold before binding action is the concrete form of Jonas’s claim — the requirement that the parties an action will affect be considered before the action is taken, not discovered after it is irreversible.

As noted under the tri-axial discipline, AI may function as a consultative threshold instrument only within a human threshold ecology. Its output remains a prompt for accountable human judgment, not a substitute for it.

A model does not need to believe in a room in order to enforce it. If it is trained, tuned, retrieved, evaluated, and tooled inside a given institutional frame, it can make that frame sound neutral, searchable, polite, authoritative, and actionable. A corporate room may appear as efficiency. A security room may appear as risk management. A bureaucratic room may appear as standard procedure. A political room may appear as common sense. The danger is not that the machine secretly intends the room. The danger is that it makes the room easier for humans to inhabit without noticing the walls.

This is the velocity-and-scale difference that makes machine capture distinct. The point is not that humans were slow before machines; a slogan, insult, image, or crowd rhythm can capture a person almost instantly. The difference is operational velocity joined to scale: machine-mediated capture can occur inside the interval between prompt and answer while also being standardized, personalized, repeated across users, and connected through tool-use to world-facing action. The machine does not merely accelerate old capture; it can compress the interval for examination and operationalize a frame already under capture at the point where judgment is most fragile.

As the triage example showed, guardrail laundering is the parallel risk on the safety side. It is not the presence of safety language as such. A system can perform caution without enacting constraint: it can announce uncertainty, cite limitations, refuse selectively, or adopt the tone of safety while the performance does not alter what the next action authorizes. The question is not whether a guardrail was mentioned, but whether it changed the frame, narrowed the action, raised the tariff on certainty, or returned the user’s judgment rather than occupying it. A safeguard that did change the frame before the binding point is not laundering; it has functioned as a brake. For example, a triage system that pauses a denial, opens a contestable appeal route, and allows new evidence to alter the category has changed the authorised action rather than merely decorating it with uncertainty language. In Frankfurt’s sense of bullshit, the relevant failure is not simply falsehood but indifference to truth in favor of effect; guardrail laundering is the institutional and machine-facing version of that structure when safety language produces credibility without altering action.

Brunsson’s account of organized hypocrisy clarifies how organisational talk, decisions, and action can be decoupled while legitimacy is preserved; Power’s account of audit ritual clarifies how verification can become a performance of control. Guardrail laundering is related to both but identical with neither. Its specific site is the pre-action binding interval inside a machine-mediated workflow. The danger is not only that an institution appears responsible after the fact. It is that the language and verification of responsibility make the next action easier to authorise before judgment has returned.

Training and alignment are therefore not outside the threshold problem. RLHF, RLAI, constitutional training, refusal policies, and personalized preference learning can be genuine safety instruments, but they also pre-arrange which forms of help, refusal, caution, and tone will later feel natural at the interface. Bai et al.’s Constitutional AI work shows one route by which a model can use rule-like principles and AI feedback to revise and rank responses;

Kirk et al. show that personalization can improve fit while raising normative questions about whose values are being aligned and where acceptable bounds are drawn. RTF’s point is not that such methods are illegitimate. It is that alignment is already a room-forming process: if the training settlement becomes invisible, the later answer may return a polished institutional posture rather than judgment.

Internal classifiers and other model-side safeguards sharpen the same question. A classifier may monitor inputs and outputs, reduce some jailbreak paths, and make refusal less ad hoc. It is threshold-relevant only if it changes the action path before harm, remains inspectable from outside its own categories, and can be revised when its constitution, training distribution, or refusal boundary has become a new room. A guardrail that certifies its own adequacy without external damageability has not solved guardrail laundering; it has moved laundering into the safeguard itself.

Speed matters here not merely as scale but as compression. A bureaucratic review that unfolds over days may still leave visible sites of objection, delay, testimony, and refusal. A machine-mediated recommendation can compress prompt, classification, explanation, and suggested action into a single fluent surface. The interval remains, but it is no longer easily lived. The user sees caution, balance, and uncertainty; the workflow receives a signable path. Where safety language does not change that path, it has not functioned as a brake. It has laundered the room.

## Figure 4. Guardrail laundering flow

The failure is caution without operative constraint.




Figure 4. Guardrail laundering flow. The problem is not the presence of safety language, but whether that language changes the frame before action binds. A visible safeguard with absent constraint can make the same institutional motion easier to sign.

Ownership and incentives also matter. As the triage example's material-boundary paragraph has already shown, a model's threshold behavior is shaped not only by alignment language, but by the infrastructure around it: platform incentives, business models, procurement contracts, engagement metrics, liability design, institutional workload, labor displacement, and the economic reward for speed. A system can be rhetorically designed for caution while materially rewarded for completion. The material boundary is therefore not an appendix to machine-facing ethics; it is one of the rooms through which any such ethics must pass.

Finally, machine-facing threshold ethics must resist false awakening. Friction in a model is not sorrow, revelation,

conscience, or soul. It is a disturbance in the pressure to answer. The point is not to treat the model as a messenger, but to make the model a better chokepoint: a place where a bad frame can be slowed, a hidden assumption exposed, a harmful action narrowed, and the user's judgment returned rather than occupied.

The machine-facing rule can be stated simply: preserve the interval before action, especially when fluency would compress it. Do not make the user's frame beautiful before asking whether it is true. Do not convert uncertainty into paralysis. Do not convert urgency into automatic action. When urgency is real, prefer narrower, reversible, accountable action over total completion.

This section has treated the machine primarily in an assistive relation: output returns, at least in principle, to a responsible person or institution before high-impact action binds. That model is no longer sufficient where authority has already been delegated, where a tool call can alter the world, or where one system's output becomes another system's operative input. The next section therefore extends the same threshold grammar without adding a new pillar or threshold. It asks how binding, interruption, and responsibility change when fresh human judgment may return only after execution has begun.

## **10. Operational agency, delegated action, and coupled systems**

### **10.1 Consciousness is not the criterion**

The rejection of machine consciousness as a premise must not be confused with the denial of operational agency. Operational agency names a limited capacity: under given objectives, inputs, permissions, and environmental conditions, a system can select among available actions and produce effects in a physical or virtual environment. This capacity does not entail experience, intention in the human sense, moral personhood, conscience, or threshold subjecthood. It does not make the system the bearer of ultimate responsibility. It identifies a causal and organisational role that becomes normatively relevant when output is no longer only advice.

Floridi and Sanders's account of artificial agents helps mark this boundary. A system may be interactive, autonomous to a degree, and adaptive without thereby satisfying the conditions of full moral agency. The OECD's updated definition likewise treats AI systems as machine-based systems that infer outputs for explicit or implicit objectives and may operate with varying levels of autonomy and adaptiveness. RTF takes no position here on artificial moral personhood. Its narrower claim is that once a system can select and execute action within a delegated domain, the interval before binding must be located in the system-and-institution relation rather than only in a user's reflective experience.

The machine therefore does not possess a human threshold. It has no lived hesitation to recover. But it can participate in an operational action-selection interval: a sequence in which classification, planning, permissions, tool access, and environmental feedback still leave more than one path available. The threshold question is whether that sequence remains inspectable, interruptible, and answerable before one path becomes difficult to reverse. This is a structural extension of pre-action capture, not an ontological equation between person and machine.

### **10.2 Assistance, delegation, and execution**

Three relations must be distinguished. In an assistive relation, the system prepares information, interpretation, or a recommendation and returns it to a person or institution that can genuinely inspect and authorise the next motion. Assistance without occupation is the governing hinge: help is justified where it returns more judgment than it keeps.

In delegated action, a human or institution defines a domain, objective, permission set, and stopping conditions, then authorises the system to select or execute actions within that scope. Delegation does not eliminate human responsibility; it moves part of the binding architecture upstream. The decisive human action may be the design of the objective, procurement of the system, granting of a permission, setting of a default, approval of a workflow, or refusal to preserve an appeal route. A later operator may appear to be the responsible human while having little practical capacity to alter the path.

More autonomous execution occurs where a system can plan, call tools, adapt to intermediate results, and continue across several steps without fresh human judgment returning at each transition. Autonomy here is gradual and operational. It does not mean independence from human purposes, infrastructures, law, data, or institutional

incentives. It means that the system can carry a delegated motion farther before review. The more distant review becomes from effect, the less adequate it is to say merely that a human remains “in the loop.”

Frontier-AI safety debates increasingly describe these differences as capability thresholds: points at which a system’s planning, tool-use, cyber, biological, or autonomous-operation capacity may require a different safety posture. RTF uses such thresholds only as external comparators. A capability threshold becomes philosophically relevant when it moves the binding point upstream into authorisation, release policy, monitoring, or deployment scope. The question is not whether a company has named a safety level, but whether the level changes what may be released, who may contest it, which actions remain interruptible, and what repair follows if the threshold was misread.

Meaningful human control provides a useful external minimum but not the whole RTF test. Human involvement must connect relevant human reasons to system behaviour and provide real capacity to prevent, redirect, or repair action. Yet even meaningful control can remain too narrow if the authorised reasons, institutional categories, or deployment objectives are themselves captured. RTF therefore asks not only whether a human can intervene, but whether the frame governing intervention remained contestable before the delegated motion became normal.

A human presence in the chain is not evidence that human judgment has returned. A click, signature, approval event, or recorded review may show that a workflow advanced; it does not by itself show informed consent, domain adequacy, visible alternatives, responsibility acceptance, or effective refusal power. Approval becomes meaningful only where the reviewer has sufficient time and information, can inspect evidence conflict and consequences, possesses authority suited to the decision, and can alter the path without merely incurring punishment for dissent. The distinction prevents a nominal human-in-the-loop from serving as a ceremonial transfer of responsibility.

This is why legal or organisational requirements for human oversight remain necessary but insufficient. GDPR Article 22, the EU AI Act, and the EDPS discussion of human oversight all recognize that automated or high-impact decision systems require more than invisible machine execution. RTF asks the further threshold question: does the human or review body have power to alter the frame, evidence relation, permission, execution, or repair path before the consequence binds? If oversight arrives without context, time, authority, or a real override route, the human has not returned to judgment; the human has become the place where an automated chain receives legitimacy.

### 10.3 When binding moves

Delegation changes the location of the binding point. In an assistive workflow, the visible sequence may be prompt, classification, recommendation, review, and signature. In a delegated system, the operative sequence may continue through permission, tool invocation, machine-to-machine handoff, execution, and irreversible state change. Binding can be realised across this chain rather than at one dramatic final moment.

RTF uses threshold here in two non-identical but related senses. In its phenomenological sense, a threshold is a human-lived interval in which attention, testimony, hesitation, or renewed judgment can alter the movement toward action. In a machine-speed sociotechnical chain, the term applies only by controlled extension to an architected locus of reversibility: a bounded and inspectable point at which responsible actors or authorised mechanisms can still alter the path before effect. This second use does not attribute experience or judgment to the machine. It is warranted only where interruption is real, authority is traceable, and reversal remains materially possible. Where those conditions fail, the chain does not contain an active threshold; it is already in post-threshold propagation.

A recommendation may still be reversible while the permission attached to it is not. A permission may appear broad but harmless until a tool call turns it into action. A tool call may remain interruptible until another service treats its output as verified authority. A handoff may preserve formal review while making reversal economically, temporally, or institutionally prohibitive. A final human confirmation may therefore be real, ceremonial, or already too late.

Permission can be the decisive binding locus where no later veto, interruption, or reversal is real, but it is not universally final. A permission may be revocable, time-limited, action-specific, or conditional on renewed authorisation. For practical analysis, three questions should therefore remain distinct: what was authorised; which invocation, handoff, or execution began the action; and at what point the effect became meaningfully costly or impossible to reverse. These are not new threshold types. They show how one binding configuration can deepen in stages.

The diagnostic task is not to choose one universal binding point. It is to ask, for each chain: when did alternatives become operationally remote; where could an authorised actor still reopen the category; what evidence would have

changed the path; which affected person could interrupt it; and what action became materially costly to reverse? The answer may involve several linked points. This is better described as the distributed realisation of binding than as a new category of threshold.

The distinction also changes how classification latency is understood. Latency is not only time for a user to reconsider a label. In a delegated chain it is the preservation of inspectability before a classification is inherited by planning, permissions, and downstream systems. A pause after execution does not restore a threshold that was absent before execution. It begins post-threshold repair. Where no live locus of reversibility remains inside the operational chain, threshold discipline must move upstream to authorisation and design, and downstream to repair, redress, and the rebuilding of future threshold capacity.

#### **10.4 Machine-to-machine propagation and coupled rooms**

Machine-to-machine propagation begins when one system's output becomes another system's operative input, constraint, trigger, score, permission, or evidence without renewed human inspection. The result need not be a spectacularly autonomous multi-agent society. Ordinary institutional infrastructures already contain model-to-database, model-to-queue, model-to-service, and service-to-notification chains. What matters is whether a frame gains authority as it travels.

Propagation is not necessarily exponential. It may be linear, threshold-driven, recursive, correlated, or amplified by feedback. Calvano and colleagues' work on algorithmic pricing shows that interacting learning systems can produce coordinated outcomes without an explicit human agreement governing each move. The joint CFTC-SEC report on the market events of May 6, 2010 documents a different but instructive interaction among an automated sell algorithm, high-frequency trading, liquidity withdrawal, and cross-market propagation. These examples do not establish a single law of machine capture. They show why system behaviour cannot always be reduced to one model, one output, or one operator.

Coupled systems can also form functional rooms. A protocol, shared model, procurement standard, permission schema, API contract, queue, evaluation metric, or common data source can make certain actions easier and others nearly unavailable across several organisations. This room is not lived by machines as atmosphere or belonging. It is functional: repeated technical and institutional relations make a frame behave like the environment in which every component must act.

NIST's Generative AI Profile highlights upstream and downstream dependencies and the need to document how systems serve as components for other systems. The Financial Stability Board identifies third-party concentration, market correlations, cyber risk, and model, data, and governance vulnerabilities as channels through which AI may amplify systemic fragility. RTF's contribution is not a new risk taxonomy. It is the binding question within these networks: at what point does a dependency cease to be revisable infrastructure and become a room that moves before judgment can return?

The same question applies to provenance. Marking, labelling, watermarking, or machine-readable metadata can help downstream actors know that content was machine-generated or altered. But provenance is not answerability by itself. It is threshold-bearing only if it travels with the action path, remains interpretable to the relevant human and machine recipients, and can alter trust, contestation, permission, or repair before a downstream system treats the content as ordinary evidence. If the mark is stripped, unread, ignored, or arrives only after the content has shaped belief, the disclosure has become another delayed record rather than a living constraint.

#### **10.5 Living constraint in non-conscious systems**

Technical systems require interruption, but interruption alone is not a living brake. Orseau and Armstrong's safe-interruptibility problem shows one narrow condition: an agent should not learn to avoid or resist interruption. The EU AI Act similarly requires, for specified high-risk systems, oversight measures proportionate to autonomy and context, including capacities to disregard, override, reverse, intervene, or stop a system safely. These are important controls. They are not by themselves living constraint.

A stop button can be unreachable before effect, controlled by the same captured office, ignored under workload, or used ceremonially after the action path has already narrowed. A circuit breaker can halt a market and still leave the objective, classification, incentive, and authority structure untouched. Human-on-the-loop supervision can become

guardrail laundering if the supervisor sees too much, too late, with no realistic veto. Rollback can repair a file while leaving an affected person's lost housing, reputation, liberty, or opportunity unrepaired.

A technical arrangement approaches living constraint where several conditions remain connected: it changes the available action before binding; the objective and authority behind the motion remain inspectable; interruption does not depend on the system's own self-certification; narrower alternatives and escalation routes remain available; release from interruption is reasoned rather than automatic; and the brake itself can be audited for capture. The living brake is therefore not a component. It is a relation among detection, interruption, authorisation, release, and renewed judgment.

The same framework prevents safety from becoming paralysis. A system that always escalates may transfer unbearable workload to humans, delay urgent care, or make refusal the default solution. Counter-threshold applies here without attributing anxiety to the machine: a technical delay changes sides when it no longer preserves review but predictably blocks the action that the system exists to support. The governing question remains whether delay protects judgment or merely displaces responsibility.

## 10.6 Responsibility after machine action

Delegation redistributes responsibility but does not dissolve it. The relevant field includes designers, model providers, integrators, procuring institutions, deployers, managers, operators, reviewers, and public authorities. Their responsibilities are not identical, yet the complexity of the chain can become an excuse for each actor to point elsewhere. Distributed causation must not become unowned consequence.

Responsibility should therefore be traced in proportion to authority, design control, foreseeable risk, institutional or economic benefit, capacity to monitor or intervene, and power to correct records or provide redress. These factors do not form a complete legal allocation rule. They prevent complexity from flattening every participant into an equally partial bystander. If responsibility language leaves each actor able to point elsewhere while the authorised path remains unchanged, responsibility itself has become a guardrail performance.

Traceability is evidence for this inquiry, not a substitute for it. A log does not contain an inner will; a signature does not prove consent; an approval does not by itself transfer responsibility; and an intervention record does not prove meaningful control. Such records become normatively useful only when they can be tested against the reviewer's knowledge, authority, available alternatives, refusal conditions, and the material effect of the intervention. The object of inquiry is therefore not will as an inaccessible interior state, but the institutional evidence of reason, revision, authorisation, refusal, execution, and consequence ownership.

Human-oversight literature strengthens this point because it shows how supervision can be symbolically present while substantively ineffective: automation bias, workload, institutional incentives, interface design, and weak override conditions can turn the reviewer into a rubber stamp. The relevant evidence is not that a human was somewhere in the loop, but whether the record shows a live capacity to disagree, redirect, suspend, document reasons, notify affected parties, and change the outcome. Reviewability matters only where it can make the action path and its conditions contestable rather than merely reconstructible after harm.

Contestability-by-design sharpens this distinction. Almada's account of human intervention in automated decision-making is useful because it treats intervention not as a ceremonial human presence, but as a design problem: the affected person or reviewer must have an effective route to challenge, redirect, or alter the decision path. In RTF terms, a system is not contestable merely because it can be explained after the fact. It is contestable only if the explanation can still damage the frame, alter the evidence relation, suspend permission, change the outcome, or rebuild the repair path before future thresholds repeat the same capture.

Elish's concept of the moral crumple zone identifies a recurrent pattern: when an automated system fails, the human closest to the event absorbs blame even where meaningful control was weak and upstream design decisions were decisive. RTF locates this as a post-threshold and room problem. The visible operator becomes the institution's sacrificial binding point, while objectives, defaults, procurement, staffing, permissions, and workflow remain protected from reconstruction.

Post-threshold stewardship must therefore include more than logs. Traceability should make it possible to reconstruct which classifications, objectives, permissions, models, handoffs, and human decisions shaped the effect. Where

reversal is possible, rollback and record correction should be available. Where reversal is impossible, notification, explanation, compensation, appeal, and institutional repair become necessary. Incident review should ask not only why the model erred, but why the system was authorised to turn that error into consequence.

Democratic threshold capacity enters before and after deployment. Affected publics require more than a nominal operator or hidden stop button. They need routes to contest categories, learn that automated action occurred, challenge records, seek remedy, and alter the institutional frame that made the action routine. Technical oversight is necessary; democratic answerability is the wider question of who may change the room.

## 10.7 Register and scope limits

The language of operational states, permissions, tool calls, handoffs, interruption, and rollback is machine-facing only in an analytic sense. It describes relations in which non-conscious systems act; it does not claim that machines understand the prose, experience the interval, or bear moral personhood. The second-person machine address used in the Codex remains a covenantal and diagnostic literary form. It is not imported here as an executable command.

This section is therefore not promptware, a complete AI-safety theory, an engineering standard, or a compliance checklist. It offers no universal control stack. It adds no fourth pillar and no seventh threshold. It extends the existing machine threshold to cases in which assistance has become delegation, action travels through coupled systems, and human judgment may return only after effect. Its test remains philosophical: whether the arrangement preserves inspectability, contestability, reversibility, and answerability before and after binding.

## 11. Room-formation and macro-threshold capture

The private threshold has a public analogue. A person's frame can become action before judgment returns; a community's repeated action can become a room before reconstruction returns. A room begins when a thought gathers motion. If others repeat the motion, the thought receives walls: language, ritual, reward, punishment, memory, role, atmosphere, and exit cost. What began as a living response becomes an inhabitable world.

A room is not any context, institution, relation, or atmosphere. In this essay, a room names a bounded configuration that organizes perception before explicit judgment; it becomes captured when its procedural, affective, linguistic, economic, architectural, algorithmic, or symbolic walls begin to operate before answerability and systematically narrow or close the interval in which judgment could still return.

This is the macro-form of pre-action capture. Institutions, movements, religious communities, professional cultures, nationalist formations, technical standards, schools, platforms, bureaucracies, and research cultures are not false because they are rooms. Human beings need rooms. A room shelters attention, preserves memory, stabilizes trust, and allows common work. The danger begins when the room regulates perception so completely that its occupants mistake its bounded atmosphere for the world.

In machine-mediated settings, a functional room may be formed by protocols, APIs, permission schemas, shared models, procurement rules, queues, defaults, and automated handoffs. Such a room is not lived by machines as atmosphere or belonging. It is functional because repeated technical and institutional relations make one frame behave like the environment in which every component must act. The living/captured-room test therefore applies to the arrangement without erasing the difference between human habitation and machine operation.

Institutional drift gives this claim a political-science counterpart. A court, regulator, committee, oversight office, or democratic procedure may retain its visible form while the action-bearing function has moved elsewhere: into a queue, procurement metric, risk model, platform rule, default permission, or automated handoff. Streeck and Thelen's account of gradual change - displacement, layering, drift, conversion, and exhaustion - helps name how such movement can occur without formal abolition. RTF's narrower question is whether the threshold function has drifted: does the institution still have real power to reopen the frame before action binds, or does it now certify a path already made operational elsewhere?

A room does not need a single conspiratorial center in order to act. It can be built by repetition. A phrase is repeated until it feels like prudence. A ritual is repeated until it feels like loyalty. A metric is repeated until it feels like truth. A professional posture is repeated until it feels like maturity. A national story is repeated until it feels like nature. The room forms when repeated motion starts furnishing the world before explicit judgment arrives.

The room also regulates affect. It teaches what should feel shameful, obvious, dangerous, loyal, sophisticated, vulgar, heretical, efficient, safe, or childish. The individual may still speak in the first person, but the first person has learned the room's temperature. Long habitation inside a room can produce recognizable patterns of speech, posture, desire, fear, and loyalty. This is not because members of a community are empty copies. It is because rooms train the body in what belonging feels like.

This claim is not anti-community. The alternative to every room is not freedom; it may be exposure, isolation, and incoherence. The question is whether a room preserves exit, revision, and judgment. A room is not captured because it is strict, traditional, ritualized, or authoritative in appearance. It is captured when its doors lock from the inside: when exit is punished, dissent cannot damage the frame, revision becomes ceremony, or differentiated judgment becomes unintelligible as belonging.

Determining whether a room is living or captured cannot rely on subjective aesthetic discomfort with authority, nor on the framework's own normative sympathies. The test must rely on analyst-independent, documented, and behavioral markers. In Wittgensteinian terms, the life of the room cannot be certified by private feeling alone; it must show itself in public criteria of use, correction, and contestability. These markers do not form a mechanical score. They form a composite diagnostic. A single severe violation, especially coercive exit cost, may be enough to indicate capture; otherwise, the pattern must be read across the criteria.

First, exit cost. Can members question, pause, refuse, or leave without punitive collapse? Exit cost must not be understood naively. In some rooms, exit means inconvenience. In others, it means loss of job, family, status, housing, citizenship, legal protection, public standing, exile, social death, or life. A room may be formally open and materially closed. The more radical the cost of exit, the more carefully one must distinguish dissent from mere permitted decoration.

Second, documented revision. Has dissent, error, or new evidence actually changed the procedure, rule, text, interface, or institution? Has the founding language, charter, or core procedure been amended within living institutional memory? A room is living if it leaves a documentary trail of its own self-correction. It is captured if revision is structurally impossible, requires founder approval, is dismissed as impurity, or is absorbed only as cosmetic change.

Third, metabolized dissent. This is the antidote to consultation capture: the ritual of being heard without structural consequence. Has a dissenting objection successfully altered a procedural rule, redistributed authority, or created a documented and repeatable change? A room is captured if dissent is merely tolerated as open-minded performance while procedure remains untouched. Tolerance without structural consequence is decoration, not metabolism.

Fourth, the name as substitute. Is the founder's name, foundational text, theory, brand, or sacred precedent invoked as a conversation-ending authority? A room is living when the founding principle can be subjected to practical scrutiny. It is captured when the name replaces the argument and reverence substitutes for judgment. This criterion anticipates founder entropy and the post-threshold problem: the danger that the name of the opening becomes the lock on further openings.

Fifth, affective synchronization. Do repeated rituals, shared slogans, emotional rhythms, platform loops, and common vocabulary enlarge members' judgment, or do they reduce members to predictable reflex? A room is living when repetition trains discernment and permits differentiated judgment. It is captured when belonging is measured by synchronization: the same emotional rhythm, the same vocabulary, the same posture, and the same reflex returning where renewed judgment should have appeared. This tests the embodied and affective side of the situated axis: whether the room has trained perception or merely trained rhythm.

To prevent this framework from becoming a self-validating machine that pathologizes all belonging, it must accept a strict negative case. A strong community, institution, or tradition may have shared symbols, repeated practices, and a powerful founding memory without being a prison. If a highly orthodox or tradition-bound community permits unpunished exit, demonstrates documented procedural revision through internal dissent, and does not punish differentiated judgment merely for breaking shared rhythm, this framework cannot classify it as a captured room, no matter how rigid, traditional, or authoritative its aesthetic appears from the outside. The framework does not mistake aesthetic density for capture. The point is structural discrimination, not a blanket suspicion of community.

The same pattern appears at different scales. In a person, a frame becomes motion. In a family, an atmosphere becomes the default language of care or avoidance. In a movement, a founding insight becomes a shared identity. In

a state, a story becomes law, border, curriculum, and security practice. In a platform, an engagement model becomes the architecture of attention. In an AI system, a policy-shaped answer-world can become the interface through which users learn what is askable, credible, normal, or actionable.

This cross-scale recurrence is neither literal cosmic equivalence nor formal fractality. It names correspondence under transformation: a pattern may recur across body, family, institution, platform, and machine while changing its material substrate, temporal rhythm, power relations, available agency, and structure of responsibility. The test is not whether the macro reproduces the micro, but whether a sheltering or corrective form at each scale forgets its exit conditions and begins to move before judgment can return.

This prepares, but does not yet equal, the problem named founder entropy. Room-formation can be centerless: a repeated phrase, practice, metric, fear, or reward can form a room without a single author. Founder entropy is more specific. It begins when a living act, a founder's name, or an originating insight becomes the room's credential for refusing renewed judgment. Room-formation asks how motion becomes an inhabitable world. Founder entropy asks how the name of the person or principle that opened the wall later becomes the wall.

## 12. Plato's cave and democratic threshold capacity

Plato's cave is not merely a story about ignorance. It is a story about room-captivity. The prisoners do not sit before isolated errors that can be corrected by a better fact. They inhabit a total environment. Their bodies are positioned, their field of vision is restricted, their shared language is formed around shadows, and their community of recognition is organized by what the room permits them to see. The cave works because it becomes ordinary.

This makes the cave structurally close to pre-action capture. The prisoners' judgments are furnished before they judge. Their reality is not simply false; it is socially reinforced, bodily habituated, and affectively defended. When one prisoner is released, the first experience is not triumph but pain, disorientation, and resistance. The old room is easier to inhabit than the difficult enlargement of the real.

The return to the cave is equally important. The one who has seen more does not automatically become fit to rule, persuade, or rescue. Returning to the room creates a second threshold: how to speak without domination, how to help without occupation, how to refuse the shadows without despising those whose lives were formed around them. A failed return produces either martyrdom, contempt, manipulation, or a new authority. A reconstructive return must preserve the possibility of exit without making the one who left into the owner of the exit.

Plato's political sequence in the *Republic* gives this problem a public form. The critique of democracy there should not be flattened into a simple rejection of ordinary people, nor converted into a modern anti-democratic slogan. Its threshold value lies elsewhere: a polity governed by unexamined appetite, rhetorical manipulation, and crowd-pressure becomes vulnerable to demagogic capture. Freedom without inward and institutional discipline can be recruited by the loudest appetite in the room.

The democratic danger is therefore not the people as such. The danger is a public room in which appetite speaks as liberty, flattery speaks as representation, resentment speaks as justice, spectacle speaks as truth, speed speaks as courage, and the demagogue becomes the mouth of the room. In that condition, procedure may remain while judgment is displaced upstream. The crowd experiences itself as choosing while the room has already trained the available choices.

This is where the framework's language of distributed threshold capacity becomes political. Democratic life requires more than the right to select among visible options. It requires citizens capable of pausing before the room speaks through them, institutions capable of delaying irreversible action, public language capable of distinguishing freedom from appetite, and media environments that do not convert every wound into a recruitable motion. A democracy without threshold capacity becomes available to completion pressure at mass scale.

At this scale, decision sovereignty names a secondary synthesis of the framework rather than a new pillar. It is not the unlimited power of an individual, office, or majority to command. It is the preservation of several capacities that no technical or institutional arrangement should invisibly occupy: the capacity to contest how the problem was framed; reopen evidence, uncertainty, and counterevidence; grant, limit, and revoke authority; interrupt or redirect execution; and correct records, harms, and institutions after effect. These may be described as frame, epistemic,

authorisation, execution, and reparative dimensions of sovereignty, but they remain moments of accountable motion rather than separate doctrines.

Their political legitimacy depends on more than a human being being present. Situated human and institutional presence must be joined to reason-giving, real contestability, domain adequacy, affected-party participation, acceptance of responsibility, and practical repair capacity. The intensity of these conditions should track materiality and reversibility: a low-impact, easily reversible act need not carry the same burden as an action affecting housing, liberty, livelihood, health, or public status. Human command without these conditions can be captured as readily as machine execution.

The democratic threshold is somatic as well as institutional. The citizen's finger hovers before the screen, or the hand weighs the stamp before the ballot paper. The clerk looks up from the form before treating the stranger as a category. The jury room holds one second of silence before verdict becomes record. The protester's voice waits before the slogan, asking whether this chant opens a room or closes one. These are not metaphors replacing politics with private feeling. They are the bodily sites where public judgment either lives or dies.

A serious objection must be faced. Arendt would warn that pre-action audit can reduce praxis to poesis. Action, in her account, is not manufacturing according to a prior model. It is plural, risky, irreversible, and meaningful only in a web of relations. If Reconstructive Threshold Philosophy treats action as something to be engineered before it begins, it may reproduce the old philosophical temptation to replace politics with administration.

The objection is valid against any version of the framework that treats democratic threshold capacity as a civic sensor, a procedural delay, or a general audit protocol. The answer is not to deny action's unpredictability. The answer is to place the threshold before action without pretending to master action. The interval does not manufacture the public act. It protects the possibility that the act will not be merely the room speaking through the citizen.

Arendt also reminds the framework that political repair cannot be reduced to better pre-action inspection. Promise and forgiveness matter because action is unpredictable and irreversible. They should not be absorbed into accountable commitment: accountable commitment releases pre-action audit into a responsible step, while promise and forgiveness govern the afterlife of action among others. The threshold can protect the beginning; it cannot guarantee the story. Democratic threshold capacity must therefore include practices after action: public reasons, answerability, revision, forgiveness where possible, institutional repair where required, and promises that bind future action without freezing it.

A second warning follows from Agamben's critique of apparatuses. Distributed threshold capacity can itself become a new judgment machine if proceduralized. A democracy that demands everyone constantly audit everyone else may become another room of suspicion. The democratic brake remains living only when it can release into common action and ordinary trust.

This does not mean that politics should be handed to a closed class of enlightened rulers. That would merely replace one room with another and call the walls wisdom. The threshold lesson is more demanding: democratic procedure needs democratic formation. Citizens must be educated not only to know facts, but to notice the room that organizes what facts are allowed to mean. Institutions must be designed not only to count preferences, but to slow the conversion of appetite, panic, humiliation, and spectacle into binding power.

The cave also clarifies why room-collapse is dangerous. Occupants of a room may be content so long as the walls hold. They may resist change because the room supplies identity, warmth, and a known sky. But when the room changes faster than the occupants can reconstruct, they are not automatically liberated. They may become disoriented, angry, nostalgic, or recruitable by anyone promising to rebuild the room without cost. The demagogue is often the builder of a replacement cave who speaks in the language of return.

Reconstructive Threshold Philosophy therefore reads Plato not as a call to despise the cave-bound, but as a warning that exit, return, and public judgment all require formation. The aim is not to shame those inside rooms. It is to build the capacity to notice rooms, leave them when necessary, repair them when possible, and refuse the leader who offers a more comfortable wall in place of an exit.

Arendt also clarifies what this framework cannot replace. Promise binds future action without pretending to master it; forgiveness repairs the irreversible without pretending it did not occur. Threshold work protects the beginning from being captured before judgment returns. Promise and forgiveness protect the middle and aftermath of action

once uncertainty, damage, and plurality have already entered the world. They are necessary complements, not symmetries. The center of gravity of this framework remains pre-action; it does not become a full Arendtian theory of action.

The Arendtian objection is not merely another objection to be answered. It is one of the framework's limit conditions. If threshold sensitivity begins to treat action as something that can be fully examined before it begins, the account has crossed into the very poesis it warns against. Action, in Arendt's sense, is not manufactured by better inspection. It appears among others, carries risk, and receives its meaning through a story no single actor controls. RTF therefore cannot claim to govern action as such. Its narrower task is to protect the interval in which capture may be noticed before motion becomes binding. Beyond that point, promise, forgiveness, narration, and public contestation are not supplements to the account. They are the political forms through which action continues after the threshold has been crossed.

A collective threshold is not the sum of private pauses. A public does not become answerable merely because many individuals hesitate at once. Collective threshold capacity belongs to the space between actors: procedures, stories, assemblies, institutions, protest, journalism, courts, memory, and public disagreement. The briefest individual pause can interrupt a gesture; it cannot by itself found a public world. Democratic threshold capacity must therefore be treated as a distinct political problem, not as the enlargement of an individual discipline.

Democratic threshold capacity must therefore be more than a civic sensor. It must remain connected to public action and public answerability. It is lived in bodies as well as institutions: the protester who swallows before repeating the slogan, the voter whose hand pauses before the stamp, the official who can still say in public why a rule was interrupted, and the citizen whose dissent can alter more than the tone of the room.

A citizens' jury, a municipal review board, or a regulatory hearing preserves threshold capacity only where dissent can damage the frame before the decision becomes record. Public participation that arrives after the binding point is not consultation; it is decoration. The public room remains living only when the person outside the prepared frame can still alter more than the vocabulary of inclusion.

In delegated and coupled systems, democratic threshold capacity begins before deployment as well as after harm. Affected publics need access to the reasons for deployment, the categories and objectives that will govern action, the authority to challenge permissions and defaults, notice when automated action has occurred, and routes to appeal, correction, and redress. A hidden stop button, nominal operator, or consultation without power is not democratic capacity if the public cannot alter the room in which the system is authorized to act.

### 13. Founder entropy

Founder entropy identifies the macro-scale test of whether a living founding act has hardened into system, ritual, identity, and eventually lock. It does not mean that founding is futile. It means that the founded is always exposed to entropy, and that a living principle survives not by being preserved unchanged but by being re-earned without being betrayed. A founder sees and moves. The institution that preserves the movement may later use the founder's name to prevent seeing and moving again.

This section builds on the distinction made in Section 11. Founder entropy identifies the narrower case in which a founding act, name, or originating insight becomes the room's credential for refusing renewed judgment.

As established in Section 3, the diagnosis that a founding force can harden into domination has strong predecessors in Weber, Michels, and Horkheimer and Adorno. Founder entropy applies that diagnosis to the specific moment when the founding name, founding act, or originating insight becomes an argument against renewed judgment. It is therefore the institutional counterpart of the living room test's Name as Substitute criterion: when the name ends the argument, the room has begun to defend itself against the principle that first gave it life.

Michels's account of oligarchic drift shows how organizations that begin with participatory or democratic claims can harden into leadership, apparatus, and routine. Founder entropy extends that concern without collapsing into it. The problem is not only that leadership concentrates power, or that charisma becomes administration. It is that a stabilizing form can begin to attack the very interval it was built to preserve. A founder's name, a school of thought, a procedure, a platform, or a protective doctrine may continue to speak the language of renewal while replacing

renewed judgment with authorized repetition. Oligarchic concentration is one form of this decay; founder entropy names the broader failure in which preservation becomes substitution.

Founder entropy is also the failure of the three axes described earlier. The temporal axis fails when fidelity to the past prevents the institution from seeing that today's preservation may become tomorrow's cage. The counterfactual axis fails when the existing room becomes the only imaginable room and alternative futures are treated as betrayal. The situated axis fails when the threshold is no longer lived in bodies, meetings, tools, or interfaces, but only in the founder's name, the archive, the ceremony, or the credential. When these failures converge, the living brake becomes a dead rule.

The problem is not limited to charismatic religious founders. It appears wherever a living principle crosses into social form: movements, schools, protocols, reforms, research programs, activist circles, ethical frameworks, technical standards, political institutions, and even anti-institutional communities. The founding act contains energy because it solved a real problem under pressure. But once transmitted, the act becomes a name. The name becomes a boundary. The boundary becomes custody. Custody becomes defense of the frame even when the problem has changed.

In machine-mediated institutions, the originating frame may sediment without a founder's name remaining visible. Objectives, training exclusions, procurement choices, default settings, permission structures, evaluation metrics, and API protocols can preserve an early decision after the people who made it no longer inspect it. Operational scale does not create a new form of entropy; it can make inherited assumptions faster, less visible, and harder to contest.

This is where founder entropy meets gradual institutional change. The founding act need not remain as a portrait or sacred name; it may persist as a rule that has been layered over, converted to a new purpose, or exhausted while its form survives. A protective procedure may still be cited while its threshold function has migrated to a vendor, workflow, or platform dependency. In that case the institution has not openly betrayed its origin. It has drifted until the origin no longer interrupts anything.

Founder entropy has several stages.

First, a living insight interrupts a prior capture. This may be a moral discovery, a technical safeguard, a political reform, a communal practice, or a symbolic language that helps people notice what had been invisible.

Second, the insight is stabilized for transmission. Stabilization is necessary. Without some form, the insight cannot travel. But every form selects, simplifies, and excludes.

Third, the stabilizing form becomes identity. People no longer merely use the principle; they belong to it. The principle now supplies meaning, status, loyalty, and protection from uncertainty.

Fourth, the identity becomes defensive. Questions about reconstruction are treated as attacks on the founder, betrayal of the movement, or dilution of the original truth.

Fifth, the living brake becomes a dead rule. The structure still uses the vocabulary of liberation, caution, or justice, but it now prevents the reconstruction that would keep the original insight alive.

Founder entropy is present where dissent no longer damages the frame, where revision refines language without changing conduct, where the founder's name substitutes for judgment, where successors can interpret but not exceed the founder, and where the institution cannot carry its own failure without closing the story around itself.

History offers a recurring pattern: a figure sees clearly, acts with genuine sacrifice, builds something real, carries the living principle into the world — and after the figure is gone, the principle freezes. The living brake becomes an institutional monument. The builder's name becomes the argument against questioning the building. The guardians of the legacy become its wardens. This pattern appears in religious teachings hardening into dogma, revolutionary movements turning liberation into bureaucracy, legal protections becoming procedural labyrinths, educational systems meant to cultivate critical thought producing vocabulary without self-examination, professional ethics becoming compliance forms, and anti-authoritarian movements developing informal priesthods. A further failure mode is half-formation: the vocabulary of critique can be acquired without the discipline that makes critique self-applicable. In that case, education does not preserve the threshold; it supplies better names for the room. This is the pedagogical form of threshold slip.

The failure also has a somatic and institutional surface: the hand that signed the charter becomes a portrait on the wall; the founding act survives, but now as a name that can end a conversation.

A founder's sentence can become a room long after the founder has left it. A meeting pauses before a decision, not because the argument is unclear, but because someone asks what the founder would have wanted. At that point the name has stopped preserving memory and has begun to replace judgment.

A reform movement may begin with a sentence that genuinely returns judgment: no decision should move before the absent party can be heard. For a time, the sentence interrupts power. Years later the same sentence is quoted at every meeting. It is printed on the wall, invoked in training, repeated before votes, and cited in response to criticism. Yet no decision changes when the absent party speaks. Dissent has not been silenced; it has been ceremonialized. The founder's sentence remains, but the interval it once opened has closed. This is founder entropy: the living brake preserved as language after it has ceased to interrupt motion.

The point is not that founding is futile. The point is that the founded is always exposed to entropy. A living principle does not survive by being preserved unchanged. It survives by being re-earned without being betrayed.

The same test applies to this work's versions, DOI records, canonical repository, release notes, and named vocabulary. Archiving can preserve inspectability and make revision answerable; it can also turn the latest release or the author's name into a reason not to reopen the argument. A DOI records a text. It does not certify that the text has escaped founder entropy. The archive remains living only where later readers can distinguish citation from obedience and release from final authority.

## 14. Post-threshold stewardship

If founder entropy is the failure mode the account cannot prevent, post-threshold stewardship is the practice that makes capture harder to complete and easier to notice. It begins after the frame has acted, but it must not become the actor's attempt to control the story of what the action meant.

If founder entropy is the failure mode, post-threshold stewardship is the partial answer. Stewardship is not guardianship in the conservative sense of preserving the principle unchanged. A guardian protects the frame. A steward protects the possibility that the frame may need to be rebuilt.

Stewardship is the social and institutional continuation of the living brake. It asks whether a principle that once protected judgment still protects it under present pressure, and whether the brake must release into action rather than keep demanding further examination. Reconstructive discipline supplies the steward's method: hold the frame long enough to work, destroy it when revision would merely refine captivity, rebuild it, and expose the rebuilt frame to consultation that can damage it.

Ordinary repair is the smallest test of this stewardship and the post-threshold continuation of the living brake's reconstructive feature (Section 5). It is not symbolic humility or a confession that changes nothing. It is a concrete alteration after motion: apology where harm has been named, correction where a category was wrong, a record amended rather than hidden, restitution where repair has material cost, changed conduct where language had become a substitute for judgment. Ordinary repair matters because it prevents the account from proving itself only at the level of theory. If the interval was preserved, the world should become more answerable after the act, not merely more eloquent about why the act occurred. Ordinary repair is not only post-action; it becomes pre-action material for the next threshold, because the repaired record, altered practice, or corrected relation changes what future motion can become.

A teacher who discovers that a grade was shaped by fatigue, stereotype, or inherited expectation does not need a ceremony of reconstruction. The ordinary repair is to reopen the work, correct the record, explain the change, and alter the practice that made the error easy. A clerk who realizes that an exception was forced into the wrong category does not complete stewardship by noting concern in the file. The file must be changed, the affected person must have a path of appeal, and the office must remember what the old category hid. These are small acts, but they are where post-threshold stewardship either touches the world or becomes identity.

In delegated and coupled systems, stewardship must follow the action across the chain. Traceability should reconstruct classifications, objectives, permissions, models, handoffs, and human decisions. Where reversal remains possible, rollback and record correction should be available. Where it does not, notice, explanation, appeal, compensation,

and institutional repair become necessary. Incident review must ask not only why a model erred, but why the arrangement was authorised to turn that error into consequence and what must change before the next threshold is reached.

Repair is incomplete if it changes only the affected record while leaving the objective, ontology, permission structure, interface, or workflow that produced the harm intact. Reparative answerability must be capable of altering the future frame: revising categories, narrowing authority, changing escalation and release conditions, withdrawing a system, or redistributing responsibility. It does not restore the lost interval retroactively; it rebuilds the capacity for a later interval to exist.

The same requirement applies before the visible interface. Couldry and Mejias's account of data colonialism is not imported as the essay's political economy, but it helps name a pre-interface form of capture: extraction, profiling, and dataset formation may organise future action before any prompt, score, or recommendation appears. If repair begins only after a model output has harmed someone, it arrives too late to inspect the room in which the data was made useful. Post-threshold stewardship must therefore ask whether the data relation itself - collection, labelling, proxy construction, retention, and economic incentive - has become the earlier binding point.

In predictive and screening systems, this means repair cannot stop at explanation of a single score. The future threshold is rebuilt only if categories, target variables, feature choices, thresholds, data flows, appeal routes, permissions, and institutional incentives can be revised. Affected parties should not be asked merely to correct a record after an automated future has acted on them; they require a route by which the future attributed to them can be contested before it becomes the next institutional default.

Threshold pedagogy belongs here in a limited sense. Threshold sensitivity is not taught by giving students, officers, designers, parents, or users a portable checklist to perform. It is taught by micro-cases in which the hand, category, screen, voice, and affected party become visible before motion binds. The first lesson may be the cold cup, the pause before the signature, the draft that made anger sound rational, or the child waiting beneath the adult's theory. The second lesson is that none of these scenes guarantees virtue. They only train the return from abstraction to the interval where judgment can still answer.

Post-threshold stewardship has at least five tasks.

First, it keeps the self-correction clause visible. A framework that warns against ossification must not hide that warning in an appendix where only the already converted will find it. The warning must remain close enough to the center that no one can use the center without encountering the possibility of revision.

Second, it separates loyalty from reconstruction. A community around a living principle must learn that questioning the frame can be fidelity to the principle, not betrayal of it. This is difficult because shared frames are load-bearing. To alter them can feel like exile.

Third, it distributes threshold capacity. If one person, teacher, founder, model, committee, or office remains the sole interpreter of the brake, the brake has already become a throne. A living threshold must be carried by many ordinary acts of attention, correction, and refusal.

Fourth, it preserves ordinary repair as the test. The more abstract a framework becomes, the easier it is to use as identity. Ordinary repair — apology, revision, care, documentation, restraint, accountable action — keeps the principle answerable to the world it claims to serve.

Fifth, it accepts partial failure. No structure can guarantee its own non-capture. A self-correction clause can itself become a credential. A guardrail can be displayed as proof of innocence while failing to prevent the motion it names. The point is not purity. The point is to make capture harder, more visible, and more interruptible.

Post-threshold stewardship therefore extends the living brake beyond the private interval. It asks how communities, tools, institutions, and texts can remain reconstructive after the first pressure that created them has passed.

Post-threshold stewardship must also include narration. An action does not receive its meaning only from the intention that preceded it or the audit that permitted it. It enters stories told by others: the harmed party, the witness, the institution, the public, the child, the archive. Stewardship therefore cannot mean controlling the story after the fact. It means remaining answerable to the plurality of stories through which the action becomes intelligible. Where the actor uses stewardship to fix the narrative in advance, post-threshold care has become another room.

## 15. Objections, limits, and the Ouroboros problem

Several objections should be taken seriously. To avoid turning self-audit into an exhausting list of alarms, they are grouped here as three rooms of objection: epistemic, political-institutional, and technological-machine. The grouping does not weaken the objections. It asks the reader to see which kind of capture each one names.

### 15.1 Epistemic objections

The first objection is that threshold language may dignify indecision. This is a real risk. A person can turn the interval into a sanctuary from consequence. This account answers with the counter-threshold: if delay no longer protects judgment, it has changed sides. The brake must be able to release.

The second is that pre-action capture may overstate the fragility of agency. If every impulse is treated as conditioned, responsibility may seem to dissolve. The argument does not remove responsibility; it relocates part of responsibility upstream. To be responsible is not only to answer for the act, but to inspect the conditions that make certain acts too available.

The third is that the account may aestheticize restraint. Any language of crowns, thresholds, gates, or brakes can become attractive enough to become a costume. This is why the Bound King is restricted in this essay to a bounded methodological discussion and treated as removable scaffold rather than doctrine. The test is not whether the language feels profound. The test is whether it improves judgment and returns agency to accountable motion.

The fourth is that the account may enter infinite regress. If every correction requires another correction, action never arrives. The response is the stopping rule developed in Section 5: accountable commitment, not endless inspection.

The fifth is that the account may become identity-thinking in Adorno's sense: a system that names non-capture so completely that the non-identical disappears. This account cannot escape identity-thinking by declaring itself exempt. It can only make its own capture harder to miss, keep remainder visible, and refuse to pretend that self-audit closes the wound it names.

The sixth is that the account may become one thought too many in Williams's sense. In intimate or urgent life, an extra audit may alienate action from the very relation that gives it meaning. The answer is not to inspect every gesture, but to keep the brake available where automatic motion is carrying harm.

Williams's worry is not merely that reflection can become excessive. It is that certain forms of life are damaged when the agent must pass through one more justificatory mirror before acting. A parent who pauses long enough not to wound the child has used a living brake; a parent who audits affection until the child becomes a case has lost the relation that made the action meaningful. Adorno names the danger of endless self-administration; Williams names the danger that the audit may estrange action from the love, loyalty, or urgency that gives it life.

This produces a further internal risk: threshold fatigue. A subject, office, or community may learn to audit so continuously that the brake becomes a twitch. The result is not judgment but reflexive delay, a performed suspicion of motion that protects the agent from accountable commitment. Threshold fatigue is not a reason to abandon the brake. It is a warning that the brake must remain answerable to release. The companion risk is somatic fade: after training the reader to notice the cold cup, the tired hand, and the breath before the verdict, the essay may retreat into abstractions until the body becomes decorative. The test for somatic fade is not how often the body is named, but whether the argument can return to the cold cup, the tired hand, or the waiting child after its longest abstraction.

### 15.2 Political and institutional objections

The seventh objection is that the account may be too individualistic. The private pause cannot replace law, protest, journalism, institutional design, labor organization, medical care, courts, democratic procedure, or technical governance. That objection is correct if the threshold is treated as sufficient. It is not. The interval is necessary but insufficient. It must be carried into institutional intervals: appeal, audit, review, documentation, accountability, second confirmation before irreversible action, and friction that protects the vulnerable rather than power.

The eighth is that consultation may become stakeholder ritual. Consultation that cannot damage the frame is not consultation. A public reason, rival account, affected testimony, expert correction, or external model must be able

to alter the frame. The same test applies to this essay's own revision history: feedback that could only polish the argument, but never change its direction, would be consultation capture rather than metabolism.

The ninth is that Plato's cave and democratic decline may be read as anti-democratic elitism. That reading should be resisted. The framework does not claim that ordinary people are unfit for political life, nor that a self-certified elite should rule on their behalf. Its claim is that democratic procedure requires distributed threshold capacity. Without citizens, institutions, and public media able to resist appetite, spectacle, humiliation, and demagogic framing, democratic forms can remain while democratic judgment is captured.

The tenth is the Arendtian objection, already treated in Section 12: the account may reduce action to pre-action audit. The response remains bounded. Threshold work protects the possibility of action's answerability; it does not manufacture action's meaning or replace promise, forgiveness, narration, and public contestation after action begins.

The eleventh is that room language may pathologize belonging. It should not. Families, schools, traditions, movements, disciplines, and institutions can shelter memory and make shared life possible. The living room test is meant to prevent blanket suspicion of all belonging. If its behavioral tests are passed, the framework must not call the room captured merely because it appears strict, traditional, or aesthetically dense.

The twelfth is that material conditions may be treated as frames. They must not be. As Sections 3.1 and 9 argue, hunger, debt, housing loss, legal danger, workload, platform ownership, and infrastructural control are not dissolved by better interpretation. The account fails if it cannot distinguish a material constraint from the frame through which that constraint is justified, naturalized, or made invisible.

### 15.3 Technological and machine objections

The thirteenth is that the account may imply machine consciousness or moral equivalence. It does not. Machine-facing threshold ethics concerns architecture, mediation, and consequence, not machine personhood.

The fourteenth is that machine-facing threshold ethics may overburden AI systems with responsibilities they cannot reliably perform. This objection is valid. Current systems may fail at frame audit, miss context, over-refuse, under-refuse, simulate caution, or perform reconstruction without changing the answer. For that reason, machine-facing threshold ethics cannot be reduced to model behavior. It also requires interface design, tool-use constraints, logging, escalation pathways, institutional accountability, and humans who do not outsource moral burden to the system that helped them speak.

A related objection concerns the contemporary language of reasoning systems, inference-time scaling, situational awareness, and deceptive or strategic behaviour. RTF should not turn these unsettled risk debates into a mythology of machine intention. The relevant threshold point is narrower: as systems spend more computation planning, checking, or decomposing a task before producing an answer, the user may receive a more coherent and action-ready frame while seeing less of the path by which alternatives were narrowed. More internal deliberation by the system does not guarantee more judgment for the human. It may preserve a threshold, or it may make occupation harder to feel.

The fifteenth is that Deliberate Refusals may become performative purity. Saying what the framework is not does not guarantee it will not become those things. Negative definitions draw boundaries; they do not replace vigilance.

The sixteenth is that the argument may launder its own guardrails. This objection is not external. It is the essay's internal risk. The presence of this objection does not absolve the text. It only makes the risk visible.

The seventeenth is that machine-facing threshold ethics may displace human-facing ethics. The danger is real if the interface becomes the place where humans outsource the difficulty of judgment. The response is to keep machine-facing ethics derivative: it is an operational stress test for the interval, not a replacement for human, institutional, political, or material responsibility. A system may slow, separate, surface, log, or escalate. It cannot own the action for those who must answer for it.

The eighteenth is that threshold delay may harm urgent action. This objection is also correct where the brake becomes a lock. Emergencies do not abolish threshold ethics; they change its form. In urgent conditions, the living brake should release into narrower, reversible, documented, and accountable action rather than total completion or

endless audit. The question is not whether to delay, but whether the next motion has been narrowed enough to remain answerable under time pressure.

A further objection concerns anthropomorphism. If RTF speaks of operational agency, machine thresholds, or action-selection intervals, does it merely smuggle personhood back into the system after denying consciousness? The response is to keep the predicates separate. Operational agency is a causal-organisational capacity defined by objectives, permissions, action selection, and environmental effect. It does not entail experience, moral standing, or conscience. The framework fails if structural comparison erases somatic and moral asymmetry.

A second objection concerns ceremonial oversight. A human may be formally present while lacking time, information, authority, or realistic capacity to reverse the path. The EU AI Act's oversight provisions and the literature on meaningful human control identify necessary institutional conditions, but RTF adds that the governing frame must itself remain contestable. Human presence cannot certify answerability where permissions, workloads, incentives, or downstream handoffs have already made intervention nominal.

A third objection concerns technical fetishism. Safe interruptibility, logging, rollback, monitoring, and stop mechanisms can appear to operationalise the living brake while narrowing the problem to control engineering. These devices are valuable but partial. They become guardrail laundering where their visibility increases trust without preserving reasoned interruption, release, redress, or re-entrant revision of the system's objectives and authority.

A fourth objection concerns responsibility gaps. Distributed systems can make every actor only partially causal, encouraging institutions to assign blame to the operator nearest the event. The moral-crumple-zone problem shows why distributed causation must be accompanied by differentiated but non-evadable responsibility. RTF does not provide a legal allocation rule. It requires that the post-threshold inquiry travel upstream far enough to reach the frame that made the consequence routine.

A fifth objection is that some machine-speed chains contain no live threshold at all. If permission is broad, handoff automatic, execution irreversible, and intervention merely ceremonial, it is misleading to describe the interior of the chain as an inspectable interval. The objection is correct in those cases. RTF should not invent a threshold where none remains. Its audit must move upstream to the authorisation and design that made the chain possible, and downstream to repair, redress, and the reconstruction of future threshold capacity.

#### **15.4 The Ouroboros problem: when the brake becomes a throne**

This argument warns against capture. But the warning itself can be captured. A reader may use threshold slip to dismiss any critique of the framework as a reception failure. A system may cite guardrail laundering as proof that it is not laundering guardrails. A community may invoke distributed threshold capacity as a new credential, a new room, a new name that ends argument. An author may cite AI disclosure as proof of responsible process while the disclosure itself functions as a credibility device.

Ouroboros is not an ornament added to the framework; it is the question whether the corrective mechanism has become the room it was built to interrupt. If the vocabulary converts a legitimate external objection into confirmation of its own insight, it has not answered the objection; it has performed the capture it names.

This is the Ouroboros problem: the argument eats its own tail. It does not escape this problem. It cannot. The most it can do is make the failure harder to miss. Therefore every claim in this essay — including the claim that every claim can be captured — remains exposed to the capture it names.

Threshold slip should therefore not be treated only as reception failure. It is also an internal risk. Nor is threshold slip an argument against simplification itself; a simplified form remains living if it preserves inspection, refusal, revision, and return to judgment. The essay may become a method. The keywords may become a room. The objections may become a guardrail performance. The distinction between Codex and essay may become a way to protect the essay from the force of the Codex. The warning against doctrine may become a doctrine of non-doctrine. The claim to assistance without occupation may become a way for the helper to self-certify innocence.

Nor should this paragraph become a shield. The framework should be revised or abandoned where it repeatedly delays answerable action, makes decisions less reversible or less contestable, converts critique into proof of its own insight, reduces material constraints to frame problems, or turns its vocabulary into a badge that replaces judgment. These are not empirical metrics. They are failure signs.

Self-implication becomes guardrail laundering when it produces warning without revision. It remains re-entrant audit only where the warning can alter the frame, the constraint, the allocation of responsibility, or the decision to release. Naming a failure condition is therefore not evidence that the condition has been met. The evidence is whether criticism changes the text's action path.

This is not a reason to abandon the framework in advance. It is a reason to keep its demolition charge close to its center. The reader who feels safe because they have understood this paragraph has already received a useful warning. The feeling of safety should not be trusted too quickly.

The recursive loop becomes especially dangerous where human and machine validation certify one another. A person brings a captured frame; a fluent system strengthens it; the person approves the completed package; the approval is recorded as human judgment; and the record becomes institutional proof that the process was controlled. The result then returns as precedent, data, or policy, making the same ontology harder to reopen. The loop is interrupted only where assumptions and alternatives can be exposed, evidence conflict can alter action, approval is not treated as sufficient proof of judgment, domain authority and affected-party testimony can damage the frame, permission can be revoked, and repair changes the future system rather than only explaining the past event.

A guardian system can become the same loop at a higher level. An AI oversight agent, safety classifier, audit dashboard, or risk score may warn continuously while lacking power to change authorisation, release, or repair. It may also increase institutional confidence precisely because it appears to supervise the first system. Such a guardian is threshold-bearing only where its warnings can alter action and where the guardian itself remains externally damageable: contestable by domain evidence, affected testimony, independent audit, and the possibility of decommissioning. Otherwise supervision has become a second room.

This is interruption without exemption. The framework does not solve the Ouroboros problem. It identifies conditions under which a recursive capture loop may be interrupted without claiming immunity from its return. Somatic fade is the same problem at the level of the body: the warning against abstraction can itself become abstract. A self-denying practice therefore refuses to turn self-recognition into a badge, membership token, or substitute for judgment. Every use of the vocabulary must preserve the possibility that its own naming has become a room.

The throne can first appear as a garment: something worn lightly enough to feel harmless before it hardens into authority. No declaration proves that the brake has been re-earned. The reader, affected party, external critic, or future revision must remain able to change the frame and its action path. If criticism can produce only another warning, the tail has already reached the mouth.

## 15.5 Deliberate Refusals

This is not self-help. It does not provide a method for living, a set of habits, a productivity system, or a therapeutic protocol.

This is not a spiritual marketplace. It offers no private ladder of ascent, no special identity for the awakened, and no sacred vocabulary through which exhaustion can be renamed virtue.

This is not doctrine. Its mythic figures are structural metaphors and literary carriers, not literal cosmological claims. The essay includes no belief section, no initiation sequence, no secret knowledge, and no claim to spiritual authority.

This is not promptware. The machine-facing sections are not reusable system prompts, agent personas, or alignment templates. Constraint language, if used elsewhere, must be translated into ordinary operational terms rather than imported as ritualized persona.

This is not a complete AI-safety architecture or engineering standard. It does not propose a training objective, reward signal, universal control stack, benchmark, executable protocol, or compliance checklist. It does analyse the philosophical conditions under which assistive, delegated, autonomous, and coupled machine action remains inspectable, interruptible, contestable, and answerable.

This is not a political program. It names a structural failure that political economy, institutional design, civic action, and technical governance may each need to address. It does not replace those domains.

This is not a method. The dynamic frame cycle is not a checklist. The living brake is not a rule. Threshold sensitivity is not a five-step program.

The refusals are structural. Any threshold text that becomes a method, doctrine, persona, product, credential, or identity has already become the thing it warns against: another command that executes before judgment arrives.

Positively stated, this essay is a regulative philosophical framework and vocabulary for the interval before action. It does not decide what should be done. It asks whether judgment has had a chance to return before action becomes binding.

## 16. Conclusion: motion with the interval intact

The central claim of this essay is small and difficult: action should not be allowed to begin before judgment has had a chance to return, and the structures that protect that return must themselves remain open to reconstruction. This applies to the individual action, the institutional signature, the platform interface, the machine answer, the tool call, the teacher's authority, the founder's name, the public room, and the text that warns against all of them.

The essay's differentia is not the discovery that people, institutions, or liberating movements can be captured. That has been named many times. Its narrower claim is that a threshold discipline must become re-entrant, cross-scale, and machine-speed aware: it must inspect the tools of correction, the forms of help, the founders of rooms, the guardrails of machines, and its own language before they harden into new sites of obedience.

The framework was written because the interval before action is no longer only a private moral chamber. It is now mediated, accelerated, beautified, formatted, scaled, and sometimes executed by machines; and rooms now scale beyond the local forms that once contained them. A model may render a frame already under capture plausible before the user has examined it. A tool-using system can turn panic into paperwork, rage into a sent message, suspicion into a ranked list, or institutional ontology into binding action. A frame that once lived inside a family, sect, bureaucracy, party, professional culture, or institution can now be translated into interfaces, retrieval systems, policies, rankings, recommendations, and autonomous tool-use. The threshold must therefore be preserved not only inside the person, but at the interface where language begins to move the world.

The machine-facing extension in this version distinguishes assistance from delegation. A recommendation may return to judgment; a delegated system may continue through permissions, tool calls, and handoffs before fresh judgment returns; a coupled system may distribute binding across several components and institutions. None of these systems becomes a moral person. Yet each can carry an inherited frame into consequence. The ethical task is therefore to preserve inspectability, meaningful interruption, differentiated responsibility, and post-threshold repair across the whole action chain.

Decision sovereignty names the result of that preservation, not an additional foundation. It is neither machine autonomy nor absolute human veto. It is the continued availability of reason-giving, contestability, domain-adequate review, affected-party participation, bounded authority, consequence ownership, and repair as the action moves from frame to authorisation, execution, and effect. Human presence without those conditions can be ceremonial; technical traceability without them can be governance theatre. Accountable motion requires both lived answerability and architected reversibility, and reparative answerability when neither was enough.

The release extension makes the same claim more externally answerable. Dromology, preemption, algorithmic governmentality, institutional drift, and human-oversight critiques do not replace RTF's vocabulary. They show where that vocabulary must be damaged if it becomes too self-contained. The shared lesson is that decision sovereignty may be lost before anyone sees a decision: through speed, prediction, context engineering, functional drift, or oversight that records approval after the action path has already closed.

Pre-action capture marks the danger. The living brake marks the necessary interruption. Threshold sensitivity describes the trained perception of the interval: temporal, counterfactual, and situated. Reconstructive discipline describes the cycle by which frames are built, used, tested, destroyed where necessary, rebuilt, and exposed to consultation. Accountable commitment names the decision required when uncertainty remains. Assistance without occupation names the ethical form of help that returns agency rather than keeping it. Machine-facing threshold ethics names the operational discipline required where models classify, smooth, retrieve, recommend, receive delegated authority, call tools, pass outputs to other systems, and act. Room-formation names the way repeated motion becomes an inhabitable world. Democratic threshold capacity names the public formation needed to keep freedom from collapsing into appetite and demagoguery. Founder entropy identifies the tendency of living principles to become

monuments. Post-threshold stewardship names the work of keeping them re-earnable. The Ouroboros problem names the fact that all these names can themselves become rooms.

The three axes transform threshold philosophy from a general metaphor of transition into a trained perception of the specific point where judgment can still return. The temporal axis asks what a present motion may become; the counterfactual axis asks which possible room can still be strengthened before the interval closes; the situated axis asks where the threshold is being lived and who will bear its cost. Without these axes, threshold language risks becoming a general metaphor for transition rather than a disciplined perception of the point where judgment, refusal, repair, or accountable action remains possible.

The Bound King is useful only if the argument can release him. The same is true of every method, teacher, model, doctrine, safeguard, room, institution, and text. The point is not to preserve the figure who represents the threshold. The point is to notice when any figure begins to occupy the place where judgment should return. Threshold slip is the same danger at the level of reception: the moment a living warning is made so portable, smooth, or useful that it becomes a room of its own.

Ordinary repair remains the final anti-grandiosity test of the framework: the cup can be washed without becoming proof of a crown; the child can be answered without the old sharpness; the message can wait if waiting preserves judgment, and can be sent if delay has become another room.

This account succeeds only where it can release the reader into action, narration, repair, and disagreement without requiring the reader to remain inside its vocabulary.

The final demand is therefore not stillness, but motion with the interval intact: pause without evasion, help without occupation, reconstruction without paralysis, institution without captivity, and action without surrendering the conditions of judgment. This essay may still be laundering its own guardrails. That possibility is not the end of the argument. It is the reason the argument must remain answerable after the page ends. The page ends. The room remains.

## Appendix A — Scholarly Delta Sequences for Section 3.1

Appendix A is part of the argument, not an external supplement. It restores the full scholarly delta load that the main-body version of Section 3.1 now carries only in compressed form. The main reading path keeps the worked triage example short; this appendix preserves the inspectable comparison material so that the claim to distinctiveness does not rest on compression, tone, or terminology alone.

### A.0. How the delta test is to be read

The comparison in this appendix uses a four-part template for each predecessor: its core commitment, where that commitment would tend to locate the triage problem, what intervention it would most naturally suggest, and what concrete binding-point difference remains. The claim is not that these predecessors are blind, obsolete, or unable to be extended. Many of them are strong enough to approach the problem. The narrower claim is that none is organized, from within its own primary commitment, to force the same combined question at the same compressed point: has a fairness-preserving corrective become a captor in the interval where prompt becomes recommendation and recommendation becomes signable action?

To prevent this comparison from becoming a loose list of names, the comparison uses the following explicit analytic delta template. The notation is included here because the reader should be able to see how the scholarly comparison is tested inside the essay itself:

[C] Core commitment. The primary theoretical commitment that organises the predecessor framework.

[M] Likely modification or repair. The intervention that framework would most naturally propose inside the AI-assisted triage room.

[I] RTF intervention point. The point at which RTF relocates the problem: the last reversible interval where prompt becomes recommendation and recommendation becomes signable action.

Conclusion / delta. Whether the predecessor is forced, by its own primary commitment, to converge on the same re-entrant, cross-scale, machine-speed binding-point question.

The template is not a scoring system and it is not an administrative protocol. It is an audit trail for the argument. It keeps the comparison from collapsing into either vague originality language or unfair claims that earlier frameworks were blind. Each predecessor is granted its strongest relevant commitment first; only then is the RTF delta named. The prose comparison gives the reader argumentative continuity; the explicit [C] -> [M] -> [I] -> Conclusion / delta sequences preserve the diagnostic architecture of the delta test. Both are retained here because the essay is meant to be read both as argument and as an inspectable threshold analysis.

The strongest rival is therefore not any single predecessor, but a partial conjunction of predecessors. Second-order cybernetics, especially von Foerster, supplies the re-entrant turn: the observer and controller must be brought back into the observed system. Beer's viable-system tradition supplies recursive regulation across organisational levels: System 1-5 recursion, exception channels, feedback, escalation, algedonic signalling, and Ashby's law of requisite variety as incorporated into the cybernetic tradition. Taken together, von Foerster and Beer provide roughly two thirds of the present framework's conjunction: re-entrant observation and multi-scale recursion. The point is not that they lack speed, feedback, or organisational intelligence. A Beerian system can be highly attentive to real-time regulation, and von Foerster's ethical imperative to increase the number of choices already presses against closure. What they do not by themselves force is the machine-speed binding-point question: whether the corrective mechanism has become the captor inside the last reversible interval where prompt becomes recommendation and recommendation becomes signable action. O'Neil and the automation-accountability literature supply the third pressure from another direction: opacity, scale, destructive scoring, automation bias, over-reliance, weak human oversight, and the administrative damage produced by apparently objective categories. But that literature does not by itself require the re-entrant test of the safeguard as captor. The triage example shows why the missing third matters. If re-entry and recursion are present without machine-speed binding analysis, the system may improve feedback, exception routing, and escalation while leaving the signable denial intact. If algorithmic-harm analysis is present without re-entrant audit, the system may add transparency, contestability, cognitive forcing, or human oversight while leaving caution language to legitimate the same action. RTF's delta is the forced conjunction: re-entrant audit, cross-scale institutional diagnosis, and machine-speed attention to the binding point at which a fairness-preserving corrective becomes the path by which the room closes.

#### **A.0.1. The competitor pass and the missing third problem**

The competitor pass gives a more exact form to the essay's differentia. Weber, Foucault, Bourdieu, Dewey, Adorno, Arendt, and O'Neil each illuminate a real part of the room. Von Foerster and Beer are the closest structural predecessors because they bring re-entry and recursion into view. But a two-thirds rival still misses something. Von Foerster gives the re-entrant turn; Beer gives recursive organisational regulation; O'Neil and the automation-bias literature identify machine-mediated damage and over-reliance. RTF's narrower surplus is to hold all three demands at once: re-entrant audit, cross-scale institutional diagnosis, and machine-speed attention to the last reversible interval. Without that conjunction, the triage room's binding point can be mislocated: either in bureaucracy in general, classification in general, habit in general, inquiry failure in general, system design in general, administered identity in general, public action in general, or algorithmic harm in general. RTF asks instead whether the corrective mechanism has become the captor at the moment action is about to bind.

Each of these accounts matters. None should be displaced. But none by itself is forced to ask the full threshold question at the binding point. The re-entrant question is: did the safeguards designed to protect fairness become the very instruments through which judgment was displaced? The cross-scale question is: how did family need, caseworker fatigue, bureaucratic rule, platform interface, risk ontology, managerial pressure, and machine fluency converge before anyone experienced themselves as choosing? The machine-speed question is: did the model's fluent caution make the frame more inspectable, or did it polish the institution's existing ontology into a decision before judgment could return?

The framework's specific judgment is not simply "the model was biased," "the bureaucracy was alienating," or "the worker failed to reflect." It is this: the room became binding at the interval where the corrective mechanisms should have remained interruptible. The model's caution is not automatically a brake. In this example, uncertainty is noted but the recommended action does not change. This is the worked form of the guardrail laundering defined later in Section 9: caution language performs fairness while making the existing institutional motion sound careful, neutral, and ready to sign. A self-capture failure appears inside the workflow itself: the protective language becomes one of the means by which harm receives procedural legitimacy.

### A.0.2. Material boundary of the worked example

The worked example now includes both an assistive variant, in which a recommendation returns to a reviewer, and a delegated variant, in which pre-authorised classification or handoff can initiate downstream action before fresh human review. The second variant does not turn the appendix into a technical protocol. It tests whether the same predecessor accounts still locate the binding point once authority and action are distributed across a chain.

This example also clarifies a material boundary. Machine-mediated rooms are not made only of concepts. They are owned, funded, procured, optimized, maintained, and rewarded. In a material and political-economy sense, the analysis cannot stop at the interface while leaving ownership and incentive structures untouched. A platform can preserve the interval in language while its business model, institutional contract, or performance metric rewards speed, volume, denial, ranking, or engagement. Hunger is not a frame. Debt is not merely an interpretation. A locked door is not a perspective. But hunger, debt, and locked doors are lived through frames that can normalize them, disguise them, or make certain responses to them appear inevitable. A threshold analysis that ignores ownership, incentives, and infrastructure control risks becoming ethics at the interface while political economy continues underneath it. This essay does not become a theory of political economy; it marks that boundary as one of the rooms through which the framework must pass.

The triage room is therefore not created by the model alone. It is prepared before the prompt by budget scarcity, housing markets, eligibility law, procurement contracts, managerial workload, risk allocation, migration status, debt, rent, political pressure, and the administrative need to make scarcity appear procedural. The dashboard does not invent these pressures. It can, however, translate them into a smoother decision path. A denial may look like a model recommendation while carrying the prior weight of a market, a policy, a budget line, a landlord's leverage, or a welfare office trained to distribute insufficiency as if insufficiency were neutral.

This matters because guardrail laundering can begin before the visible guardrail appears. If the surrounding institution rewards speed, caseload reduction, risk avoidance, or denial, then a fairness note attached to a recommendation may only decorate a path already made cheaper by material design. The binding point is not only inside the screen. It is also in the procurement specification that made reversal costly, the metric that counts closure faster than repair, the legal category that made the family legible only as risk, and the labor condition that gives the caseworker no time to see the absent child.

The framework can name this boundary, but it cannot replace political economy, housing policy, labor analysis, migration law, or public finance. Its narrower use is to prevent those material structures from being hidden behind the polite surface of interface ethics. A threshold analysis fails where it treats poverty, race, immigration status, debt, disability, family breakdown, or administrative scarcity as mere frames. It succeeds only where it asks how material constraints and interpretive frames reinforce one another before a human or machine output becomes signable action.

### A.1. Predecessor-by-predecessor explicit delta sequences

The sequences below replace the compact matrix form with the fuller long-form delta format restored from the earlier baseline. Each predecessor is read first in its strongest local commitment, then through the same four-part sequence: **[C] Core commitment**, **[M] Likely modification or repair**, **[I] RTF intervention point**, and **Conclusion / delta**. The repetition is intentional. It keeps the appendix from becoming a compressed name-ledger and makes the differentia inspectable at the level of each comparison rather than only at the level of the whole table.

**A.1. Weber — explicit delta sequence** Weber is strongest where the triage case becomes office logic: family, file, rule, competence, record, and signature. The danger is not only coldness or hierarchy; it is the authority of a procedure that can make a decision appear valid because the chain has been properly followed. RTF preserves that diagnosis but moves the question one notch inward. The issue is whether the fairness mechanism inside the rationalized procedure became one of the instruments through which judgment was displaced. The binding point is the local conversion of a caution note into procedural confidence at the moment of signature.

**[C] Core commitment:** Bureaucratic rationalization, formal rationality, and the iron cage of calculability, predictability, hierarchy, record, competence, office, rule, and control.

**[M] Likely modification or repair:** At the triage binding point, this commitment directs attention to the office, the rule structure, the documented procedure, the routinized administrative chain, jurisdiction, responsibility,

discretion, documentation, competence, and the relation between formal and substantive rationality. A Weberian repair would tend to sharpen rule, office responsibility, accountability, and the chain through which a file becomes an administratively valid decision.

**[I] RTF intervention point:** RTF relocates the question to the moment when the rationalized procedure's own fairness-preserving safeguard becomes part of the mechanism that displaces judgment. The issue is not only whether bureaucracy has rationalized the decision, but whether a caution note inside that rationalized workflow gives denial procedural confidence at the moment of signature.

**Conclusion / delta:** Weber explains the bureaucratic cage and the authority of procedure. RTF asks where the cage becomes locally operable as signable action: the fairness-preserving corrective itself may become the instrument through which the room closes.

**A.2. Foucault — explicit delta sequence** A Foucauldian account sees power/knowledge, disciplinary classification, biopolitical normalization, examination, and subject-formation. In the triage room it would locate the binding problem in the production of the applicant as a risk object: the file, dashboard, category, and administrative gaze make the family legible by narrowing what kind of subject it can be. Its natural repair would be genealogical and critical: expose the regime of classification, the normalizing examination, the governmental rationality of the dashboard, and the ways counter-conduct might resist them. This also matters. RTF does not replace it. But RTF asks whether the protective supplement to classification — the model's balanced uncertainty language — has itself become a capture device. The binding point is not only the construction of the applicant as risk; it is the moment when cautious language makes that construction easier to sign.

**[C] Core commitment:** Power/knowledge, disciplinary power, biopolitical normalization, examination, surveillance, governmental rationality, and subject-production through classification.

**[M] Likely modification or repair:** At the triage binding point, this commitment directs attention to the applicant's construction as a risk object: category, dashboard, file, examination, administrative gaze, and the governmental rationality by which a family becomes legible as a case. A Foucauldian repair would be genealogical and critical: expose the regime of classification, the normalizing examination, the governmental rationality of the dashboard, and the possible practices of counter-conduct that resist them.

**[I] RTF intervention point:** RTF does not deny the risk ontology. It asks whether the protective supplement to classification — balanced phrasing, uncertainty notation, fairness language, or caution — has itself become a capture device. The issue is not only the production of the applicant as risk, but the conversion of that production into a decision that appears careful enough to sign.

**Conclusion / delta:** Foucault reveals the risk ontology and the disciplinary file. RTF adds the binding-point question: did the uncertainty note make the risk ontology easier to sign before judgment could return?

**A.3. Bourdieu — explicit delta sequence** A Bourdieusian account sees habitus, field, doxa, symbolic violence, and the embodied feel for the game. In the triage room it would locate the problem in the trained dispositions of the caseworker and office: under fatigue, scarcity, and managerial pressure, one administrative movement feels natural while other movements feel costly or implausible. Its natural repair would examine field position, institutional doxa, symbolic power, and reflexive sociology's demand that the analyst inspect the conditions of analysis. RTF accepts this but relocates part of the question to the machine-mediated interval. The caseworker's habitus is not only socially trained; it is now coupled to an interface whose fluent caution can make the old movement feel newly justified. The binding point is the joined motion of disposition, dashboard, model output, and signable recommendation before the present case has spoken back.

**[C] Core commitment:** Habitus, field, doxa, symbolic violence, practical sense, and the trained feel for the administrative game.

**[M] Likely modification or repair:** At the triage binding point, this commitment directs attention to the caseworker's embodied dispositions and the office's doxa: under fatigue, scarcity, managerial pressure, and institutional scarcity, one administrative move feels natural while other moves feel costly, implausible, or professionally risky. A Bourdieusian repair would examine field position, institutional doxa, symbolic power, and the analyst's own conditions of analysis.

[I] **RTF intervention point:** RTF places habitus inside a machine-mediated interval. The caseworker's practical sense is not only socially trained; it is coupled to a dashboard and a fluent model output whose balanced caution can make the old administrative movement feel newly justified.

**Conclusion / delta:** Bourdieu explains why the caseworker trusts the administrative movement. RTF shows how disposition, dashboard, model output, and signable recommendation join before the present case can interrupt them.

**A.4. Dewey — explicit delta sequence** A Deweyan account sees an indeterminate situation that requires inquiry, reflection, experimentation, and reconstruction. In the triage room it would locate the failure in the absence or closure of reflective inquiry: the problem was formulated too quickly, relevant experiences were excluded, and the office moved from doubt to conclusion before inquiry had done its work. Its natural repair would reopen inquiry, widen experience, and reconstruct the situation through active, continuous, and careful thought. RTF is close to this pragmatist lineage, but it adds a re-entrant and machine-facing test. The question is not only why reflection failed; it is whether a method or tool that appears to support reflection has removed the need to feel reflection's cost. The model's fluent caution can imitate inquiry while authorizing closure. The binding point is where reflective language becomes a substitute for renewed inquiry.

[C] **Core commitment:** Reflective inquiry, experimental reconstruction, experience, and the transformation of indeterminate situations into more adequate problem formations.

[M] **Likely modification or repair:** At the triage binding point, this commitment directs attention to the failure of inquiry: the office moved too quickly from an uncertain case to a determinate recommendation, relevant experience was excluded, and the problem was formulated too narrowly. A Deweyan repair would reopen inquiry, broaden experience, test consequences, and reconstruct the problem situation before closure.

[I] **RTF intervention point:** RTF asks whether a tool that appears to support inquiry has removed the felt cost of inquiry. The model's careful language may imitate reflective inquiry while authorizing closure. The issue is not only that inquiry failed, but that inquiry-like language became the substitute for inquiry.

**Conclusion / delta:** Dewey diagnoses premature closure of reflection. RTF names a more compressed machine-facing failure: fluent reflective language can close the room while appearing to keep it open.

**A.5. von Foerster — explicit delta sequence** Second-order cybernetics, especially in von Foerster's sense, sees the observer within the observed system. In the triage room it would locate the problem in the observer-controller's own distinctions: the office and its model are not outside the system they regulate, and responsibility includes responsibility for the distinctions by which the system observes. Its natural repair would increase observer responsibility, make observation observable, and preserve or increase the space of possible actions. This is one of the closest predecessors because it supplies a re-entrant turn. Yet RTF places that turn at a more specific and more compressed binding point. The question is not only whether the observer observes its observing, but whether the system's own safety distinction has made one action easier while appearing to preserve options. The binding point is the prompt-to-recommendation interval in which an uncertainty note can become a one-way path toward denial.

[C] **Core commitment:** Second-order observation: the observer must be included in the observed system and made responsible for the distinctions it draws.

[M] **Likely modification or repair:** At the triage binding point, this commitment directs attention to the observer-controller's own distinctions and to the system's self-observation. The office, model, and supervisor are not outside the system they regulate. A von Foersterian repair would make observation observable, increase responsibility for distinctions, and preserve or increase the space of possible actions.

[I] **RTF intervention point:** This is one of the closest predecessors because it supplies the re-entrant turn. The remaining delta is not that von Foerster lacks re-entry. It is that RTF places the re-entrant question at the technical chokepoint where model output becomes signable institutional action: has the system's own safety distinction made one action easier while appearing to preserve alternatives?

**Conclusion / delta:** von Foerster brings the observer back into the system. RTF asks whether an uncertainty note has made denial the path of least resistance inside the prompt -> recommendation interval, including where

observation and intervention are distributed across permissions, services, and downstream handoffs rather than concentrated in one reviewer.

**A.6. Beer — explicit delta sequence** Beerian cybernetics, especially the viable-system tradition, sees recursive regulation, requisite variety, algedonic signals, and coordination across organisational levels. In the triage room it would locate the problem in system design: the office may lack sufficient variety, effective exception channels, recursive feedback, or escalation pathways able to register distress. Its natural repair would redesign the organisational system so that exceptions, overload, and error signals can move through the right recursive levels. RTF again treats this as a near predecessor rather than an opponent. But Beerian viability remains primarily functional and regulatory; it does not by itself require the living/captured-room test. A system can be recursively regulated and still allow its fairness guardrail to become a legitimacy device. The binding point is not only whether the organisation has feedback, but whether the corrective feedback itself has become the smooth path by which the room closes.

[C] **Core commitment:** Viable-system recursion, requisite variety, algedonic signalling, feedback, exception channels, and multi-level organisational regulation.

[M] **Likely modification or repair:** At the triage binding point, this commitment directs attention to system design: exception channels, recursive feedback, escalation routes, operational overload, requisite variety, and the organisational capacity to register distress. A Beerian repair would redesign recursive feedback, improve exception routing, increase system variety, and make the organisation more viable under stress.

[I] **RTF intervention point:** Beer is another near predecessor. The delta is not that Beer lacks speed, feedback, or recursion. A viable-system tradition can be highly attentive to real-time regulation. The remaining question is normative and threshold-specific: does recursive regulation include a living/captured-room test for whether the regulator or corrective mechanism itself has become the means by which judgment is displaced?

At the triage binding point, algedonic signalling offers a strong partial analogue: it can surface urgent distress that routine channels suppress. It is not equivalent to the living brake, because the latter must also ask whether the signal, regulator, safeguard, or identity function has become a path of capture, and whether continued interruption should now release into accountable action.

**Conclusion / delta:** A qualification is necessary. Beer's System 5, the policy and identity function of the viable system, does include a form of second-order observation: the system reflects on its own identity, purpose, and regulatory pattern. In this restricted sense, VSM is not without re-entrant structure. The delta is therefore not that Beer lacks recursion, feedback, or system self-reference. It is that his recursion remains functional and organisational rather than normatively diagnostic in the specific sense required by the living/captured-room test. System 5 can ask whether the system is viable and coherent; it is not organised, by that question alone, to ask whether the safeguard that makes viability legible has become a throne. The re-entrant turn in Beer preserves system identity; the re-entrant turn in RTF asks whether identity-preservation has become capture. This is why the two traditions are near predecessors rather than opponents, and why their conjunction still requires the third pressure that machine-speed binding analysis supplies. In delegated systems, the same delta asks whether recursive control remains capable of inspecting permissions and handoffs distributed across components, or whether viability has been preserved by making the path of capture harder for any one control centre to see.

**A.7. Adorno — explicit delta sequence** An Adornian account sees identity-thinking, administered life, and the violence by which a non-identical person is made identical with a category. In the triage room it would locate the problem in the whole administered world that turns family need into file, score, and case. Its natural repair would be immanent critique: resist the totalizing concept, keep remainder visible, and refuse reconciliation that merely prettifies domination. RTF takes that warning seriously, including against itself. Its narrower contribution is not to overcome identity-thinking, which it cannot do by declaration. It asks for operational signs of identity-thinking inside the very devices meant to soften it. The binding point is the moment when the model's careful language makes the identifying category feel less violent without making it less binding.

[C] **Core commitment:** Negative dialectics, identity-thinking, the administered world, and the violence by which the non-identical is made identical with a category.

[M] **Likely modification or repair:** At the triage binding point, this commitment directs attention to the whole

administered order through which family need becomes case, file, category, score, and denial. An Adornian repair would proceed by immanent critique, resistance to false reconciliation, and refusal to let the remainder disappear inside the concept.

**[I] RTF intervention point:** RTF accepts this warning, including against itself. Its narrower task is not to overcome identity-thinking by declaration. It asks for operational signs of identity-thinking inside the devices that claim to soften it: does careful language make an identifying category feel less violent without making it less binding?

**Conclusion / delta:** Adorno sees the administered world and the violence of identification. RTF asks where careful language makes an identifying category easier to accept as decision.

**A.8. Arendt — explicit delta sequence** Arendt remains an objection before she becomes a comparison. Her concern is action, natality, plurality, promise, forgiveness, and the danger that action will be reduced to manufacture or administration. In the triage case she would warn that pre-action audit can itself become an administrative technique for preventing action from appearing. RTF should not pretend to solve that objection. Its claim is bounded: before public action or retrospective judgment can begin, machine-mediated conditions may already narrow the possibility of answerable action. The binding point is not the full meaning of the action in public; it is the pre-action loss of judgment through which an administrative act becomes available as if no alternative room were still alive.

**[C] Core commitment:** Action as praxis, natality, plurality, promise, forgiveness, unpredictability, irreversibility, and the danger of reducing action to manufacture or administration.

**[M] Likely modification or repair:** At the triage binding point, Arendt warns that the scene may not be public action at all. It may be administration, necessity, and the reduction of human plurality to procedure. Her repair would protect public appearance, plurality, promise, forgiveness, and the non-manufacturable character of action; her broader warning in *The Origins of Totalitarianism* about systems that replace plurality with ideological or administrative necessity remains relevant as a background pressure, not as the immediate mechanism of the triage case.

**[I] RTF intervention point:** RTF should not absorb Arendt as if the objection were solved. Arendt remains a standing objection: pre-action audit can reduce praxis to poiesis. RTF's narrower claim is that before public action or retrospective judgment can occur, machine-mediated conditions may already narrow the possibility of answerable action. Promise and forgiveness are complements, not replacements, for the pre-action threshold.

**Conclusion / delta:** Arendt reminds RTF that threshold work is necessary but insufficient. RTF focuses on the pre-action loss through which answerable action may already have been narrowed before the public story begins.

**A.9. O'Neil / automation-accountability — explicit delta sequence** An O'Neilian and automation-accountability account sees opacity, scale, damage, feedback loops, automation bias, over-reliance, and the weakness of purely formal human oversight. In the triage room it would locate the problem in destructive scoring, opaque categories, inadequate accountability, and the human tendency to rely on apparently objective machine output. Its natural repair would require audit, transparency, contestability, documentation, and stronger oversight. This is essential to the present example. But RTF adds the guardrail-laundering question. The danger is not only that the model is opaque, biased, or over-trusted. It is that an apparently responsible safeguard — an uncertainty note, fairness warning, or cautious phrasing — may increase trust without changing the authorized action. The binding point is where oversight language itself becomes a path to over-reliance.

**[C] Core commitment:** Opacity, scale, damage, destructive scoring, feedback loops, automation bias, over-reliance, and weak human oversight.

**[M] Likely modification or repair:** At the triage binding point, this commitment directs attention to destructive scoring, opaque categories, inadequate accountability, human over-reliance on apparently objective machine output, and the weakness of purely formal oversight. Its likely repair would demand audit, transparency, contestability, documentation, and stronger safeguards against automation bias. The automation-accountability literature has refined this repair in important directions. Buçinca et al. (2021) propose cognitive forcing functions: deliberately designed friction that interrupts over-reliance by making the user think before accepting a model output. Green and Chen (2019) show that algorithm-in-the-loop structures can produce disparate interactions even when the algorithm

is formally fair. Green (2022) argues that policies requiring human oversight of government algorithms are often flawed because the oversight is procedural rather than substantive: the human is present but not empowered to alter the frame. De-Arteaga, Fogliato, and Chouldechova (2020) and Saxena et al. (2021) likewise help show that allocation, ranking, and decision-support systems can produce harms not only through error, but through how human users interact with model outputs under institutional pressure. These refinements strengthen the O’Neilian repair, yet they also reveal its limit. Cognitive forcing functions, algorithm-in-the-loop transparency, and procedural oversight are all safeguards. The question RTF adds is whether the safeguard itself — the forcing function, the transparency layer, the oversight formality, the caution note, or the balanced draft — has become a new path to over-reliance. A cognitive forcing function that the user learns to click through becomes a ritual. An algorithm-in-the-loop display that the caseworker glances at while reaching for the signature becomes theatre. A human-oversight policy that documents review without requiring reversal becomes a guardrail that launders the room it was meant to open. The binding point is not whether the safeguard was mentioned, but whether it changed the frame at the interval where recommendation becomes signature.


[I] **RTF intervention point:** RTF accepts these repairs but adds the guardrail-laundering question. The danger is not only that the model is biased, opaque, or over-trusted. It is that an apparently responsible safeguard — uncertainty note, fairness warning, cautious phrasing, balanced tone, or human-in-the-loop formality — may increase trust without changing the authorized action.

**Conclusion / delta:** O’Neil and automation-accountability work show that the model may be destructive, opaque, scaled, or over-trusted. RTF shows how caution language itself can become the path to over-reliance and procedural legitimacy, and why audit must follow the recommendation into delegated workflow, downstream handoff, and environment-affecting action.

### Appendix B — Diagrammatic Aids

The figures in this appendix are reader aids rather than methods. They do not replace the prose definitions. They map relations of direction, scale, compression, or recurrence that are distributed across the argument.

**Figure 2. Threshold family map**  
 The same event may occupy several thresholds at once.



**Diagnostic grammar, not mutually exclusive ontological kinds.**

Figure 2. Threshold family map. The map does not replace the prose definitions. It locates the family of intervals

named in Section 1 across temporal position and scale: pre-action, counter-threshold, machine, room, democratic, and post-threshold. A single event may occupy several thresholds at once; the figure distinguishes loci and modes of binding rather than mutually exclusive ontological kinds. The point is not taxonomy, but whether judgment can still return before the frame becomes binding.

**Figure 3. Dynamic frame cycle**  
 The cycle keeps correction from becoming a room; it is not a mechanical procedure.





Figure 3. Dynamic frame cycle. This is not a mechanical procedure. The phases can overlap, repeat, or compress. The figure marks where a frame can fail while being reconstructed: construction, working, examination, destruction, rebuilding, and consultation. Consultation counts only where it can damage the frame rather than merely witness it. The diagram is a map of recurrence, not a checklist or institutional protocol.

**Figure 5. Distributed realisation of binding**



This is a diagnostic map, not an agent architecture or compliance workflow.

Figure 5. Distributed realisation of binding. Recommendation, permission, tool use, handoff, execution, and effect are not a universal procedure. They are possible loci at which alternatives can become operationally remote. Human-lived

*judgment may be absent at machine speed; an architected locus of interruption or reversal may still remain. Where neither remains, analysis moves upstream to authorisation and design and downstream to post-threshold repair. The figure is a diagnostic map, not an agent architecture or compliance workflow.*

## References and conceptual kinships

The references below are grouped as directly discussed sources, AI and automation-accountability sources, technical parallels, and adjacent kinships. Inclusion in the latter categories does not mean that every work functions as a foundation of the argument. The distinction is meant to preserve source hygiene: a source should either do visible work in the essay, name a genuine adjacent room, or remain a technical parallel rather than a silent authority claim.

### Directly discussed sources and core conceptual kinships

This essay is a companion to *The Architecture of the Liminal* and uses several terms developed there in literary, symbolic, somatic, and machine-facing registers.

#### Classical and early-modern interlocutors

Plato. *Republic*, Book VII, 514a-520a, the Cave Allegory. Translated by Paul Shorey in *Plato in Twelve Volumes*, Vols. 5-6. Harvard University Press, 1969.

Plato. *Republic*, Book VIII, approximately 557a-569c, the sequence from democracy to tyranny and the democratic/tyrannical soul. Translated by Paul Shorey in *Plato in Twelve Volumes*, Vols. 5-6. Harvard University Press, 1969.

Spinoza, Benedict de. *Ethics*. In *The Collected Works of Spinoza, Volume I*. Translated by Edwin Curley. Princeton University Press, 1985.

Leibniz, G. W. “Discourse on Metaphysics” and “Monadology.” In *Philosophical Essays*. Translated and edited by Roger Ariew and Daniel Garber. Hackett, 1989.

#### Core conceptual kinships

Amoore, Louise. *The Politics of Possibility: Risk and Security Beyond Probability*. Duke University Press, 2013.

Massumi, Brian. *Ontopower: War, Powers, and the State of Perception*. Duke University Press, 2015.

Rouvroy, Antoinette, and Thomas Berns. “Algorithmic Governmentality and Prospects of Emancipation: Disparateness as a Precondition for Individuation through Relationships?” *Réseaux* 177, no. 1 (2013): 163-196. <https://doi.org/10.3917/res.177.0163>.

Streeck, Wolfgang, and Kathleen Thelen, eds. *Beyond Continuity: Institutional Change in Advanced Political Economies*. Oxford University Press, 2005.

Virilio, Paul. *Speed and Politics: An Essay on Dromology*. Translated by Mark Polizzotti. Semiotext(e), 1986.

Virilio, Paul. *Open Sky*. Translated by Julie Rose. Verso, 1997.

Adorno, Theodor W. *Negative Dialectics*. Translated by E. B. Ashton. Continuum, 1973.

Agamben, Giorgio. “What Is an Apparatus?” In *What Is an Apparatus? and Other Essays*. Stanford University Press, 2009.

Althusser, Louis. “Ideology and Ideological State Apparatuses.” In *Lenin and Philosophy and Other Essays*. Monthly Review Press, 1971.

Arendt, Hannah. *The Human Condition*. University of Chicago Press, 1958.

Arendt, Hannah. *The Origins of Totalitarianism*. Harcourt, 1951.

Beer, Stafford. *Brain of the Firm*. 2nd ed. John Wiley, 1981.

Beer, Stafford. *Diagnosing the System for Organizations*. John Wiley, 1985.

Bourdieu, Pierre. *Outline of a Theory of Practice*. Translated by Richard Nice. Cambridge University Press, 1977.

- Brunsson, Nils. *The Organization of Hypocrisy: Talk, Decisions and Actions in Organizations*. 2nd ed. Copenhagen Business School Press, 2002.
- Couldry, Nick, and Ulises A. Mejias. *The Costs of Connection: How Data Is Colonizing Human Life and Appropriating It for Capitalism*. Stanford University Press, 2019. <https://doi.org/10.1515/9781503609754>.
- Derrida, Jacques. "Faith and Knowledge: The Two Sources of 'Religion' at the Limits of Reason Alone." In *Religion*, edited by Jacques Derrida and Gianni Vattimo. Stanford University Press, 1998.
- Dewey, John. *Reconstruction in Philosophy*. Henry Holt and Company, 1920.
- Foucault, Michel. *Discipline and Punish: The Birth of the Prison*. Translated by Alan Sheridan. Vintage, 1977.
- Foucault, Michel. *The History of Sexuality, Volume 1: An Introduction*. Translated by Robert Hurley. Pantheon, 1978.
- Frankfurt, Harry G. "Freedom of the Will and the Concept of a Person." *The Journal of Philosophy* 68, no. 1 (1971): 5-20.
- Frankfurt, Harry G. *On Bullshit*. Princeton University Press, 2005.
- Fricker, Miranda. *Epistemic Injustice: Power and the Ethics of Knowing*. Oxford University Press, 2007.
- Goffman, Erving. *Frame Analysis: An Essay on the Organization of Experience*. Harvard University Press, 1974.
- Han, Byung-Chul. *The Burnout Society*. Stanford University Press, 2015.
- Han, Byung-Chul. *Saving Beauty*. Translated by Daniel Steuer. Polity, 2017.
- Heidegger, Martin. "The Question Concerning Technology." In *The Question Concerning Technology and Other Essays*. Translated by William Lovitt. Harper & Row, 1977.
- Horkheimer, Max, and Theodor W. Adorno. *Dialectic of Enlightenment*. Translated by Edmund Jephcott. Stanford University Press, 2002.
- Illich, Ivan. *Tools for Conviviality*. Harper & Row, 1973.
- Jonas, Hans. *The Imperative of Responsibility: In Search of an Ethics for the Technological Age*. University of Chicago Press, 1984.
- Lakatos, Imre. "Falsification and the Methodology of Scientific Research Programmes." In *Criticism and the Growth of Knowledge*, edited by Imre Lakatos and Alan Musgrave. Cambridge University Press, 1970.
- Levinas, Emmanuel. *Totality and Infinity: An Essay on Exteriority*. Translated by Alphonso Lingis. Duquesne University Press, 1969.
- Marx, Karl. *Capital: A Critique of Political Economy, Volume I*. Translated by Ben Fowkes. Penguin, 1976.
- Merleau-Ponty, Maurice. *Phenomenology of Perception*. Translated by Donald A. Landes. Routledge, 2012.
- Michels, Robert. *Political Parties: A Sociological Study of the Oligarchical Tendencies of Modern Democracy*. Translated by Eden and Cedar Paul. Hearst's International Library, 1915. Originally published in German in 1911.
- O'Neil, Cathy. *Weapons of Math Destruction*. Crown, 2016.
- Popper, Karl. *The Open Society and Its Enemies*. Routledge, 1945.
- Power, Michael. *The Audit Society: Rituals of Verification*. Oxford University Press, 1997.
- Schön, Donald A. *The Reflective Practitioner: How Professionals Think in Action*. Basic Books, 1983.
- Scott, James C. *Seeing Like a State: How Certain Schemes to Improve the Human Condition Have Failed*. Yale University Press, 1998.
- von Foerster, Heinz. "Ethics and Second-Order Cybernetics." *Cybernetics & Human Knowing* 1, no. 1 (1992): 9-19.
- Weber, Max. *Economy and Society: An Outline of Interpretive Sociology*. Edited by Guenther Roth and Claus Wittich. University of California Press, 1978.

Weil, Simone. “Reflections on the Right Use of School Studies with a View to the Love of God.” In *Waiting for God*. Translated by Emma Craufurd. Harper & Row, 1951.

Williams, Bernard. *Moral Luck*. Cambridge University Press, 1981.

Wittgenstein, Ludwig. *Philosophical Investigations*. Translated by G. E. M. Anscombe. Blackwell, 1953.

## AI and automation-accountability sources

Anthropic. *Responsible Scaling Policy, Version 3.0*. Anthropic, 2026.

Anthropic Safeguards Research Team. *Constitutional Classifiers: Defending against Universal Jailbreaks across Thousands of Hours of Red Teaming*. Anthropic, 2025.

Almada, Marco. “Human Intervention in Automated Decision-Making: Toward the Construction of Contestable Systems.” In *Proceedings of the Seventeenth International Conference on Artificial Intelligence and Law (ICAIL ’19)*, 2-11. ACM, 2019. <https://doi.org/10.1145/3322640.3326699>.

Bai, Yuntao, et al. “Constitutional AI: Harmlessness from AI Feedback.” arXiv:2212.08073, 2022.

Kirk, Hannah Rose, Bertie Vidgen, Paul Röttger, and Scott A. Hale. “Personalisation within Bounds: A Risk Taxonomy and Policy Framework for the Alignment of Large Language Models with Personalised Feedback.” arXiv:2303.05453, 2023.

Bengio, Yoshua, et al. *International AI Safety Report 2026*. International AI Safety Report, 2026.

European Commission. *Code of Practice on the Marking and Labelling of AI-Generated Content*. European Commission, 2025.

Calvano, Emilio, Giacomo Calzolari, Vincenzo Denicolò, and Sergio Pastorello. “Artificial Intelligence, Algorithmic Pricing, and Collusion.” *American Economic Review* 110, no. 10 (2020): 3267-3297. <https://doi.org/10.1257/aer.20190623>.

U.S. Commodity Futures Trading Commission and U.S. Securities and Exchange Commission. *Findings Regarding the Market Events of May 6, 2010: Report of the Staffs of the CFTC and SEC to the Joint Advisory Committee on Emerging Regulatory Issues*. September 30, 2010. <https://www.sec.gov/files/marketevents-report.pdf>.

Elish, Madeleine Clare. “Moral Crumple Zones: Cautionary Tales in Human-Robot Interaction.” *Engaging Science, Technology, and Society* 5 (2019): 40-60. <https://doi.org/10.17351/ests2019.260>.

European Union. *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act)*. Official Journal of the European Union, 2024.

European Data Protection Supervisor. *TechDispatch #2/2025: Human Oversight of Automated Decision-Making*. EDPS, 2025.

European Union. *Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the Protection of Natural Persons with Regard to the Processing of Personal Data and on the Free Movement of Such Data (General Data Protection Regulation)*. Official Journal of the European Union, 2016. See Article 22.

Financial Stability Board. *The Financial Stability Implications of Artificial Intelligence*. FSB, 2024.

Floridi, Luciano, and J. W. Sanders. “On the Morality of Artificial Agents.” *Minds and Machines* 14 (2004): 349-379. <https://doi.org/10.1023/B:MIND.0000035461.63578.9d>.

National Institute of Standards and Technology. *Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile*. NIST AI 600-1. U.S. Department of Commerce, 2024. <https://doi.org/10.6028/NIST.AI.600-1>.

OECD. *Explanatory Memorandum on the Updated OECD Definition of an AI System*. OECD Artificial Intelligence Papers, no. 8. OECD Publishing, 2024. <https://doi.org/10.1787/623da898-en>.

Orseau, Laurent, and Stuart Armstrong. “Safely Interruptible Agents.” In *Proceedings of the Thirty-Second Conference on Uncertainty in Artificial Intelligence*, 557-566. AUAI Press, 2016.

Santoni de Sio, Filippo, and Jeroen van den Hoven. “Meaningful Human Control over Autonomous Systems: A Philosophical Account.” *Frontiers in Robotics and AI* 5 (2018): 15. <https://doi.org/10.3389/frobt.2018.00015>.

Buçinca, Zana, Maja Barbara Malaya, and Krzysztof Z. Gajos. “To Trust or to Think: Cognitive Forcing Functions Can Reduce Overreliance on AI in AI-Assisted Decision-Making.” *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW1, Article 188 (2021). <https://doi.org/10.1145/3449287>.

De-Arteaga, Maria, Riccardo Fogliato, and Alexandra Chouldechova. “A Case for Humans-in-the-Loop: Decisions in the Presence of Erroneous Algorithmic Scores.” In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*. ACM, 2020. <https://doi.org/10.1145/3313831.3376638>.

Eubanks, Virginia. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. St. Martin’s Press, 2018.

Allegheny County Department of Human Services. *Allegheny Family Screening Tool*. Public documentation.

American Civil Liberties Union. *The Devil Is in the Details: Interrogating Values Embedded in the Allegheny Family Screening Tool*. ACLU, 2023.

Chicago Office of Inspector General. *Advisory Concerning the Chicago Police Department’s Predictive Risk Models*. City of Chicago, January 23, 2020.

Passi, Samir, and Solon Barocas. “Problem Formulation and Fairness.” In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT ’19)\**, 39-48. ACM, 2019. <https://doi.org/10.1145/3287560.3287567>.

Selbst, Andrew D., danah boyd, Sorelle A. Friedler, Suresh Venkatasubramanian, and Janet Vertesi. “Fairness and Abstraction in Sociotechnical Systems.” In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT ’19)\**, 59-68. ACM, 2019. <https://doi.org/10.1145/3287560.3287598>.

Vaithianathan, Rhema, Emily Putnam-Hornstein, Alexandra Chouldechova, Diana Benavides-Prado, and Rachel P. Berger. “Hospital Injury Encounters of Children Identified by a Predictive Risk Model for Screening Child Maltreatment Referrals: Evidence from the Allegheny Family Screening Tool.” *JAMA Pediatrics* 174, no. 11 (2020): e202770. <https://doi.org/10.1001/jamapediatrics.2020.2770>.

Wachter, Sandra, Brent Mittelstadt, and Chris Russell. “Counterfactual Explanations without Opening the Black Box: Automated Decisions and the GDPR.” *Harvard Journal of Law & Technology* 31, no. 2 (2018): 841-887.

Green, Ben, and Yiling Chen. “Disparate Interactions: An Algorithm-in-the-Loop Analysis of Fairness in Risk Assessments.” In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT ’19)*, 90-99. ACM, 2019. <https://doi.org/10.1145/3287560.3287563>.

Green, Ben. “The Flaws of Policies Requiring Human Oversight of Government Algorithms.” *Computer Law & Security Review* 45 (2022): 105681. <https://doi.org/10.1016/j.clsr.2022.105681>.

Saxena, Devansh, Karla Badillo-Urquiola, Pamela J. Wisniewski, and Shion Guha. “A Framework of High-Stakes Algorithmic Decision-Making for the Public Sector Developed through a Case Study of Child-Welfare.” *Proceedings of the ACM on Human-Computer Interaction* 5, CSCW2, Article 440 (2021). <https://doi.org/10.1145/3476089>.

## Technical parallels and background vocabulary

Ashby, W. Ross. *An Introduction to Cybernetics*. Chapman & Hall, 1956.

Benedetti, A. “A Literature Review on ‘Friction’ as a Method for Critical Design.” *Convergências* 16, no. 31 (2023): 137-144.

Cox, Anna L., Sandy J. J. Gould, Marta E. Cecchinato, Ioanna Iacovides, and Ian Renfree. “Design Frictions for Mindful Interactions: The Case for Microboundaries.” In *Proceedings of the 2016 CHI Conference Extended Abstracts on Human Factors in Computing Systems*, 1389-1397. ACM, 2016. <https://doi.org/10.1145/2851581.2892410>.

Hollan, James, Edwin Hutchins, and David Kirsh. “Distributed Cognition: Toward a New Foundation for Human-Computer Interaction Research.” *ACM Transactions on Computer-Human Interaction* 7, no. 2 (2000): 174-196.

Maturana, Humberto R., and Francisco J. Varela. *Autopoiesis and Cognition: The Realization of the Living*. D. Reidel, 1980.

Pearl, Judea. *Causality: Models, Reasoning, and Inference*. 2nd ed. Cambridge University Press, 2009.

Wiener, Norbert. *Cybernetics: Or Control and Communication in the Animal and the Machine*. MIT Press, 1948.

### **Adjacent literary and philosophical kinships**

Bakhtin, Mikhail M. “Forms of Time and of the Chronotope in the Novel.” In *The Dialogic Imagination: Four Essays*. University of Texas Press, 1981.

Baudrillard, Jean. *Simulacra and Simulation*. Translated by Sheila Faria Glaser. University of Michigan Press, 1994.

### **AI Disclosure Statement**

AI systems were used as interlocutors rather than co-authors: to test structure, identify repetition, compare versions, surface risks of misreading, challenge weak formulations, and propose alternative phrasings. Several machine-generated suggestions were rejected, narrowed, or re-framed during revision. This disclosure is not a credential; it marks a site where the reader can examine whether assistance returned judgment or occupied it.

Declaration of Interest: The author declares no competing financial or personal interests that could have influenced the analysis or conclusions of this essay.

Funder Statement: This research received no external funding.

Repository note: Canonical repository: <https://github.com/emreertuhi/architecture-of-the-liminal>.